

Identification Capacity and Rate-Query Tradeoffs in Classification Systems

Tristan Simas

Abstract—We extend classical rate-distortion theory to a discrete classification setting with three resources: *tag rate* L (bits of storage per entity), *identification cost* W (queries to determine class membership), and *distortion* D (misidentification probability). The fundamental question is: how do these resources trade off when an observer must identify an entity’s class from limited observations?

Information barrier (zero-error identifiability). When distinct classes share identical attribute profiles, no algorithm, regardless of computational power, can identify the class from attribute queries alone. Formally: if π is not injective on classes, then zero-error identification from attribute queries alone is impossible.

Rate-identification tradeoff. We identify the unique Pareto-optimal zero-error point in the (L, W, D) tradeoff space and describe the induced tradeoff geometry. A nominal tag of $L = \lceil \log_2 k \rceil$ bits for k classes yields $W = O(1)$ with $D = 0$. Without tags ($L = 0$), zero-error identification requires $W = \Omega(d)$ attribute queries, where d is the distinguishing dimension; in the worst case $d = n$ (the ambient attribute count), giving $W = \Omega(n)$. In the presence of attribute collisions, any tag-free scheme incurs $D > 0$.

Converse. In any information-barrier domain, any scheme achieving $D = 0$ requires $L \geq \log_2 k$ bits. This is tight: nominal tagging achieves the bound with $W = O(1)$.

Matroid structure. Minimal sufficient query sets form the bases of a matroid. The *distinguishing dimension* (the common cardinality of all minimal query sets) is well-defined, connecting to zero-error source coding via graph entropy.

Applications. The theory instantiates to type systems (duck vs. nominal typing), databases (attribute vs. key lookup), and biological taxonomy (phenotype vs. species identifier). The unbounded gap $\Omega(d)$ vs. $O(1)$ (with a worst-case family where $d = n$) explains convergence toward hybrid systems combining structural observation with nominal tagging.

All results are machine-checked in Lean 4 (6,000+ lines, 0 sorry).

Keywords: rate-distortion theory, identification capacity, zero-error source coding, query complexity, matroid structure, classification systems

I. INTRODUCTION

A. The Identification Problem

Consider an encoder-decoder pair communicating about entities from a large universe \mathcal{V} . The decoder must *identify* each entity, determining which of k classes it belongs to, using only:

- A *tag* of L bits stored with the entity, and/or
- *Queries* to a binary oracle: “does entity v satisfy attribute I ?“

T. Simas is with McGill University, Montreal, QC, Canada (e-mail: tristan.simas@mail.mcgill.ca).

Manuscript received January 15, 2026.

© 2026 Tristan Simas. This work is licensed under CC BY 4.0. License: <https://creativecommons.org/licenses/by/4.0/>

This is not reconstruction (the decoder need not recover v), but *identification* in the sense of Ahlswede and Dueck [1]: the decoder must answer “which class?” with zero or bounded error. Our work extends this framework to consider the trade-off between tag storage, query complexity, and identification accuracy.

We prove three results:

- 1) **Information barrier (identifiability limit).** When the attribute profile $\pi : \mathcal{V} \rightarrow \{0,1\}^n$ is not injective on classes, zero-error identification via queries alone is impossible: any decoder produces identical output on colliding classes, so cannot be correct for both.
- 2) **Optimal tagging (achievability).** A tag of $L = \lceil \log_2 k \rceil$ bits achieves zero-error identification with $W = O(1)$ query cost. This is the unique Pareto-optimal point in the (L, W, D) tradeoff space for $D = 0$ in the information-barrier regime.
- 3) **Matroid structure (query complexity).** Minimal sufficient query sets form the bases of a matroid. The *distinguishing dimension* (the common cardinality of all minimal sets) is well-defined and lower-bounds the query cost W for any tag-free scheme.

These results are universal: the theory applies to type systems, databases, biological taxonomy, and knowledge graphs. We develop the mathematics in full generality, then exhibit concrete instantiations.

B. The Observation Model

We formalize the observational constraint as a family of binary predicates. The terminology is deliberately abstract; concrete instantiations follow in Section VIII.

Definition I.1 (Entity space and attribute family). Let \mathcal{V} be a set of entities (program objects, database records, biological specimens, library items). Let \mathcal{I} be a finite set of *attributes*: observable properties that partition the entity space.

Remark I.2 (Terminology). We use “attribute” for the abstract concept. In type systems, attributes are *interfaces* or *method signatures*. In databases, they are *columns*. In taxonomy, they are *phenotypic characters*. In library science, they are *facets*. The mathematics is identical.

Definition I.3 (Interface observation family). For each $I \in \mathcal{I}$, define the interface-membership observation $q_I : \mathcal{V} \rightarrow \{0,1\}$:

$$q_I(v) = \begin{cases} 1 & \text{if } v \text{ satisfies interface } I \\ 0 & \text{otherwise} \end{cases}$$

Let $\Phi_{\mathcal{I}} = \{q_I : I \in \mathcal{I}\}$ denote the interface observation family.

Remark I.4 (Notation for size parameters). We write $n := |\mathcal{I}|$ for the ambient number of available attributes (interfaces). We write d for the distinguishing dimension (the common size of all minimal distinguishing query sets; Definition V.9), so $d \leq n$ and there exist worst-case families with $d = n$. We write m for the number of *query sites* (call sites) that perform attribute checks in a program or protocol (used only in the complexity-of-maintenance discussion). When discussing a particular identification/verification task, we may write s for the number of attributes actually queried/traversed by the procedure (e.g., members/fields checked in a structural type test, phenotypic characters checked in taxonomy), with $s \leq n$.

Definition I.5 (Interface profile). The interface profile function $\pi : \mathcal{V} \rightarrow \{0, 1\}^{|\mathcal{I}|}$ maps each value to its complete interface signature:

$$\pi(v) = (q_I(v))_{I \in \mathcal{I}}$$

Definition I.6 (Interface indistinguishability). Values $v, w \in \mathcal{V}$ are *interface-indistinguishable*, written $v \sim w$, iff $\pi(v) = \pi(w)$.

The relation \sim is an equivalence relation. We write $[v]_\sim$ for the equivalence class of v .

Definition I.7 (Interface-only observer). An *interface-only observer* is any procedure whose interaction with a value $v \in \mathcal{V}$ is limited to queries in $\Phi_{\mathcal{I}}$. Formally, the observer interacts with v only via primitive interface queries $q_I \in \Phi_{\mathcal{I}}$; hence any transcript (and output) factors through $\pi(v)$.

C. The Central Question

The central question is: **what semantic properties can an interface-only observer compute?**

A semantic property is a function $P : \mathcal{V} \rightarrow \{0, 1\}$ (or more generally, $P : \mathcal{V} \rightarrow Y$ for some codomain Y). We say P is *interface-computable* if there exists a function $f : \{0, 1\}^{|\mathcal{I}|} \rightarrow Y$ such that $P(v) = f(\pi(v))$ for all v .

D. The Information Barrier

Theorem I.8 (Information barrier). Let $P : \mathcal{V} \rightarrow Y$ be any function. If P is interface-computable, then P is constant on \sim -equivalence classes:

$$v \sim w \implies P(v) = P(w)$$

Equivalently: no interface-only observer can compute any property that varies within an equivalence class.

Proof. Suppose P is interface-computable via f , i.e., $P(v) = f(\pi(v))$ for all v . Let $v \sim w$, so $\pi(v) = \pi(w)$. Then:

$$P(v) = f(\pi(v)) = f(\pi(w)) = P(w)$$

■

Remark I.9 (Information-theoretic nature). The barrier is *informational*, not computational. Given unlimited time, memory, and computational power, an interface-only observer still cannot distinguish v from w when $\pi(v) = \pi(w)$. The constraint is on the evidence itself.

Corollary I.10 (Provenance is not interface-computable). Let $\text{type} : \mathcal{V} \rightarrow \mathcal{T}$ be the type assignment function. If there exist values v, w with $\pi(v) = \pi(w)$ but $\text{type}(v) \neq \text{type}(w)$, then type identity is not interface-computable.

Proof. Direct application of Theorem I.8 to $P = \text{type}$. ■

E. The Positive Result: Nominal Tagging

We now show that augmenting interface observations with a single primitive, nominal-tag access, achieves constant witness cost.

Definition I.11 (Nominal-tag access). A *nominal tag* is a value $\tau(v) \in \mathcal{T}$ associated with each $v \in \mathcal{V}$, representing the type identity of v . The *nominal-tag access* operation returns $\tau(v)$ in $O(1)$ time.

Definition I.12 (Primitive query set). The extended primitive query set is $\Phi_{\mathcal{I}}^+ = \Phi_{\mathcal{I}} \cup \{\tau\}$, where τ denotes nominal-tag access.

Definition I.13 (Witness cost). Let $W(P)$ denote the minimum number of primitive queries from $\Phi_{\mathcal{I}}^+$ required to compute property P . We distinguish two tasks:

- W_{id} : Cost to identify the class of a single entity.
- W_{eq} : Cost to determine if two entities have the same class (type identity).

Unless specified, W refers to W_{eq} (type identity checking).

Theorem I.14 (Constant witness for type identity). Under nominal-tag access, type identity checking has constant witness cost:

$$W(\text{type-identity}) = O(1)$$

Specifically, the witness procedure is: return $\tau(v_1) = \tau(v_2)$.

Proof. The procedure makes exactly 2 primitive queries (one τ access per value) and one comparison. This is $O(1)$ regardless of the number of interfaces $|\mathcal{I}|$. ■

Theorem I.15 (Interface-only lower bound). For interface-only observers, type identity checking requires:

$$W(\text{type-identity}) = \Omega(d)$$

in the worst case, where d is the distinguishing dimension (Definition V.9).

Proof. By definition of the distinguishing dimension, any witness procedure must query at least one minimal distinguishing set of size d . Therefore the worst-case query cost is $\Omega(d)$. ■

F. Main Contributions

This paper establishes the following results:

- 1) **Information Barrier Theorem** (Theorem I.8): Interface-only observers cannot compute any property that varies within \sim -equivalence classes. This is an information-theoretic impossibility, not a computational limitation.
- 2) **Constant-Witness Theorem** (Theorem I.14): Nominal-tag access achieves $W(\text{type-identity}) = O(1)$, with matching lower bound $\Omega(d)$ for interface-only observers

- (Theorem I.15), where d is the distinguishing dimension (Definition V.9).
- 3) **Complexity Separation** (Section III): We establish $O(1)$ vs $O(k)$ vs $\Omega(d)$ complexity bounds for error localization under different observation regimes (where d is the distinguishing dimension).
 - 4) **Matroid Structure** (Section V): Minimal distinguishing query sets form the bases of a matroid. All such sets have equal cardinality, establishing a well-defined “distinguishing dimension.”
 - 5) **(L, W, D) Optimality** (Section VII): Nominal-tag observers achieve the unique Pareto-optimal point in the (L, W, D) tradeoff space (tag length, witness cost, distortion).
 - 6) **Machine-Checked Proofs:** All results formalized in Lean 4 (6,086 lines, 265 theorems, 0 sorry placeholders).

G. Related Work and Positioning

Identification via channels. Our work extends the identification paradigm introduced by Ahlswede and Dueck [1], [2]. In their framework, a decoder need not reconstruct a message but only answer “is the message m ?” for a given hypothesis. This yields dramatically different capacity: double-exponential codebook sizes become achievable. Our setting differs in three ways: (1) we consider zero-error identification rather than vanishing error, (2) queries are adaptive rather than block codes, and (3) we allow auxiliary tagging (rate L) to reduce query cost. The (L, W, D) tradeoff generalizes Ahlswede-Dueck to a multi-dimensional operating regime.

Rate-distortion theory. The (L, W, D) framework connects to Shannon’s rate-distortion theory [3], [4] with an important twist: the “distortion” D is semantic (class misidentification), and there is a second resource W (query cost) alongside rate L . Classical rate-distortion asks: what is the minimum rate to achieve distortion D ? We ask: given rate L , what is the minimum query cost W to achieve distortion $D = 0$? Theorem VII.3 identifies the unique Pareto-optimal point in this three-dimensional tradeoff at $D = 0$.

Rate-distortion-perception tradeoffs. Blau and Michaeli [5] extended rate-distortion theory by adding a perception constraint, creating a three-way tradeoff. Our query cost W plays an analogous role: it measures the interactive cost of achieving low distortion rather than a distributional constraint. This parallel suggests that (L, W, D) tradeoffs may admit similar geometric characterizations. Section IX develops this connection further.

Zero-error information theory. The matroid structure (Section V) connects to zero-error capacity and graph entropy. Körner [6] and Witsenhausen [7] studied zero-error source coding where confusable symbols must be distinguished. Our distinguishing dimension (Definition V.9) is the minimum number of binary queries to separate all classes, which is precisely the zero-error identification cost when $L = 0$.

extbf{Query complexity} and communication complexity. The $\Omega(d)$ lower bound for interface-only identification relates to decision tree complexity [8] and interactive communication

[9]. The key distinction is that our queries are constrained to a fixed attribute family \mathcal{I} , not arbitrary predicates. This constraint models practical systems where the observer’s interface to entities is architecturally fixed.

Compression in classification systems. Our framework instantiates to type systems, where the compression question becomes: how many bits must be stored per object to enable $O(1)$ type identification? The answer ($\lceil \log_2 k \rceil$ bits for k classes) matches the converse bound (Theorem II.27). This provides an information-theoretic foundation for the nominal-vs-structural typing debate in programming language theory [10], [11].

extbf{Historical context.} The “duck typing” philosophy (“if it walks like a duck and quacks like a duck, it’s a duck”) advocates structural observation over nominal tagging. Within our model, we prove this incurs $\Omega(d)$ witness cost where tagging achieves $O(1)$ (see Definition V.9). The result does not “resolve” the broader debate (which involves usability and tooling concerns beyond this model) but establishes that the tradeoff has a precise information-theoretic component.

Practical convergence. Modern systems have converged on hybrid classification: Python’s Abstract Base Classes, TypeScript’s branded types, Rust’s trait system, DNA barcoding in taxonomy [12]. This convergence is consistent with the rate-query tradeoff: nominal tags provide $O(1)$ identification at cost $O(\log k)$ bits. The contribution is not advocacy for any design, but a formal framework for analyzing identification cost in classification systems.

H. Paper Organization

Section II formalizes the compression framework and defines the (L, W, D) tradeoff. Section III establishes complexity bounds for error localization. Section V proves the matroid structure of type axes. Section VI analyzes witness cost in detail. Section VII proves Pareto optimality. Section VIII instantiates the theory in real runtimes. Section X concludes. Appendix A describes the Lean 4 formalization.

II. COMPRESSION FRAMEWORK

A. Semantic Compression: The Problem

The fundamental problem of *semantic compression* is: given a value v from a large space \mathcal{V} , how can we represent v compactly while preserving the ability to answer semantic queries about v ? This differs from classical source coding in that the goal is not reconstruction but *identification*: determining which equivalence class v belongs to.

Classical rate-distortion theory [3] studies the tradeoff between representation size and reconstruction fidelity. We extend this to a discrete classification setting with three dimensions: *tag length* L (bits of storage), *witness cost* W (queries or bits of communication required to determine class membership), and *distortion* D (semantic fidelity).

B. Universe of Discourse

Definition II.1 (Classification scheme). A *classification scheme* is any procedure (deterministic or randomized), with

arbitrary time and memory, whose only access to a value $v \in \mathcal{V}$ is via:

- 1) The *observation family* $\Phi = \{q_I : I \in \mathcal{I}\}$, where $q_I(v) = 1$ iff v satisfies attribute I ; and optionally
- 2) A *nominal-tag primitive* $\tau : \mathcal{V} \rightarrow \mathcal{T}$ returning an opaque type identifier.

All theorems in this paper quantify over all such schemes.

This definition is intentionally broad: schemes may be adaptive, randomized, or computationally unbounded. The constraint is *observational*, not computational.

Theorem II.2 (Information barrier). *For all classification schemes with access only to Φ (no nominal tag), the output is constant on \sim_Φ -equivalence classes. Therefore, no such scheme can compute any property that varies within a \sim_Φ -class.*

Proof. Let $v \sim_\Phi w$, meaning $q_I(v) = q_I(w)$ for all $I \in \mathcal{I}$. Any scheme's execution trace depends only on query responses. Since all queries return identical values for v and w , the scheme cannot distinguish them. Any output must therefore be identical. ■

Proposition II.3 (Model capture). *Any real-world classification protocol whose evidence consists solely of attribute-membership queries is representable as a scheme in the above model. Conversely, any additional capability corresponds to adding new observations to Φ .*

This proposition forces any objection into a precise form: to claim the theorem does not apply, one must name the additional observation capability not in Φ . “Different universe” is not a coherent objection; it must reduce to “I have access to oracle $X \notin \Phi$.”

C. The Two-Axis Model

We adopt a two-axis model of semantic structure, where each value is characterized by:

- **Lineage axis (B):** The provenance chain of the value’s class (which classes it derives from, in what order)¹
- **Profile axis (S):** The observable interface (which attributes/methods the value provides)

Definition II.4 (Two-axis representation). A value $v \in \mathcal{V}$ has representation $(B(v), S(v))$ where:

$$B(v) = \text{lineage(class}(v)\text{)} \quad (\text{class derivation chain}) \quad (1)$$

$$S(v) = \pi(v) = (q_I(v))_{I \in \mathcal{I}} \quad (\text{interface profile}) \quad (2)$$

The lineage axis captures *nominal* identity: where the class comes from. The profile axis captures *structural* identity: what the value can do.

In the PL instantiation, B is carried by the runtime lineage order (e.g., C3/MRO output), while OpenHCS additionally uses a separate normalization registry R applied before lookup and not defining inheritance (Appendix A).

¹In the Lean formalization (Appendix A), the lineage axis is denoted Bases , reflecting its instantiation as the inheritance chain in object-oriented languages.

Theorem II.5 (Fixed-axis completeness). *Let a fixed-axis domain be specified by an axis map $\alpha : \mathcal{V} \rightarrow \mathcal{A}$ and an observation interface Φ such that each primitive query $q \in \Phi$ factors through α . Then every in-scope semantic property (i.e., any property computable by an admissible Φ -only strategy) factors through α : there exists \tilde{P} with*

$$P(v) = \tilde{P}(\alpha(v)) \quad \text{for all } v \in \mathcal{V}.$$

In the PL instantiation, $\alpha(v) = (B(v), S(v))$, so in-scope semantic properties are functions of (B, S) .

Proof. An admissible Φ -only strategy observes v solely through responses to primitive queries $q_I \in \Phi$. By hypothesis each such response is a function of $\alpha(v)$. Therefore every query transcript, and hence any strategy's output, depends only on $\alpha(v)$, so the computed property factors through α . ■

D. Interface Equivalence and Observational Limits

Recall from Section 1 the interface equivalence relation:

Definition II.6 (Interface equivalence (restated)). Values $v, w \in \mathcal{V}$ are interface-equivalent, written $v \sim w$, iff $\pi(v) = \pi(w)$, i.e., they satisfy exactly the same interfaces.

Proposition II.7 (Equivalence class structure). *The relation \sim partitions \mathcal{V} into equivalence classes. Let \mathcal{V}/\sim denote the quotient space. An interface-only observer effectively operates on \mathcal{V}/\sim , not \mathcal{V} .*

Corollary II.8 (Information loss quantification). *The information lost by interface-only observation is:*

$$H(\mathcal{V}) - H(\mathcal{V}/\sim) = H(\mathcal{V}|\pi)$$

where H denotes entropy. This quantity is positive whenever multiple types share the same interface profile.

E. Identification Capacity

We now formalize the identification problem in channel-theoretic terms. Let $C : \mathcal{V} \rightarrow \{1, \dots, k\}$ denote the class assignment function, and let $\pi : \mathcal{V} \rightarrow \{0, 1\}^n$ denote the attribute profile.

Definition II.9 (Identification channel). The *identification channel* induced by observation family Φ is the mapping $C \rightarrow \pi(V)$ for a random entity V drawn from distribution P_V over \mathcal{V} . The channel output is the attribute profile; the channel input is implicitly the class $C(V)$.

Theorem II.10 (Identification capacity). *Let $\mathcal{C} = \{1, \dots, k\}$ be the class space. The zero-error identification capacity of the observation channel is:*

$$C_{id} = \begin{cases} \log_2 k & \text{if } \pi \text{ is injective on classes} \\ 0 & \text{otherwise} \end{cases}$$

That is, zero-error identification of all k classes is achievable if and only if every class has a distinct attribute profile. When π is not class-injective, no rate of identification is achievable with zero error.

Proof. Achievability: If π is injective on classes, then observing $\pi(v)$ determines $C(v)$ uniquely. The decoder simply inverts the class-to-profile mapping.

Converse (deterministic): Suppose two distinct classes $c_1 \neq c_2$ share a profile: $\exists v_1 \in c_1, v_2 \in c_2$ with $\pi(v_1) = \pi(v_2)$. Then any decoder $g(\pi(v))$ outputs the same class label on both v_1 and v_2 , so it cannot be correct for both. Hence zero-error identification of all classes is impossible. ■

Remark II.11 (Information-theoretic corollary). Under any distribution with positive mass on both colliding classes, $I(C; \pi(V)) < H(C)$. This is an average-case consequence of the deterministic barrier above.

Remark II.12 (Relation to Ahlswede-Dueck). In the identification paradigm of [1], the decoder asks “is the message m ?” rather than “what is the message?” This yields double-exponential codebook sizes. Our setting is different: we require zero-error identification of the *class*, not hypothesis testing. The one-shot zero-error identification feasibility threshold (π must be class-injective) is binary rather than a rate.

The key insight is that tagging provides a *side channel* that restores identifiability when the attribute channel fails:

Theorem II.13 (Tag-restored capacity). A tag of length $L \geq \lceil \log_2 k \rceil$ bits restores zero-error identification capacity, regardless of whether π is class-injective. The tag provides a noiseless side channel with capacity L bits.

Proof. A nominal tag $\tau : \mathcal{V} \rightarrow \{1, \dots, k\}$ assigns a unique identifier to each class. Reading $\tau(v)$ determines $C(v)$ in $O(1)$ queries, independent of the attribute channel. ■

F. Witness Cost: Query Complexity for Semantic Properties

Definition II.14 (Witness procedure). A witness procedure for property $P : \mathcal{V} \rightarrow Y$ is an algorithm A that:

- 1) Takes as input a value $v \in \mathcal{V}$ (via query access only)
- 2) Makes queries to the primitive set $\Phi_{\mathcal{I}}^+$
- 3) Outputs $P(v)$

Definition II.15 (Witness cost). The witness cost of property P is:

$$W(P) = \min_{A \text{ computes } P} c(A)$$

where $c(A)$ is the worst-case number of primitive queries made by A .

Remark II.16 (Relationship to query complexity). Witness cost is a form of query complexity [8] specialized to semantic properties. Unlike Kolmogorov complexity, W is computable and depends on the primitive set, not a universal machine.

Lemma II.17 (Witness cost lower bounds). For any property P :

- 1) If P is interface-computable: $W(P) \leq |\mathcal{I}|$
- 2) If P varies within some \sim -class: $W(P) = \infty$ for interface-only observers
- 3) With nominal-tag access: $W(\text{type-identity}) = O(1)$

G. The (L, W, D) Tradeoff

We now define the three-dimensional tradeoff space that characterizes observation strategies, using information-theoretic units.

Definition II.18 (Tag rate). For a set of type identifiers (tags) \mathcal{T} with $|\mathcal{T}| = k$, the tag rate L is the minimum number of bits required to encode a type identifier:

$$L \geq \log_2 k \text{ bits per value}$$

For nominal-tag observers, $L = \lceil \log_2 k \rceil$ (optimal prefix-free encoding). For interface-only observers, $L = 0$ (no explicit tag stored). Under a distribution P over types, the expected tag length is $\mathbb{E}[L] \geq H(P)$ by Shannon’s source coding theorem [3].

Definition II.19 (Witness cost (Query/Communication complexity)). The witness cost W is the minimum number of primitive queries (or bits of interactive communication) required for type identity checking:

$$W = \min_{A \text{ decides type-identity}} c(A)$$

where $c(A)$ is the worst-case query count. This is a form of query complexity [8] or interactive identification cost.

Definition II.20 (Behavioral equivalence). Let behavior : $\mathcal{V} \rightarrow \mathcal{B}$ be a function mapping each entity to its observable behavior (the set of responses to all possible operations). Two entities v, w are *behaviorally equivalent*, written $v \equiv w$, iff $\text{behavior}(v) = \text{behavior}(w)$.

In type systems, $\text{behavior}(v)$ is the denotational semantics: the function computed by v . In databases, it is the set of query results. In taxonomy, it is the phenotype. The formalism is parametric in this choice.

Definition II.21 (Distortion function). Let $d : \mathcal{V} \times \mathcal{V} \rightarrow \{0, 1\}$ be the misclassification indicator:

$$d(v, \hat{v}) = \begin{cases} 0 & \text{if } \text{type}(v) = \text{type}(\hat{v}) \Rightarrow v \equiv \hat{v} \\ 1 & \text{otherwise (semantic error)} \end{cases}$$

That is, $d = 0$ when type equality implies behavioral equivalence (soundness), and $d = 1$ when the observer conflates behaviorally distinct entities.

Definition II.22 (Expected and worst-case distortion). Given a distribution P over values, the expected distortion is:

$$D = \mathbb{E}_{v \sim P}[d(v, \hat{v})]$$

The zero-error regime requires $D = 0$ (no semantic errors for any v). All theorems in this paper are proved in the zero-error regime, the strongest case where the separation is sharpest.

Remark II.23 (Distortion interpretation). $D = 0$ means the observation strategy is *sound*: type equality (as computed by the observer) implies behavioral equivalence. $D > 0$ means the strategy may conflate behaviorally distinct values with positive probability. We use D as expected misclassification rate, and in the core theorems we focus on the zero-error regime $D = 0$ vs $D > 0$; richer distortion structures are discussed in Section IX.

H. The (L, W, D) Tradeoff Space

Admissible schemes. To make the Pareto-optimality claim precise, we specify the class of admissible observation strategies:

- **Deterministic or randomized:** Schemes may use randomness; W is worst-case query count.
- **Computationally unbounded:** No time/space restrictions; the constraint is observational.
- **No preprocessing over type universe:** The scheme cannot precompute a global lookup table indexed by all possible types.
- **Tags are injective on classes:** A nominal tag $\tau(v)$ uniquely identifies the type of v . Variable-length or compressed tags are permitted; L counts bits.
- **No amortization across queries:** W is per-identification cost, not amortized over a sequence.

Justification. The “no preprocessing” and “no amortization” constraints exclude trivializations:

- *Preprocessing:* With unbounded preprocessing over the type universe \mathcal{T} , one could build a lookup table mapping attribute profiles to types. This reduces identification to $O(1)$ table lookup, but the table has size $O(|\mathcal{T}|)$, hiding the complexity in space rather than eliminating it. The constraint models systems that cannot afford $O(|\mathcal{T}|)$ storage per observer.
- *Amortization:* If W were amortized over a sequence of identifications, one could cache earlier results. This again hides complexity in state. The per-identification model captures stateless observers (typical in type checking, database queries, and biological identification).

Dropping these constraints changes the achievable region but not the qualitative separation: nominal tags still dominate for $D = 0$ because they provide $O(1)$ worst-case identification without requiring $O(|\mathcal{T}|)$ preprocessing.

Under these rules, “dominance” means strict improvement on at least one of (L, W, D) with no regression on others.

Definition II.24 (Achievable region). A point (L, W, D) is *achievable* if there exists an admissible observation strategy realizing those values. Let $\mathcal{R} \subseteq \mathbb{R}_{\geq 0} \times \mathbb{R}_{\geq 0} \times [0, 1]$ denote the achievable region.

Definition II.25 (Pareto optimality). A point (L^*, W^*, D^*) is *Pareto-optimal* if there is no achievable (L, W, D) with $L \leq L^*$, $W \leq W^*$, $D \leq D^*$, and at least one strict inequality.

The main result of Section VII is that nominal-tag observation achieves the unique Pareto-optimal point with $D = 0$ in the information-barrier regime.

Definition II.26 (Information-barrier domain). A classification domain has an *information barrier* (relative to Φ) if there exist distinct classes $c_1 \neq c_2$ with identical Φ -profiles. Equivalently, π is not injective on classes.

I. Converse: Tag Rate Lower Bound

Theorem II.27 (Converse). In any information-barrier domain, any scheme achieving $D = 0$ (zero-error identification)

TABLE I
IDENTIFICATION STRATEGIES FOR 1000 CLASSES WITH 50 METHODS.

Strategy	Tag L	Witness W
Nominal (class ID)	$\lceil \log_2 1000 \rceil = 10$ bits	$O(1)$
Duck typing (query all)	0	≤ 50 queries
Adaptive duck typing	0	$\geq d$ queries

requires tag length $L \geq \log_2 k$, where k is the number of distinct classes.

Proof. Zero error requires that distinct classes map to distinct decoder outputs. With k classes, at least $\log_2 k$ bits are needed to encode k distinct values. The converse is tight: $L = \lceil \log_2 k \rceil$ suffices. ■

J. Rate-Distortion Tradeoff

When $D > 0$ is permitted, we obtain a classic rate-distortion tradeoff:

Proposition II.28 (Lossy identification). For any $\epsilon \in (0, 1)$, there exist schemes with $L = 0$ (no tags) achieving $D \leq \epsilon$ with query cost $W = O(\log(1/\epsilon) \cdot d)$, where d is the distinguishing dimension. The tradeoff: smaller D requires larger W .

The zero-error corner ($D = 0$) is special: nominal tagging is the unique Pareto optimum in the information-barrier regime. Relaxing to $D > 0$ enables tag-free schemes at the cost of increased query complexity. This mirrors the classical distinction between zero-error and ϵ -error capacity in channel coding.

K. Concrete Example

Consider a type system with $k = 1000$ classes, each characterized by a subset of $n = 50$ interface methods. Table I compares the strategies.

Here d is the distinguishing dimension, the size of any minimal distinguishing query set. For typical hierarchies, $d \approx 5\text{--}15$. The gap between 10 bits of storage vs. 5–50 queries per identification is the cost of forgoing nominal tagging.

III. COMPLEXITY BOUNDS

IV. CORE THEOREMS

A. The Error Localization Theorem

Definition IV.1 (Error location count). Let $E(\mathcal{O})$ be the number of locations that must be inspected to find all potential violations of a constraint under observation family \mathcal{O} .

Theorem IV.2 (Nominal-tag localization). $E(\text{nominal-tag}) = O(1)$.

Proof. Under nominal-tag observation, the constraint “ v must be of class A ” is satisfied iff $\tau(v) \in \text{subtypes}(A)$. This is determined at a single location: the definition of $\tau(v)$ ’s class. One location. ■

Theorem IV.3 (Declared interface localization). $E(\text{interface-only, declared}) = O(k)$ where $k = \text{number of entity classes}$.

Proof. With declared interfaces, the constraint “ v must satisfy interface I ” requires verifying that each class implements all attributes in I . For k classes, $O(k)$ locations. ■

Theorem IV.4 (Attribute-only localization). $E(\text{attribute-only}) = \Omega(m)$ where $m = \text{number of query sites}$.

Proof. Under attribute-only observation, each query site independently checks “does v have attribute a ?” with no centralized declaration. For m query sites, each must be inspected. Lower bound is $\Omega(m)$. ■

Corollary IV.5 (Strict dominance). *Nominal-tag observation strictly dominates attribute-only: $E(\text{nominal-tag}) = O(1) < \Omega(m) = E(\text{attribute-only})$ for all $m > 1$.*

B. The Information Scattering Theorem

Definition IV.6 (Constraint encoding locations). Let $I(\mathcal{O}, c)$ be the set of locations where constraint c is encoded under observation family \mathcal{O} .

Theorem IV.7 (Attribute-only scattering). *For attribute-only observation, $|I(\text{attribute-only}, c)| = O(m)$ where $m = \text{query sites using constraint } c$.*

Proof. Each attribute query independently encodes the constraint. No shared reference exists. Constraint encodings scale with query sites. ■

Theorem IV.8 (Nominal-tag centralization). *For nominal-tag observation, $|I(\text{nominal-tag}, c)| = O(1)$.*

Proof. The constraint “must be of class A ” is encoded once in the definition of A . All tag checks reference this single definition. ■

Corollary IV.9 (Maintenance entropy). *Attribute-only observation maximizes maintenance entropy; nominal-tag observation minimizes it.*

V. MATROID STRUCTURE

A. Model Contract (Fixed-Axis Domains)

Model contract (fixed-axis domain). A domain is specified by a fixed observation interface Φ derived from a fixed axis map $\alpha : \mathcal{V} \rightarrow \mathcal{A}$ (e.g., $\alpha(v) = (B(v), S(v))$). An observer is permitted to interact with v only through primitive queries in Φ , and each primitive query factors through α : for every $q \in \Phi$, there exists \tilde{q} such that $q(v) = \tilde{q}(\alpha(v))$. A property is in-scope semantic iff it is computable by an admissible strategy that uses only responses to queries in Φ (under our admissibility constraints: no global preprocessing tables, no amortized caching, etc.).

We adopt Φ as the complete observation universe for this paper: to claim applicability to a concrete runtime one must either (i) exhibit mappings from each runtime observable into Φ , or (ii) enforce the admissibility constraints (no external registries, no reflection, no preprocessing/amortization). Under either condition the theorems apply without qualification.

Proposition V.1 (Observational Quotient). *For any admissible strategy using only Φ , the entire interaction transcript (and hence the output) depends only on $\alpha(v)$. Equivalently, any in-scope semantic property P factors through α : there exists \tilde{P} with $P(v) = \tilde{P}(\alpha(v))$ for all v .*

Corollary V.2 (Why “ad hoc” = adding an axis/tag). *If two values v, w satisfy $\alpha(v) = \alpha(w)$, then no admissible Φ -only strategy can distinguish them with zero error. Any mechanism that does distinguish such pairs must introduce additional information not present in α (equivalently, refine the axis map by adding a new axis/tag).*

B. Query Families and Distinguishing Sets

The classification problem is: given a set of queries, which subsets suffice to distinguish all entities?

Definition V.3 (Query family). Let \mathcal{Q} be the set of all primitive queries available to an observer. For a classification system with interface set \mathcal{I} , we have $\mathcal{Q} = \{q_I : I \in \mathcal{I}\}$ where $q_I(v) = 1$ iff v satisfies interface I .

In this section, “queries” are the primitive interface predicates $q \in \Phi$ (equivalently, each q factors through the axis map: $q = \tilde{q} \circ \alpha$). See the Convention above where $\Phi := \mathcal{Q}$.

Convention: $\Phi := \mathcal{Q}$. All universal quantification over “queries” ranges over $q \in \Phi$ only.

Definition V.4 (Distinguishing set). A subset $S \subseteq \mathcal{Q}$ is *distinguishing* if, for all values v, w with $\text{type}(v) \neq \text{type}(w)$, there exists $q \in S$ such that $q(v) \neq q(w)$.

Definition V.5 (Minimal distinguishing set). A distinguishing set S is *minimal* if no proper subset of S is distinguishing.

C. Matroid Structure of Query Families

extbfScope and assumptions. The matroid theorem below is unconditional within the fixed-axis observational theory defined above. In this section, “query” always means a primitive predicate $q \in \Phi$ (equivalently, q factors through α as in the Model Contract). It depends only on:

- $E = \Phi$ is the ground set of primitive queries (interface predicates).
- “Distinguishing”: for all values v, w with $\text{type}(v) \neq \text{type}(w)$, there exists $q \in S$ such that $q(v) \neq q(w)$ (Def. above).
- “Minimal” means inclusion-minimal: no proper subset suffices.

No further assumptions are required within this theory (i.e., beyond the fixed interface Φ already specified). The proof constructs a closure operator satisfying extensivity, monotonicity, and idempotence, from which basis exchange follows (see Lean formalization).

Definition V.6 (Bases family). Let $E = \Phi (= \mathcal{Q})$ be the ground set of primitive queries (interface predicates). Let $\mathcal{B} \subseteq 2^E$ be the family of minimal distinguishing sets.

Lemma V.7 (Basis exchange). *For any $B_1, B_2 \in \mathcal{B}$ and any $q \in B_1 \setminus B_2$, there exists $q' \in B_2 \setminus B_1$ such that $(B_1 \setminus \{q\}) \cup \{q'\} \in \mathcal{B}$.*

Proof sketch. Define the closure operator $\text{cl}(X) = \{q : X\text{-equivalence implies } q\text{-equivalence}\}$. We verify the matroid axioms:

- 1) **Closure axioms:** cl is extensive, monotone, and idempotent. These follow directly from the definition of logical implication.
- 2) **Exchange property:** If $q \in \text{cl}(X \cup \{q'\}) \setminus \text{cl}(X)$, then $q' \in \text{cl}(X \cup \{q\})$.

The exchange property is the non-trivial step. It follows from the symmetry of indistinguishability. If adding q' to X allows distinguishing v, w that were previously X -equivalent (thus determining q), then v, w must differ on q' . This same pair v, w witnesses that adding q to X allows distinguishing q' . Thus the dependency is symmetric.

Minimal distinguishing sets are exactly the bases of the matroid defined by this closure operator. Full machine-checked proof: `proofs/abstract_class_system.lean`, namespace `AxisClosure`. ■

Theorem V.8 (Matroid bases). \mathcal{B} is the set of bases of a matroid on ground set E .

Proof. By the basis-exchange lemma and the standard characterization of matroid bases [13]. ■

Definition V.9 (Distinguishing dimension). The *distinguishing dimension* of a classification system is the common cardinality of all minimal distinguishing sets.

Remark V.10 (Ambient attribute count vs. distinguishing dimension). Let $n := |\mathcal{I}|$ be the ambient number of available attributes (interfaces). Clearly $d \leq n$, and there exist worst-case families with $d = n$.

Corollary V.11 (Well-defined distinguishing dimension). All minimal distinguishing sets have equal cardinality. Thus the distinguishing dimension (Definition V.9) is well-defined.

D. Implications for Witness Cost

Corollary V.12 (Lower bound on interface-only witness cost). For any interface-only observer, $W(\text{type-identity}) \geq d$ where d is the distinguishing dimension.

Proof. Any witness procedure must query at least one minimal distinguishing set. ■

The key insight: the distinguishing dimension is invariant across all minimal query strategies. The difference between nominal-tag and interface-only observers lies in *witness cost*: a nominal tag achieves $W = O(1)$ by storing the identity directly, bypassing query enumeration.

VI. WITNESS COST ANALYSIS

A. Witness Cost for Type Identity

Recall from Section 2 that the witness cost $W(P)$ is the minimum number of primitive queries required to compute property P . For type identity, we ask: what is the minimum number of queries to determine if two values have the same type?

Theorem VI.1 (Nominal-Tag Observers Achieve Minimum Witness Cost). *Nominal-tag observers achieve the minimum witness cost for type identity:*

$$W_{eq} = O(1)$$

Specifically, the witness is a single tag read: compare $\text{tag}(v_1) = \text{tag}(v_2)$.

Interface-only observers require $W_{eq} = \Omega(d)$ where d is the distinguishing dimension (and $d \leq n$, with worst-case $d = n$).

Proof. See `Lean` formalization: `proofs/nominal_resolution.lean`. The proof shows:

- 1) Nominal-tag access is a single primitive query
 - 2) Interface-only observers must query at least d interfaces in the worst case (a generic strategy queries all n)
 - 3) No shorter witness exists for interface-only observers (by the information barrier)
-

B. Witness Cost Comparison

Observer Class	Witness Procedure	Witness Cost W
Nominal-tag	Single tag read	$O(1)$
Interface-only	Query a distinguishing set	$\Omega(d)$

TABLE II
WITNESS COST FOR TYPE IDENTITY BY OBSERVER CLASS.

The Lean 4 formalization (Appendix A) provides a machine-checked proof that nominal-tag access minimizes witness cost for type identity.

VII. (L, W, D) OPTIMALITY

A. Three-Dimensional Tradeoff: Tag Length, Witness Cost, Distortion

Recall from Section 2 that observer strategies are characterized by three dimensions:

- **Tag length** L : bits required to encode a type identifier ($L \geq \log_2 k$ for k types)
- **Witness cost** W : minimum number of primitive queries for type identity checking
- **Distortion** D : probability of misclassification, $D = \Pr[\hat{C} \neq C]$.

We compare two observer classes:

Definition VII.1 (Interface-only observer). An observer that queries only interface membership ($q_I \in \Phi_I$), with no access to explicit type tags.

Definition VII.2 (Nominal-tag observer). An observer that may read a single type identifier (nominal tag) per value, in addition to interface queries.

Theorem VII.3 (Pareto Optimality of Nominal-Tag Observers). *In any information-barrier domain, nominal-tag observers achieve the unique Pareto-optimal point in the (L, W, D) space with $D = 0$:*

- **Tag length:** $L = \lceil \log_2 k \rceil$ bits for k types
- **Witness cost:** $W = O(1)$ queries (one tag read)




Fig. 1. Schematic illustration of the (L, W, D) tradeoff. For a concrete example with $k = 1000$ classes and distinguishing dimension $d = 10$, the nominal-tag strategy achieves $L = 10$ bits, $W = O(1)$, $D = 0$, while the interface-only strategy requires $W = 10$ queries and incurs $D > 0$ due to attribute collisions.

- **Distortion:** $D = 0$ (zero misclassification probability)

In such domains, interface-only observers achieve:

- **Tag length:** $L = 0$ bits (no explicit tag)
- **Witness cost:** $W = \Omega(d)$ queries (must query at least one minimal distinguishing set of size d , see Definition V.9)
- **Distortion:** $D > 0$ (probability of misclassification is strictly positive due to collisions)

Proof. See Lean formalization: `proofs/python_instantiation.lean`. The proof verifies:

- 1) `nominal_cost_constant`: Nominal-tag achieves $(L, W, D) = (O(1), O(1), 0)$
- 2) `interface_cost_linear`: Interface-only admits the generic upper bound $W \leq n$ (query all n attributes); combined with the lower bound $W = \Omega(d)$, this yields an unbounded separation from nominal tagging
- 3) `python_gap_unbounded`: The cost gap is unbounded in the limit
- 4) Interface observations alone cannot distinguish provenance; nominal tags can

■

B. Pareto Frontier

The three-dimensional frontier shows:

- Nominal-tag observers dominate interface-only observers on all three dimensions
- Interface-only observers trade tag length for distortion (zero L , but $D > 0$)

Figure 1 visualizes the (L, W, D) tradeoff space. The key observation: *nominal tags trade storage for query cost*, achieving the optimal $(L, W, D) = (\log_2 k, O(1), 0)$ point.

The Lean 4 formalization (Appendix A) provides a machine-checked proof of Pareto optimality for nominal-tag observers in the (L, W, D) tradeoff.

Remark VII.4 (Programming language instantiations). In programming language terms: *nominal typing* corresponds to nominal-tag observers (e.g., CPython’s `isinstance`, Java’s `.getClass()`). *Duck typing* corresponds to interface-only

observers (e.g., Python’s `hasattr`). *Structural typing* is an intermediate case with $D = 0$ but $W = O(n)$.

Remark VII.5 (Structural-typing cost parameter). Using the notation above, a structural check that traverses s members/fields has $W = O(s)$ with $s \leq n$.

Remark VII.6 (Structural-check cost parameter). When structural typing checks traverse s members/fields (rather than ranging over the full attribute universe), the natural bound is $W = O(s)$ with $s \leq n$.

Using the notation above, a structural check that traverses s members/fields has $W = O(s)$ with $s \leq n$.

Remark VII.7 (Structural-check cost parameter). When structural typing checks traverse s members/fields (rather than ranging over the full attribute universe), the natural bound is $W = O(s)$ with $s \leq n$.

VIII. INSTANTIATIONS IN REAL RUNTIMES

The preceding sections established abstract results about observer classes and witness cost. We now ground these in concrete systems across multiple domains, showing that real classification systems instantiate the theoretical categories and that the complexity bounds are not artifacts of the model but observable properties of deployed implementations.

A. Biological Taxonomy: Phenotype vs Genotype

Linnean taxonomy classifies organisms by observable phenotypic characters: morphology, behavior, habitat. This is attribute-only observation. The information barrier applies: phenotypically identical organisms from distinct species are indistinguishable.

The cryptic species problem: Cryptic species share identical phenotypic profiles but are reproductively isolated and genetically distinct. Attribute-only observation (morphology) cannot distinguish them: $\pi(A) = \pi(B)$ but $\text{species}(A) \neq \text{species}(B)$.

The nominal tag: DNA barcoding provides the resolution [12]. A short genetic sequence (e.g., mitochondrial COI) acts as the nominal tag: $O(1)$ identity verification via sequence comparison. This reduced cryptic species identification from $\Omega(s)$ phenotypic examination (checking s characters) to constant-time molecular lookup.

B. Library Classification: Subject vs ISBN

Library classification systems like Dewey Decimal observe subject matter, a form of attribute-only classification. Two books on the same subject are indistinguishable by subject code alone.

The nominal tag: The ISBN (International Standard Book Number) is the nominal tag [14]. Given two physical books, identity verification is $O(1)$: compare ISBNs. Without ISBNs, distinguishing two copies of different editions on the same subject requires $O(s)$ attribute inspection (publication date, page count, publisher, etc.).

C. Database Systems: Columns vs Primary Keys

Relational databases observe entities via column values. The information barrier applies: rows with identical column values, excluding the key, are indistinguishable.

The nominal tag: The primary key is the nominal tag [15]. Entity identity is $O(1)$: compare keys. This is why database theory requires keys—without them, the system cannot answer “is this the same entity?”

Natural vs surrogate keys: Natural keys (composed of attributes) are attribute-only observation and inherit its limitations. Surrogate keys (auto-increment IDs, UUIDs) are pure nominal tags: no semantic content, pure identity.

D. Programming Language Runtimes

Type systems are the motivating example for this work. We survey four runtimes.

1) **CPython: The `ob_type` Pointer:** Every CPython heap object begins with a `PyObject` header containing an `ob_type` pointer to its type object [16]. This is the nominal tag: a single machine word encoding complete type identity.

Witness procedure: Given objects `a` and `b`, type identity is `type(a) is type(b)`—two pointer dereferences and one pointer comparison. Cost: $O(1)$ primitive operations, independent of interface count.

Contrast with `hasattr`: Interface-only observation in Python uses `hasattr(obj, name)` for each required method. To verify an object satisfies a protocol with s required methods requires s attribute lookups. Worse: different call sites may check different subsets, creating $\Omega(m)$ total checks where m is the number of call sites. The nominal tag eliminates this entirely.

2) **Java: `.getClass()` and the Method Table:** Java’s object model stores a pointer to the class object in every instance header [17]. The `.getClass()` method exposes this [18], and `instanceof` checks traverse the class hierarchy.

Key observation: Java’s `instanceof` is $O(d)$ where d is inheritance depth, not $O(|\mathcal{I}|)$ where $|\mathcal{I}|$ is the number of interfaces. This is because `instanceof` walks the inheritance hierarchy (a nominal-tag query), not the interface list (an attribute query).

3) **TypeScript: Structural Equivalence:** TypeScript uses attribute-only (declared) observation [19]: the compiler checks structural compatibility, not nominal identity. Two types are assignment-compatible iff their structures match.

Implication: Type identity checking requires traversing the structure. For a type with s fields/methods, $W(\text{type-identity}) = O(s)$. This is inherent to the observation model: no compilation strategy can reduce this to $O(1)$ without adding nominal tags.

4) **Rust: Static Nominal Tags:** Rust resolves type identity at compile time via its nominal type system. At runtime, `std::any::TypeId` provides nominal-tag access [20].

The dyn Trait case: Rust’s trait objects include a vtable pointer but not a type tag [21]. This is attribute-only observation: the vtable encodes which methods exist, not which type provided them.

Domain	Attribute-Only	Nominal Tag	W
Biology	Phenotype (morphology)	DNA barcode (COI)	$O(1)$
Libraries	Subject (Dewey)	ISBN	$O(1)$
Databases	Column values	Primary key	$O(1)$
CPython	<code>hasattr</code> probing	<code>ob_type</code> pointer	$O(1)$
Java	Interface check	<code>.getClass()</code>	$O(1)$
TypeScript	Structural check	(none at runtime)	$O(s)$
Rust (static)	Trait bounds	<code>TypeId</code>	$O(1)$

TABLE III
WITNESS COST FOR IDENTITY ACROSS CLASSIFICATION SYSTEMS.
NOMINAL TAGS ACHIEVE $O(1)$; ATTRIBUTE-ONLY PAYS $O(s)$ PER
STRUCTURAL CHECK (OR $O(k)$ WHEN ENUMERATING
CLASSES/DECLARED INTERFACES).

E. Cross-Domain Summary

The pattern is universal: systems with nominal tags achieve $O(1)$ witness cost; systems without them pay $O(s)$ or $O(k)$. This is not domain-specific; it is the information barrier theorem instantiated across classification systems.

IX. EXTENSIONS

A. Noisy Query Model

Throughout this paper, queries are deterministic: $q_I(v) \in \{0, 1\}$ is a fixed function of v . In practice, observations may be corrupted. We sketch an extension to noisy queries and state the resulting open problems.

Definition IX.1 (Noisy observation channel). A *noisy observation channel* with crossover probability $\epsilon \in [0, 1/2]$ returns:

$$\tilde{q}_I(v) = \begin{cases} q_I(v) & \text{with probability } 1 - \epsilon \\ 1 - q_I(v) & \text{with probability } \epsilon \end{cases}$$

Each query response is an independent $\text{BSC}(\epsilon)$ corruption of the true value.

Definition IX.2 (Noisy identification capacity). The ϵ -noisy identification capacity is the supremum rate (in bits per entity) at which zero-error identification is achievable when all attribute queries pass through a $\text{BSC}(\epsilon)$.

In the noiseless case ($\epsilon = 0$), Theorem II.10 shows the capacity is binary: $\log_2 k$ if π is class-injective, 0 otherwise. For $\epsilon > 0$, several questions arise.

Open problem (noisy identification cost). For $\epsilon > 0$ and class-injective π , zero-error identification is impossible with finite queries (since BSC has nonzero error probability). With bounded error $\delta > 0$, we expect the identification cost to scale as $W = \Theta\left(\frac{\log(1/\delta)}{(1-2\epsilon)^2}\right)$ queries per entity. A key observation is that a nominal tag of $L \geq \lceil \log_2 k \rceil$ bits (transmitted noiselessly) should restore $O(1)$ identification regardless of query noise.

The third point is the key insight: *nominal tags provide a noise-free side channel*. Even when attribute observations are corrupted, a clean tag enables $O(1)$ identification. This strengthens the case for nominal tagging in noisy environments, precisely the regime where “duck typing” would require many repeated queries to achieve confidence.

Connection to identification via channels. The noisy model connects more directly to Ahlsweide-Dueck identification [1]. In their framework, identification capacity over

a noisy channel can exceed Shannon capacity (double-exponential codebook sizes). Our setting differs: we have *adaptive queries* rather than block codes, and the decoder must identify a *class* rather than test a hypothesis. Characterizing the interplay between adaptive query strategies and channel noise is an open problem.

B. Rate-Distortion-Query Tradeoff Surface

The (L, W, D) tradeoff admits a natural geometric interpretation. We have identified the unique Pareto-optimal point at $D = 0$ (Theorem VII.3), but the full tradeoff surface contains additional structure.

Fixed- W slices. For fixed query budget W , what is the minimum tag rate L to achieve distortion D ? When $W \geq d$ (the distinguishing dimension), zero distortion is achievable with $L = 0$ via exhaustive querying. When $W < d$, the observer cannot distinguish all classes, and either:

- Accept $D > 0$ (misidentification), or
- Add tags ($L > 0$) to compensate for insufficient queries.

Fixed- L slices. For fixed tag rate $L < \log_2 k$, the tag partitions the k classes into 2^L groups. Within each group, the observer must use queries to distinguish. The query cost is determined by the distinguishing dimension *within each group*, which is potentially much smaller than the global dimension.

Open problem (subadditivity of query cost). For a tag of rate L partitioning classes into groups G_1, \dots, G_{2^L} , we expect $W(L) \leq \max_i d(G_i)$, where $d(G_i)$ is the distinguishing dimension within group G_i . Optimal tag design should minimize this maximum. Characterizing the optimal partition remains open.

C. Semantic Distortion Measures

We have treated distortion D as binary (correct identification or not). Richer distortion measures are possible:

- **Hierarchical distortion:** Misidentifying a class within the same genus (biological) or module (type system) is less severe than cross-genus errors.
- **Weighted distortion:** Some misidentifications have higher cost than others (e.g., type errors causing security vulnerabilities vs. benign type confusion).

D. Connection to Rate-Distortion-Perception Theory

Blau and Michaeli [5] extended classical rate-distortion theory by adding a *perception* constraint: the reconstructed distribution must match a target distribution under some divergence measure. This creates a three-way tradeoff between rate, distortion, and perceptual quality.

Our (L, W, D) framework admits a parallel interpretation. The query cost W plays a role analogous to the perception constraint: it measures the *interactive cost* of achieving low distortion, rather than a distributional constraint. Just as rate-distortion-perception theory asks “what is the minimum rate to achieve distortion D while satisfying perception constraint P ”, we ask “what is the minimum tag rate L to achieve distortion D with query budget W ?”

The analogy suggests several directions:

- **Perception as identification fidelity:** In classification systems that must preserve statistical properties (e.g., sampling from a type distribution), a perception constraint would require the observer’s class estimates to have the correct marginal distribution, not just low expected error.
- **Three-resource tradeoffs:** The (L, W, D) Pareto frontier (Theorem VII.3) is a discrete analogue of the rate-distortion-perception tradeoff surface. Characterizing this surface for specific classification problems would extend the geometric rate-distortion program to identification settings.

Formalizing these connections would unify identification capacity with the broader rate-distortion-perception literature.

X. CONCLUSION

This paper presents an information-theoretic analysis of classification under observational constraints. We prove three main results:

- 1) **Information Barrier:** Observers limited to attribute-membership queries cannot compute properties that vary within indistinguishability classes. This is universal: it applies to biological taxonomy, database systems, library classification, and programming language runtimes alike.
- 2) **Witness Optimality:** Nominal-tag observers achieve $W(\text{identity}) = O(1)$, the minimum witness cost. The gap from attribute-only observation ($\Omega(d)$, with a worst-case family where $d = n$) is unbounded.
- 3) **Matroid Structure:** Minimal distinguishing query sets form the bases of a matroid. The distinguishing dimension of a classification problem is well-defined and computable.

A. The Universal Pattern

Across domains, the same structure recurs:

- **Biology:** Phenotypic observation cannot distinguish cryptic species. DNA barcoding (nominal tag) resolves them in $O(1)$.
- **Databases:** Column-value queries cannot distinguish rows with identical attributes. Primary keys (nominal tag) provide $O(1)$ identity.
- **Type systems:** Interface observation cannot distinguish structurally identical types. Type tags provide $O(1)$ identity.

The information barrier is not a quirk of any particular domain; it is a mathematical necessity arising from the quotient structure induced by limited observations.

B. Implications

- **The necessity of nominal tags is a theorem, not a preference.** In information-barrier domains (Definition II.26), any scheme achieving zero-error identification ($D = 0$) requires tag length $L \geq \log_2 k$ (Theorem II.27). Nominal-tag observation achieves the unique Pareto-optimal $D = 0$ point with $W = O(1)$ at this bound (Theorem VII.3).

- **The barrier is informational, not computational:** even with unbounded resources, attribute-only observers cannot overcome it.
- **Classification system design is constrained:** the choice of observation family determines which properties are computable.

C. Future Work

- 1) **Other classification domains:** What is the matroid structure of observation spaces in chemistry (molecular fingerprints), linguistics (phonetic features), or machine learning (feature embeddings)?
- 2) **Witness complexity of other properties:** Beyond identity, what are the witness costs for provenance, equivalence, or subsumption?
- 3) **Hybrid observers:** Can observer strategies that combine tags and attributes achieve better (L, W, D) tradeoffs for specific query distributions?

D. Conclusion

Classification under observational constraints admits a clean information-theoretic analysis. In information-barrier domains (Definition II.26), nominal-tag observation achieves the unique Pareto-optimal $D = 0$ point in the (L, W, D) tradeoff (Theorem VII.3), and any $D = 0$ scheme necessarily has $L \geq \log_2 k$ (Theorem II.27). The results are universal within the stated observation model, and all proofs are machine-verified in Lean 4.

AI Disclosure

This work was developed with AI assistance (Claude, Anthropic). The AI contributed to exposition, code generation, and proof exploration. All mathematical claims were verified by the authors and machine-checked in Lean 4. The Lean proofs are the authoritative source; no theorem depends solely on AI-generated reasoning.

REFERENCES

- [1] R. Ahlswede and G. Dueck, “Identification via channels,” *IEEE Transactions on Information Theory*, vol. 35, no. 1, pp. 15–29, 1989.
- [2] ——, “Identification in the presence of feedback—a discovery of new capacity formulas,” *IEEE Transactions on Information Theory*, vol. 35, no. 1, pp. 30–36, 1989.
- [3] C. E. Shannon, “Coding theorems for a discrete source with a fidelity criterion,” *IRE National Convention Record*, vol. 7, pp. 142–163, 1959.
- [4] T. Berger, *Rate Distortion Theory: A Mathematical Basis for Data Compression*. Prentice-Hall, 1971.
- [5] Y. Blau and T. Michaeli, “Rethinking lossy compression: The rate-distortion-perception tradeoff,” *Proceedings of the 36th International Conference on Machine Learning*, pp. 675–685, 2019.
- [6] J. Körner, “Coding of an information source having ambiguous alphabet and the entropy of graphs,” *6th Prague Conference on Information Theory*, pp. 411–425, 1973.
- [7] H. S. Witsenhausen, “The zero-error side information problem and chromatic numbers,” *IEEE Transactions on Information Theory*, vol. 22, no. 5, pp. 592–593, 1976.
- [8] H. Buhrman, L. Fortnow, M. Koucký, and B. Loff, “Computational complexity of discrete problems,” *Bulletin of the EATCS*, vol. 77, pp. 42–56, 2002.
- [9] A. Orlitsky, “Worst-case interactive communication i: Two messages are almost optimal,” *IEEE Transactions on Information Theory*, vol. 37, no. 4, pp. 994–1009, 1991.

- [10] L. Cardelli and P. Wegner, “On understanding types, data abstraction, and polymorphism,” *ACM Computing Surveys*, vol. 17, no. 4, pp. 471–523, 1985.
- [11] W. R. Cook, W. Hill, and P. S. Canning, “Inheritance is not subtyping,” in *Proceedings of the 17th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, 1990, pp. 125–135.
- [12] P. D. N. Hebert, A. Cywinski, S. L. Ball, and J. R. deWaard, “Biological identifications through dna barcodes,” pp. 313–321, 2003.
- [13] D. J. A. Welsh, *Matroid Theory*. Academic Press, 1976.
- [14] International ISBN Agency, *ISBN Users’ Manual*, 7th ed. International ISBN Agency, 2017.
- [15] E. F. Codd, *The Relational Model for Database Management: Version 2*. Addison-Wesley, 1990.
- [16] P. S. Foundation, “Cpython object implementation: Object structure,” <https://github.com/python/cpython/blob/main/Include/object.h>, 2024.
- [17] Oracle, “The java virtual machine specification: Chapter 2.5 runtime data areas,” <https://docs.oracle.com/javase/specs/jvms/se21/html/jvms-2.html>, 2024.
- [18] ——, “Java object: getClass()”, <https://docs.oracle.com/javase/8/docs/api/java/lang/Object.html#getClass-->, 2024.
- [19] Microsoft, “TypeScript handbook: Type compatibility,” <https://www.typescriptlang.org/docs/handbook/type-compatibility.html>, 2024.
- [20] R. Foundation, “Rust std::any::type_id,” https://doc.rust-lang.org/std/any/fn.type_id.html, 2024.
- [21] ——, “The rust reference: Trait objects,” <https://doc.rust-lang.org/reference/types/trait-object.html>, 2024.

APPENDIX

We provide machine-checked proofs of our core theorems in Lean 4. The complete development (6,100+ lines across nine modules, 0 `sorry` placeholders) is organized as follows:

1. **Language-agnostic layer** (Section 6.12): The two-axis model (B, S) , axis lattice metatheorem, and strict dominance: proving nominal-tag observation dominates interface-only observation in **any** class system with explicit inheritance. These proofs require no Python-specific axioms.
2. **Python instantiation layer** (Sections 6.1–6.11): The dual-axis resolution algorithm, provenance preservation, and OpenHCS-specific invariants: proving that Python’s `type(name, bases, namespace)` and C3 linearization correctly instantiate the abstract model.
3. **Complexity bounds layer** (Section 6.13): Formalization of $O(1)$ vs $O(k)$ vs $\Omega(m)$ complexity separation (where m is the number of query sites). Proves that nominal error localization is $O(1)$, interface-only (declared) is $O(k)$, interface-only is $\Omega(m)$, and the gap grows without bound.

The abstract layer establishes that our theorems apply to Java, C#, Ruby, Scala, and any language with the (B, S) structure. The Python layer demonstrates concrete realization. The complexity layer proves the asymptotic dominance is machine-checkable, not informal.

A. Type Universe, Normalization Registry, and Lineage

We represent nominal types as atoms (`Typ := Nat`). The model uses two distinct structures: (i) a lineage order `mro : MRO` (most-specific-first) capturing multi-step inheritance (e.g., the result of C3 linearization), and (ii) an OpenHCS-specific normalization registry `R : Registry` that optionally rewrites certain nominal types to a canonical base type used for resolution and provenance reporting. Importantly, `R` is not the inheritance relation; inheritance is carried by `mro`. The

TABLE IV
LEAN 4 FORMALIZATION MODULES

Module	Lines	Theorems	Purpose
abstract_class_system.lean	3082	90+	Core: two-axis model, dominance
axis_framework.lean	1667	40+	Matroid structure
nominal_resolution.lean	556	21	Resolution, capabilities
discipline_migration.lean	142	11	Discipline vs migration
context_formalization.lean	215	7	Greenfield/retrofit
python_instantiation.lean	247	12	Python instantiation
typescript_instantiation.lean	65	3	TypeScript instantiation
java_instantiation.lean	63	3	Java instantiation
rust_instantiation.lean	64	3	Rust instantiation
Total	6100+	190+	

registry invariant `R.wellFormed` states that normalization targets are fixed points (canonical bases do not normalize further), yielding normalization idempotence.

```
-- Types as natural numbers (nominal identity)
abbrev Typ := Nat

-- Normalization registry: optional rewrite to a
-- canonical base type
def Registry := Typ -> Option Typ

-- Well-formed normalization: rewrite targets are
-- fixed points
def Registry.wellFormed (R : Registry) : Prop :=
  forall L B, R L = some B -> R B = none

-- One-step normalization used during resolution /
-- provenance
def normalizeType (R : Registry) (T : Typ) : Typ :=
  match R T with | some B => B | none => T
```

Invariant (Normalization Idempotence). For well-formed registries:

```
theorem normalizeType_idempotent (R : Registry) (T
  : Typ)
  (h_wf : R.wellFormed) :
  normalizeType R (normalizeType R T) =
  normalizeType R T := by
simp only [normalizeType]; cases hR : R T with
| none => simp only [hR]
| some B => have h_base := h_wf T B hR; simp
only [h_base]
```

B. MRO and Scope Stack

```
abbrev MRO := List Typ           -- Most specific
first
abbrev ScopeStack := List ScopeId

structure ConfigInstance where
  typ : Typ
  fieldValue : FieldValue

def ConfigContext := ScopeId -> List ConfigInstance
```

Here, `mro : MRO` denotes the multi-step inheritance linearization (most-specific-first); our mechanized results assume it is a valid linearization (e.g., produced by C3 when it succeeds).

C. The RESOLVE Algorithm

```
structure ResolveResult where
  value : FieldValue; scope : ScopeId; sourceType
  : Typ
deriving DecidableEq

def findConfigByType (configs : List
  ConfigInstance) (T : Typ) :=
  match configs.find? (fun c => c.typ == T) with
  | some c => some c.fieldValue | none => none

def resolve (R : Registry) (mro : MRO)
  (scopes : ScopeStack) (ctx : ConfigContext) :
  Option ResolveResult :=
  scopes.findSome? fun scope =>          -- X-axis:
    scopes
    mro.findSome? fun mroType =>          -- Y-axis: MRO
      let normType := normalizeType R mroType
      match findConfigByType (ctx scope) normType
      with
      | some v => if v != 0 then some (v, scope,
        normType) else none
      | none => none
```

D. GETATTRIBUTE Implementation

```
def rawFieldValue (obj : ConfigInstance) :=
  obj.fieldValue

def getattribute (R : Registry) (obj :
  ConfigInstance) (mro : MRO)
  (scopes : ScopeStack) (ctx : ConfigContext)
  (isLazy : Bool) :=
  let raw := rawFieldValue obj
  if raw != 0 then raw
  else if isLazy then
    match resolve R mro scopes ctx with
    | some result => result.value | none => 0
  else raw
```

E. Theorem 6.1: Resolution Completeness

Theorem 6.1 (Completeness). The `resolve` function is complete: it returns value v if and only if either no resolution occurred ($v = 0$) or a valid resolution result exists.

```
theorem resolution_completeness (R : Registry)
  (mro : MRO)
  (scopes : ScopeStack) (ctx : ConfigContext) (v
  : FieldValue) :
  (match resolve R mro scopes ctx with
  | some r => r.value | none => 0) = v <->
  (v = 0 \& resolve R mro scopes ctx = none) \/
  (exists r, resolve R mro scopes ctx = some r
  \& r.value = v) := by
cases hr : resolve R mro scopes ctx <;> simp [hr]
```

F. Theorem 6.2: Provenance Preservation

Theorem 6.2a (Uniqueness). Resolution is deterministic.

```
theorem provenance_uniqueness (R : Registry) (mro : MRO)
  (scopes : ScopeStack) (ctx : ConfigContext)
  (r1 r2 : ResolveResult)
  (hr1 : resolve R mro scopes ctx = some r1)
  (hr2 : resolve R mro scopes ctx = some r2) :
  r1 = r2 := by simp only [hr1,
  Option.some.injEq] at hr2; exact hr2

theorem resolution_determinism (R : Registry) (mro : MRO)
  (scopes : ScopeStack) (ctx : ConfigContext) :
  forall r1 r2, resolve R mro scopes ctx = r1 ->
    resolve R mro scopes ctx = r2 ->
    r1 = r2 := by
  intros r1 r2 h1 h2; rw [← h1, ← h2]
```

G. Duck Typing Formalization

We formalize interface-only observation and prove it cannot provide provenance.

```
structure InterfaceValue where
  fields : List (String * Nat)
deriving DecidableEq

def getField (obj : InterfaceValue) (name : String) : Option Nat :=
  match obj.fields.find? (fun p => p.1 == name) with
  | some p => some p.2 | none => none

-- Interface equivalence: identity = structure
def interfaceEquivalent (a b : InterfaceValue) : Prop :=
  forall name, getField a name = getField b name
```

We prove this is an equivalence relation:

```
theorem structEq_refl (a : InterfaceValue) :
  interfaceEquivalent a a := by intro name; rfl

theorem structEq_symm (a b : InterfaceValue) :
  interfaceEquivalent a b -> interfaceEquivalent b a := by
  intro h name; exact (h name).symm

theorem structEq_trans (a b c : InterfaceValue) :
  interfaceEquivalent a b -> interfaceEquivalent b c ->
  interfaceEquivalent a c := by
  intro hab hbc name; rw [hab.name, hbc.name]

def InterfaceRespecting (f : InterfaceValue -> a) : Prop :=
  forall a b, interfaceEquivalent a b -> f a = f b
```

H. Corollary 6.3: Provenance Impossibility

Under interface-only observation, equivalent objects are indistinguishable, so provenance must be constant on equivalent objects.

```
structure DuckProvenance where
  value : Nat; source : InterfaceValue
deriving DecidableEq

theorem interface_provenance_indistinguishable (getProvenance : InterfaceValue -> Option DuckProvenance)
```

```
(h_interface : InterfaceRespecting)
  getProvenance)
  (obj1 obj2 : InterfaceValue)
  (h_equiv : interfaceEquivalent obj1 obj2) :
  getProvenance obj1 = getProvenance obj2 :=
  h_interface obj1 obj2 h_equiv

theorem interface_provenance_absurdity (getProvenance : InterfaceValue -> Option DuckProvenance)
  (h_interface : InterfaceRespecting)
  getProvenance)
  (obj1 obj2 : InterfaceValue) (h_equiv : interfaceEquivalent obj1 obj2)
  (prov1 prov2 : DuckProvenance)
  (h1 : getProvenance obj1 = some prov1)
  (h2 : getProvenance obj2 = some prov2) :
  prov1 = prov2 := by
  have h_eq := h_interface obj1 obj2 h_equiv
  rw [h1, h2] at h_eq; exact Option.some.inj h_eq
```

Contrast with nominal-tag observation: Types are distinguished by identity:

```
def WellFilterConfigType : Nat := 1
def StepWellFilterConfigType : Nat := 2

theorem nominal_types_distinguishable :
  WellFilterConfigType != StepWellFilterConfigType := by decide
```

Therefore, `ResolveResult.sourceType` is meaningful: it tells you WHICH type provided the value, even if types have the same structure.

I. Verification Status

All proofs compile without warnings or `sorry` placeholders. The verification covers:

- **AbstractClassSystem** (475 lines): Two-axis model, strict dominance, axis lattice metatheorem
- **NominalResolution** (157 lines): Resolution algorithm, completeness (Thm. 6.1), uniqueness (Thm. 6.2)
- **DuckTyping** (127 lines): Structural equivalence, impossibility (Cor. 6.3)
- **MetaprogrammingGap** (156 lines): Declaration/query model, hook requirements
- **CapabilityExhaustiveness** (42 lines): Capability completeness theorems
- **AdapterAmortization** (60 lines): Cost model, amortization proofs

Total: 6,100+ lines, 190+ theorems, 0 `sorry`, 0 warnings.

J. What the Lean Proofs Guarantee

The machine-checked verification establishes:

1. **Algorithm correctness:** `resolve` returns value `v` iff resolution found a config providing `v` (Theorem 6.1).
2. **Determinism:** Same inputs always produce same `(value, scope, sourceType)` tuple (Theorem 6.2).
3. **Idempotence:** Normalizing an already-normalized type is a no-op (`normalization_idempotent`).
4. **Interface-only observation impossibility:** Any function respecting interface equivalence cannot distinguish

between interface-identical objects, making provenance tracking impossible (Corollary 6.3).

What the proofs do NOT guarantee:

- **C3 correctness:** We assume MRO is well-formed. Python’s C3 algorithm can fail on pathological diamonds (raising `TypeError`). Our proofs apply only when C3 succeeds.
- **Normalization-registry invariant:** `Registry.wellFormed` is assumed as an OpenHCS invariant: canonical base types are fixed points of normalization (normalization does not chain). We prove all results conditional on this invariant, but do not derive it from more primitive foundations.
- **Termination and complexity:** We use Lean’s termination checker to verify `resolve` terminates. The complexity bound $O(|\text{scopes}| \times |\text{MRO}|)$ is also mechanically verified via `resolution_complexity_bound` and related lemmas proving linearity in each dimension.

This is standard practice in mechanized verification: CompCert assumes well-typed input, seL4 assumes hardware correctness. Our proofs establish that *given* a well-formed registry and MRO, the resolution algorithm is correct and provides provenance that interface-only observation cannot.

K. On Proof Structure

Most proofs in the Lean formalization are short (1-3 lines). This brevity reflects the nature of *definitional* impossibilities: our core theorems establish information-theoretic barriers, not computational lower bounds. When we prove that no shape-respecting function can compute provenance, the proof follows immediately from definitions. A definitional impossibility is stronger than a complexity bound; it closes all loopholes rather than leaving room for alternative algorithms.

The semantic contribution is threefold: (1) *precision forcing*, where formalizing “interface-only observation” requires stating exactly what shape-respecting means; (2) *completeness*, where the query space partition proves every query is either shape-respecting or lineage-dependent; and (3) *universal scope*, where the impossibility applies to any system that discards the lineage axis. The Lean compiler verifies every proof step; reviewers can run `lake build` to confirm the 6000+ lines compile with zero `sorry` placeholders.

L. External Provenance Map Rebuttal

Objection: “Interface-only observation could provide provenance via an external map: `provenance_map : Dict[id(obj), SourceType]`.”

Rebuttal: This objection conflates *object identity* with *type identity*. The external map tracks which specific object instance came from where (not which *type* in the MRO provided a value).

Consider:

```
class A:
    x = 1
class B(A):
    pass # Inherits x from A
b = B()
print(b.x) # Prints 1. Which type provided this?
```

An external provenance map could record `provenance_map[id(b)] = B`. But this doesn’t answer the question “which type in B’s MRO provided x?” The answer is A, and this requires MRO traversal, which requires the Bases axis.

Formal statement: Let `ExternalMap : ObjectId → SourceType` be any external provenance map. Then `ExternalMap` cannot answer: “Which type in $\text{MRO}(\text{type}(v))$ provided attribute a ?”

Proof. The question asks about MRO position. MRO is derived from Bases. `ExternalMap` has no access to Bases (it maps object IDs to types, not types to MRO positions). Therefore `ExternalMap` cannot answer MRO-position queries. ■

The deeper point: Provenance is not about “where did this object come from?” It’s about “where did this *value* come from in the inheritance hierarchy?” The latter requires MRO, which requires Bases, which interface-only observation discards.

M. Abstract Model Lean Formalization

The abstract class system model (Section 2.4) is formalized in Lean 4 with complete proofs (no `sorry` placeholders):

```
-- The two axes of a class system
inductive Axis where
| Bases      -- B: inheritance hierarchy
| Namespace   -- S: attribute declarations
deriving DecidableEq, Repr

abbrev AxisSet := List Axis
def shapeAxes : AxisSet := [.Namespace] -- S-only
def nominalAxes : AxisSet := [.Bases, .Namespace]

inductive UnifiedCapability where
| interfaceCheck | identity | provenance
| enumeration | conflictResolution
deriving DecidableEq, Repr

def axisCapabilities (a : Axis) : List UnifiedCapability :=
match a with
| .Bases => [.identity, .provenance,
              .enumeration, .conflictResolution]
| .Namespace => [.interfaceCheck]

def axisSetCapabilities (axes : AxisSet) :=
axes.flatMap axisCapabilities |>.eraseDups
```

Definition 6.3a (Axis Projection — Lean). A type over axis set A is a dependent tuple of axis values:

```
def AxisProjection (A : List Axis) :=
(a : Axis) -> a in A -> AxisCarrier a

structure Typ where
  ns : Finset AttrName
  bs : List Typ
```

Theorem 6.3b (Isomorphism). `Typ` is isomorphic to the 2-axis projection:

```
noncomputable def Typ.equivProjection :
Typ <~> AxisProjection canonicalAxes where
  toFun := Typ.toProjection
  invFun := Typ.fromProjection
  left_inv := Typ.projection_roundtrip
  right_inv := Typ.projection_roundtrip_inv

theorem n_axis_types_are_projections (A : List Axis) :
GenericTyp A = AxisProjection A := rfl
```

This formalizes Definition 2.10 (Typing Disciplines as Axis Projections): a typing discipline using axis set A has exactly the information contained in the A -projection of the full type.

Theorem 6.4 (Axis Lattice — Lean). Shape capabilities are a strict subset of nominal capabilities:

```
theorem axis_shape_subset_nominal :
  forall c in axisSetCapabilities shapeAxes,
    c in axisSetCapabilities nominalAxes := by
  intro c hc; rw [h_shape] at hc; simp at hc; rw
  [hc]
  exact h_nominal

theorem axis_nominal_exceeds_shape :
  exists c in axisSetCapabilities nominalAxes,
    c notin axisSetCapabilities shapeAxes := by
  use UnifiedCapability.provenance; decide

theorem lattice_dominance :
  (forall c in shapeCapabilities, c in
    nominalCapabilities) /\ 
  (exists c in nominalCapabilities, c notin
    shapeCapabilities) :=
<axis_shape_subset_nominal,
  axis_nominal_exceeds_shape>
```

This formalizes Theorem 2.15: using more axes provides strictly more capabilities. The proofs are complete and compile without any sorry placeholders.

Theorem 6.11 (Capability Completeness — Lean). The Bases axis provides exactly four capabilities, no more:

```
inductive Capability where
| interfaceCheck | typeNaming | valueAccess
| methodInvocation | provenance | identity
| enumeration | conflictResolution
deriving DecidableEq, Repr

def basesRequiredCapabilities : List Capability :=
[.provenance, .identity, .enumeration,
 .conflictResolution]

def nonBasesCapabilities : List Capability :=
[.interfaceCheck, .typeNaming, .valueAccess,
 .methodInvocation]

theorem capability_partition : forall c :
  Capability,
  (c in basesRequiredCapabilities \/ c in
    nonBasesCapabilities) /\ 
  ~(c in basesRequiredCapabilities /\ c in
    nonBasesCapabilities) := by
  intro c; cases c <;> simp
  [basesRequiredCapabilities,
  nonBasesCapabilities]

theorem bases_capabilities_count :
  basesRequiredCapabilities.length = 4 := rfl
```

This formalizes Theorem 2.17 (Capability Completeness): the capability set \mathcal{C}_B is **exactly** four elements, proven by exhaustive enumeration with machine-checked partition. The capability_partition theorem proves that every capability falls into exactly one category (Bases-required or not) with no overlap and no gaps.

Scope as observational quotient. We model “scope” as a set of allowed observers $\text{Obs} \subseteq (W \rightarrow O)$ and define observational equivalence $x \approx y : \iff \forall f \in \text{Obs}, f(x) = f(y)$. The induced quotient W/\approx is the canonical object for that scope, and every in-scope observer factors through it (see `observer_factors` in

`abstract_class_system.lean`). Once the observer set is fixed, no argument can appeal to information outside that quotient; adding a new observable is literally expanding Obs.

Protocol runtime observer (shape-only). We also formalize the restricted Protocol/isinstance observer that checks only for required members. The predicate `protoCheck` ignores protocol identity and is proved `shape-respecting` (`protoCheck_in_shapeQuerySet` in `abstract_class_system.lean`), so two protocols with identical member sets are indistinguishable to that observer. Distinguishing them requires adding an observable discriminator (brand/tag/nominality), i.e., moving to another axis.

All Python object-model observables factor through axes. In the Python instantiation we prove that core runtime discriminators are functions of (B, S) : metaclass selection depends only on bases (`metaclass_depends_on_bases`); attribute presence and dispatch depend only on the namespace (`getattr_depends_on_ns`); together they yield `observer_factors_through_axes` in `python_instantiation.lean`.