

Optimal Encoding Under Coherence Constraints: Rate-Complexity Tradeoffs in Multi-Location Representation Systems

Tristan Simas

Abstract—Consider an information system encoding a fact F at multiple locations. When can such a system guarantee coherence—the impossibility of disagreement among encoding locations? We prove that exactly one independent encoding ($\text{DOF} = 1$, where DOF counts independent storage locations) is the unique rate achieving guaranteed coherence. This optimal point generalizes Rissanen’s Minimum Description Length principle to interactive encoding systems with modification constraints.

Main Results.

- 1) Optimal Rate (Theorem IV.24): $\text{DOF} = 1$ is the unique rate guaranteeing coherence. $\text{DOF} = 0$ means F is not encoded; $\text{DOF} > 1$ permits incoherent configurations where locations disagree on F ’s value.
- 2) Resolution Impossibility (Theorem IV.4): Under incoherence ($\text{DOF} > 1$), no resolution procedure is information-theoretically justified—any oracle selecting a value leaves another value disagreeing. This parallels zero-error capacity constraints: without sufficient side information, error-free decoding is impossible.
- 3) Rate-Complexity Tradeoff (Theorem XII.4): Achieving coherence via manual synchronization requires $\Omega(n)$ update operations; $\text{DOF} = 1$ systems achieve $\mathcal{O}(1)$. The gap grows without bound as $n \rightarrow \infty$.
- 4) Realization Requirements (Theorem VIII.11): Computational systems achieving $\text{DOF} = 1$ via derivation require: (a) definition-time computation hooks, and (b) introspectable derivation results. These are information-theoretic necessities, not implementation preferences.

Instantiations. The abstract encoding model applies across domains:

- Distributed databases: Replica consistency under update constraints
- Configuration systems: Multi-file settings with coherence requirements
- Software systems: Type definitions, class registries, interface contracts
- Version control: Merge resolution under conflicting branches

We evaluate computational realizations (programming languages) against formal criteria derived from the encoding model. Among mainstream systems, Python uniquely satisfies both requirements; most (Java, C++, Rust, Go, TypeScript) cannot achieve $\text{DOF} = 1$ within their design constraints.

Theoretical Foundation. The derivation framework (axes, observational quotients, provenance) is established in the companion paper on classification systems. This work proves the coherence optimality theorem and derives realizability requirements from first principles.

T. Simas is with McGill University, Montreal, QC, Canada (e-mail: tristan.simas@mail.mcgill.ca).

Manuscript received January 16, 2026.

© 2026 Tristan Simas. This work is licensed under CC BY 4.0. License: <https://creativecommons.org/licenses/by/4.0/>

All theorems machine-checked in Lean 4 (9,351 lines, 541 theorems, 0 sorry placeholders). Realizability claims grounded in formalized operational semantics for computational systems.

Index Terms—Encoding theory, coherence constraints, minimum description length, distributed source coding, rate-complexity tradeoffs, zero-error capacity

I. INTRODUCTION

II. INTRODUCTION

A. The Encoding Problem

Consider an information system storing a fact F (e.g., a threshold value, a configuration parameter, or a structural relationship) at n locations. When can such a system guarantee **coherence**—the impossibility of disagreement among locations?

If location L_1 encodes “threshold = 0.5” while location L_2 encodes “threshold = 0.7”, which is correct? No information internal to the system determines this. Any resolution requires external side information (timestamps, priority orderings, external oracles) not present in the encodings themselves.

This is an **information-theoretic** problem: what rate (number of independent encoding locations) guarantees zero-error decoding under modification constraints? We prove that exactly one independent encoding ($\text{DOF} = 1$, where DOF counts independent storage locations) is necessary and sufficient for guaranteed coherence.

Theorem II.1 (Resolution Impossibility, informal). *For any incoherent encoding system ($\text{DOF} > 1$ with divergent values) and any resolution procedure selecting a value, there exists an equally-present value that disagrees. No resolution is information-theoretically justified.*

This parallels zero-error capacity constraints [?], [?]: without sufficient side information, error-free decoding is impossible. Our contribution extends this to **interactive encoding systems** with modification requirements.

B. The Optimal Rate Theorem

We prove that $\text{DOF} = 1$ is the **unique optimal rate** for coherent encoding:

- **DOF = 0:** Fact F is not encoded (no information stored)
- **DOF = 1:** Coherence guaranteed (unique independent source)
- **DOF > 1:** Incoherent configurations reachable (multiple independent sources can diverge)

This generalizes Rissanen’s Minimum Description Length (MDL) principle [1], [2] to systems with update constraints. MDL optimizes description length for static data; we optimize encoding rate for modifiable facts. The singleton solution space—exactly one rate achieves coherence—makes this a **forcing theorem**: given coherence as a requirement, $\text{DOF} = 1$ is mathematically necessary.

C. Applications Across Domains

The abstract encoding model applies wherever facts are stored redundantly:

- **Distributed databases:** Replica consistency under partition constraints [?]
- **Version control:** Merge resolution when branches diverge [?]
- **Configuration systems:** Multi-file settings with coherence requirements [?]
- **Software systems:** Class registries, type definitions, interface contracts [3]

In each domain, the question is identical: given multiple encoding locations, which is authoritative? Our theorems characterize when this question has a unique answer ($\text{DOF} = 1$) versus when it requires arbitrary external resolution ($\text{DOF} > 1$).

D. Connection to Classical Information Theory

Our results extend classical encoding theory in three ways:

1. From static to interactive encoding. Shannon’s source coding theorem [?] characterizes optimal encoding for static data. Slepian-Wolf [?] extends this to distributed sources with side information. We extend to **interactive systems** where encodings can be modified and must remain coherent across modifications.

2. Zero-error requirement with modification constraints.

Classical zero-error capacity [?], [?] characterizes communication without errors. We characterize **encoding without incoherence**—a storage analog where errors are disagreements among locations, not bit flips.

3. Rate-complexity tradeoffs.

Rate-distortion theory [?] trades encoding rate against distortion. We trade encoding rate (DOF) against modification complexity: $\text{DOF} = 1$ achieves $O(1)$ updates; $\text{DOF} > 1$ requires $\Omega(n)$ synchronization.

E. Realizability in Computational Systems

A key question: can the abstract optimality ($\text{DOF} = 1$) be **realized** in computational systems? We prove realizability requires two capabilities:

- 1) **Definition-time derivation:** Derived encodings must be computed when the source is established, not lazily on demand.
- 2) **Introspectable results:** The system must support queries about what was derived and from what source.

These are information-theoretic requirements: systems lacking either cannot achieve $\text{DOF} = 1$ via derivation mechanisms. The proof constructs explicit incoherent configurations reachable in systems missing either capability.

Evaluation across computational systems. We evaluate programming languages, distributed databases, and configuration systems against these criteria. Among mainstream programming languages (TIOBE top-10, 5+ year consistency), Python uniquely satisfies both requirements. Most (Java, C++, Rust, Go, TypeScript) satisfy neither and cannot achieve $\text{DOF} = 1$ for structural facts.

F. Paper Organization and Main Results

This paper establishes four theorems characterizing optimal encoding under coherence constraints. All results are machine-checked in Lean 4 [4] (9,351 lines, 541 theorems, 0 `sorry` placeholders).

Section IV—Encoding Model. We formalize multi-location encoding systems: facts stored at multiple locations with independence and derivability relations. DOF (Degrees of Freedom) counts independent locations. Coherence means all locations agree.

Section VI—Optimal Rate. We prove $\text{DOF} = 1$ is the unique rate guaranteeing coherence (Theorem IV.24). The proof constructs incoherent configurations for all $\text{DOF} > 1$ and shows $\text{DOF} = 1$ makes disagreement impossible.

Section VIII—Realizability. We derive necessary and sufficient conditions for computational systems to achieve $\text{DOF} = 1$ via derivation (Theorem VIII.11). Both definition-time computation and introspectable results are required.

Section X—System Evaluation. We evaluate computational realizations (programming languages, distributed databases) against formal criteria. Python, CLOS, and Smalltalk satisfy requirements; Java, C++, Rust, Go, TypeScript do not.

Section XII—Complexity Bounds. We prove the rate-complexity tradeoff: $\text{DOF} = 1$ achieves $O(1)$ modification cost; $\text{DOF} > 1$ requires $\Omega(n)$ (Theorem XII.4). The gap grows without bound.

Section XIV—Empirical Validation. Case studies from production systems (OpenHCS [5]) demonstrate DOF reduction ($47 \rightarrow 1$) via Python’s definition-time hooks.

Connection to software engineering. The “Don’t Repeat Yourself” (DRY) principle [3] in software engineering is the practitioner recognition of encoding optimality. We prove DRY = $\text{DOF} = 1$, formalizing an informal principle with information-theoretic foundations.

G. Core Theorems

We establish four theorems characterizing optimal encoding under coherence constraints:

1. **Theorem IV.4 (Resolution Impossibility):** For any incoherent encoding system ($\text{DOF} > 1$ with divergent values) and any resolution procedure selecting a value, there exists an equally-present value disagreeing with the selection. No resolution is information-theoretically justified.

Proof: By incoherence, at least two values exist. Any selection leaves another value disagreeing. No side information distinguishes them.

2. **Theorem IV.6 (Optimal Rate):** $\text{DOF} = 1$ guarantees coherence. Exactly one independent encoding makes disagreement impossible.

Proof: All other locations are derived from the single source. Derivation enforces agreement. Single source determines all values.

3. **Theorem VIII.11 (Realizability Requirements):** A computational system achieves $\text{DOF} = 1$ via derivation if and only if it provides: (a) definition-time computation, and (b) introspectable derivation results.

Proof: Necessity by constructing incoherent configurations when either is missing. Sufficiency by exhibiting derivation mechanisms achieving $\text{DOF} = 1$.

4. **Theorem XII.4 (Rate-Complexity Tradeoff):** Modification complexity scales as: $\text{DOF} = 1$ achieves $O(1)$; $\text{DOF} = n > 1$ requires $\Omega(n)$. The ratio grows without bound: $\lim_{n \rightarrow \infty} \frac{n}{1} = \infty$.

Proof: $\text{DOF} = 1$ updates single source (constant). $\text{DOF} = n$ must synchronize n independent locations (linear).

Forcing property. $\text{DOF} = 1$ is the **unique** rate guaranteeing coherence. $\text{DOF} = 0$ means unencoded; $\text{DOF} > 1$ permits incoherence. Given coherence as a requirement, there is no design freedom—the solution is mathematically forced.

H. Scope

This work characterizes SSOT for *structural facts* (class existence, method signatures, type relationships) within *single-language* systems. The complexity analysis is asymptotic, applying to systems where n grows. External tooling can approximate SSOT behavior but operates outside language semantics.

Multi-language systems. When a system spans multiple languages (e.g., Python backend + TypeScript frontend + protobuf schemas), cross-language SSOT requires external code generation tools. The analysis in this paper characterizes single-language SSOT; multi-language SSOT is noted as future work (Section XVIII).

I. Contributions

This paper makes six contributions:

1. Epistemic foundations (Section IV-A):

- Definition of coherence and incoherence for encoding systems
- **Theorem IV.4 (Oracle Arbitrariness):** Under incoherence, no resolution is principled. The epistemic core.
- **Theorem IV.6 (Coherence Forcing):** $\text{DOF} = 1$ guarantees coherence
- **Theorem IV.7:** $\text{DOF} > 1$ permits incoherence
- **Corollary IV.8:** Given coherence requirement, $\text{DOF} = 1$ is necessary and sufficient

2. Software instantiation (Section IV-B):

- Mapping: encoding systems → codebases, facts → structural specifications
- Definition of SSOT as $\text{DOF} = 1$ for software
- Theorem VI.2: SSOT eliminates indeterminacy

3. Language requirements (Section VIII):

- Theorem VIII.7: Definition-time hooks are necessary
- Theorem VIII.9: Introspection is necessary
- Theorem VIII.11: Both together are sufficient
- Proof that these requirements are forced by the structure of the problem

4. Language evaluation (Section X):

- Exhaustive evaluation of 10 mainstream languages
- Extended evaluation of 3 non-mainstream languages (CLOS, Smalltalk, Ruby)
- Theorem X.3: Exactly three languages satisfy requirements

5. Complexity bounds (Section XII):

- Theorem XII.2: SSOT achieves $O(1)$ coherence restoration
- Theorem XII.3: Non-SSOT requires $\Omega(n)$ modifications
- Theorem XII.4: The gap is unbounded

6. Practical demonstration (Section XIV):

- Before/after examples from OpenHCS (production Python codebase)
- PR #44 [6]: Migration from 47 scattered checks to 1 ABC ($\text{DOF } 47 \rightarrow 1$)
- Empirical validation that coherence is achievable in practice

J. Empirical Context: OpenHCS

What it does: OpenHCS [5] is an open-source bioimage analysis platform for high-content screening (45K LoC Python). It processes microscopy images through configurable pipelines, with GUI-based design and Python code export. The system requires:

- Automatic registration of analysis components
- Type-safe configuration with inheritance
- Runtime enumeration of available processors
- Provenance tracking for reproducibility

Why it matters for this paper: OpenHCS requires SSOT for structural facts. When a new image processor is added (by subclassing `BaseProcessor`), it must automatically appear in:

- The GUI component palette
- The configuration schema
- The serialization registry
- The documentation generator

Without SSOT, adding a processor requires updating 4+ locations. With SSOT, only the class definition is needed. Python’s `__init_subclass__` and `__subclasses__()` handle the rest.

Key finding: PR #44 [6] migrated from duck typing (`hasattr()` checks) to nominal typing (ABC contracts). This eliminated 47 scattered checks, reducing DOF from 47 to 1. The migration validates both:

1. The theoretical prediction: DOF reduction is achievable
2. The practical benefit: Maintenance cost decreased measurably

K. Decision Procedure, Not Preference

The contribution of this paper is not the theorems alone, but their consequence: *language selection for coherent representation becomes a decision procedure*.

Given coherence as a requirement:

1. $\text{DOF} = 1$ is necessary (Corollary IV.8)
2. $\text{DOF} = 1$ for structural facts requires definition-time hooks AND introspection (Theorem VIII.11)
3. Languages lacking these features cannot achieve coherence for structural facts

Implications:

1. **Language design.** Future languages should include definition-time hooks and introspection if coherent structural representation is a design goal.
2. **Architecture.** When coherence is required, language selection is constrained. “I prefer Go” is not valid when structural coherence is required.
3. **Tooling.** External tools can work around language limitations but operate outside language semantics.
4. **Pedagogy.** DRY should be taught as epistemic necessity with language requirements, not as a vague guideline.

L. Paper Structure

Section IV establishes epistemic foundations (coherence, oracle arbitrariness) and instantiates them to software. Section VI defines SSOT as the coherent representation and proves its properties. Section VIII derives language requirements with necessity proofs. Section X evaluates mainstream languages exhaustively. Section XII proves complexity bounds. Section XIV demonstrates practical application. Section XVI surveys related work. Appendix P addresses anticipated objections. Appendix A contains complete Lean 4 proof listings.

M. Anticipated Objections

Before proceeding, we address objections readers are likely forming. Each is refuted in detail in Appendix P; here we summarize the key points.

a) “*The model doesn’t capture real Python/Rust semantics.*”: The model is validated through instantiation proofs (§P, Part I). `PythonInstantiation.lean` proves that all Python observables factor through the (B, S) axes. `LangPython.lean` directly encodes Python’s datamodel specification. The model is falsifiable: produce Python code where two types with identical `__bases__` and `__dict__` behave differently, or where `__init_subclass__` fails to fire.

b) “*Rust can achieve SSOT with proc macros and static registries.*”: No. Proc macros are per-item isolated—they cannot see other items during expansion. Registration is bypassable: you can `impl Trait` without any `#[derive]` annotation. The `inventory` crate uses linker tricks external to language semantics. Contrast Python: `__init_subclass__` fires automatically and *cannot be bypassed*. This is enforcement vs. enablement (§P, Part III).

c) “*The requirements are circular—you define structural facts as fixed at definition time, then prove you need definition-time hooks.*”: No. We define structural facts by their *syntactic locus* (encoded in type definitions). The observation that they are fixed at definition time is a *consequence* of this locus. The theorem that hooks are required is *derived* from the observation. The circularity objection mistakes consequence for premise (§P, Part II).

d) “*Build.rs / external tools can achieve SSOT.*”: External tools operate outside language semantics. They can fail, be misconfigured, or be bypassed. They provide no runtime verification—the program cannot confirm derivation occurred. Build tool configuration becomes a second source ($\text{DOF} \geq 2$). This approximates SSOT but does not achieve it (§P, Parts II–III).

e) “*The $\text{DOF} = 1$ definition is too restrictive.*”: The definition is *derived*, not chosen. $\text{DOF} = 0$ means the fact is unrepresented. $\text{DOF} > 1$ means multiple sources can diverge. $\text{DOF} = 1$ is the unique optimal point. Systems with $\text{DOF} > 1$ may be pragmatically acceptable but do not satisfy SSOT (§P, Part II).

f) “*You just need discipline, not language features.*”: Discipline is the external oracle. The theorem states: with $\text{DOF} > 1$, consistency requires an external oracle. “Code review and documentation” are exactly that oracle—human-maintained, fallible, bypassable. Language enforcement cannot be forgotten; human discipline can (§P, Part IV).

g) “*The proofs are trivial (rf1).*”: When modeling is correct, theorems become definitional. This is a feature. Not all proofs are `rf1`: `rust_lacks_introspection` is 40 lines of actual reasoning. The contribution is making the right definitions so that consequences follow structurally (§P, Part V).

If you have an objection not listed above, check Appendix P (16 objections addressed) before concluding it has not been considered.

III. ENCODING SYSTEMS AND COHERENCE

IV. ENCODING SYSTEMS AND COHERENCE

We formalize encoding systems with modification constraints and prove fundamental limits on coherence. The core results apply universally to any domain where facts are encoded at multiple locations and modifications must preserve correctness. Software systems are one instantiation; distributed databases, configuration management, and version control are others.

A. The Encoding Model

We begin with the abstract encoding model: locations, values, and coherence constraints.

Definition IV.1 (Encoding System). An *encoding system* for a fact F is a collection of locations $\{L_1, \dots, L_n\}$, each capable of holding a value for F .

Definition IV.2 (Coherence). An encoding system is *coherent* iff all locations hold the same value:

$$\forall i, j : \text{value}(L_i) = \text{value}(L_j)$$

Definition IV.3 (Incoherence). An encoding system is *incoherent* iff some locations disagree:

$$\exists i, j : \text{value}(L_i) \neq \text{value}(L_j)$$

The Resolution Problem. When an encoding system is incoherent, no resolution procedure is information-theoretically justified. Any oracle selecting a value leaves another value disagreeing, creating an unresolvable ambiguity.

Theorem IV.4 (Oracle Arbitrariness). *For any incoherent encoding system S and any oracle O that resolves S to a value $v \in S$, there exists a value $v' \in S$ such that $v' \neq v$.*

Proof. By incoherence, $\exists v_1, v_2 \in S : v_1 \neq v_2$. Either O picks v_1 (then v_2 disagrees) or O doesn't pick v_1 (then v_1 disagrees).

Interpretation. This theorem parallels zero-error capacity constraints in communication theory. Just as insufficient side information makes error-free decoding impossible, incoherence makes truth-preserving resolution impossible. The encoding system does not contain sufficient information to determine which value is correct. Any resolution requires external information not present in the encodings themselves.

Definition IV.5 (Degrees of Freedom). The *degrees of freedom* (DOF) of an encoding system is the number of locations that can be modified independently.

Theorem IV.6 (DOF = 1 Guarantees Coherence). *If $\text{DOF} = 1$, then the encoding system is coherent in all reachable states.*

Proof. With $\text{DOF} = 1$, exactly one location is independent. All other locations are derived (automatically updated when the source changes). Derived locations cannot diverge from their source. Therefore, all locations hold the value determined by the single independent source. Disagreement is impossible.

Theorem IV.7 (DOF > 1 Permits Incoherence). *If $\text{DOF} > 1$, then incoherent states are reachable.*

Proof. With $\text{DOF} > 1$, at least two locations are independent. Independent locations can be modified separately. A sequence of edits can set $L_1 = v_1$ and $L_2 = v_2$ where $v_1 \neq v_2$. This is an incoherent state. ■

Corollary IV.8 (Coherence Forces DOF = 1). *If coherence must be guaranteed (no incoherent states reachable), then $\text{DOF} = 1$ is necessary and sufficient.*

This is the information-theoretic foundation of optimal encoding under coherence constraints.

Connection to Minimum Description Length. The $\text{DOF} = 1$ optimum directly generalizes Rissanen's MDL principle [1]. MDL states that the optimal representation minimizes total description length: $|\text{model}| + |\text{data given model}|$. In encoding systems:

- **DOF = 1:** The single source is the minimal model. All derived locations are “data given model” with zero additional description length (fully determined by the source). Total encoding rate is minimized.

- **DOF > 1:** Redundant independent locations require explicit synchronization. Each additional independent location adds description length with no reduction in uncertainty—pure overhead serving no encoding purpose.

Grünwald [2] proves that MDL-optimal representations are unique under mild conditions. Theorem IV.24 establishes the analogous uniqueness for encoding systems under modification constraints: $\text{DOF} = 1$ is the unique coherence-guaranteeing rate.

Generative Complexity. Heering [7] formalized this for computational systems: the *generative complexity* of a program family is the length of the shortest generator. $\text{DOF} = 1$ systems achieve minimal generative complexity—the single source is the generator, derived locations are generated instances. This connects our framework to Kolmogorov complexity while remaining constructive (we provide the generator, not just prove existence).

The following sections show how computational systems instantiate this encoding model.

B. Computational Realizations

The abstract encoding model (Definitions IV.1–IV.5) applies to any system where:

- 1) Facts are encoded at multiple locations
- 2) Locations can be modified
- 3) Correctness requires coherence across modifications

Domains satisfying these constraints:

- **Software codebases:** Type definitions, registries, configurations
- **Distributed databases:** Replica consistency under updates
- **Configuration systems:** Multi-file settings (e.g., infrastructure-as-code)
- **Version control:** Merge resolution under concurrent modifications

We focus on *computational realizations*—systems where locations are syntactic constructs manipulated by tools or humans. Software codebases are the primary example, but the encoding model is not software-specific.

Definition IV.9 (Codebase (Software Realization)). A *codebase* C is a finite collection of source files, each containing syntactic constructs (classes, functions, statements, expressions). This is the canonical computational encoding system.

Definition IV.10 (Location). A *location* $L \in C$ is a syntactically identifiable region: a class definition, function body, configuration value, type annotation, database field, or configuration entry.

Definition IV.11 (Modification Space). For encoding system C , the *modification space* $E(C)$ is the set of all valid modifications. Each edit $\delta \in E(C)$ transforms C into $C' = \delta(C)$.

The modification space is large (exponential in system size). But we focus on modifications that *change a specific fact*.

C. Facts: Atomic Units of Specification

Definition IV.12 (Fact). A *fact* F is an atomic unit of program specification: a single piece of knowledge that can be independently modified. Facts are the indivisible units of meaning in a specification.

The granularity of facts is determined by the specification, not the implementation. If two pieces of information must always change together, they constitute a single fact. If they can change independently, they are separate facts.

Examples of facts:

Fact	Description
F_1 : “threshold = 0.5”	A configuration value
F_2 : “PNGLoader handles .png”	A type-to-handler mapping
F_3 : “validate() returns bool”	A method signature
F_4 : “Detector is a subclass of Processor”	An inheritance relationship
F_5 : “Config has field name: str”	A dataclass field
	# Location L3: Documentation
	# Supported formats: png, jpg
	There fact F = “PNGLoader handles png” is encoded at: L1: The class definition (primary encoding)
	L2: The registry dictionary (secondary encoding)
	• L3: The documentation comment (tertiary encoding)

Definition IV.13 (Structural Fact). A fact F is *structural* with respect to encoding system C iff the locations encoding F are fixed at definition time:

$$\text{structural}(F, C) \iff \forall L : \text{encodes}(L, F) \rightarrow L \in \text{DefinitionSyntax}(C)$$

where $\text{DefinitionSyntax}(C)$ comprises declarative constructs that cannot change post-definition without recreation.

Examples across domains:

- **Software:** Class declarations, method signatures, inheritance clauses, attribute definitions
- **Databases:** Schema definitions, table structures, foreign key constraints
- **Configuration:** Infrastructure topology, service dependencies
- **Version control:** Branch structure, merge policies

Key property: Structural facts are fixed at *definition time*. Once defined, their structure cannot change without recreation. This is why structural coherence requires definition-time computation: the encoding locations are only mutable during creation.

Non-structural facts (runtime values, mutable state) have encoding locations modifiable post-definition. Achieving $\text{DOF} = 1$ for non-structural facts requires different mechanisms (reactive bindings, event systems) and is outside this paper’s scope. We focus on structural facts because they demonstrate the impossibility results most clearly.

D. Encoding: The Correctness Relationship

Definition IV.14 (Encodes). Location L encodes fact F , written $\text{encodes}(L, F)$, iff correctness requires updating L when F changes.

Formally:

$$\text{encodes}(L, F) \iff \forall \delta_F : \neg \text{updated}(L, \delta_F) \rightarrow \text{incorrect}(\delta_F(C))$$

where δ_F is an edit targeting fact F .

Key insight: This definition is **forced** by correctness, not chosen. We do not decide what encodes what. Correctness

requirements determine it. If failing to update location L when fact F changes produces an incorrect program, then L encodes F . This is an objective, observable property.

Example IV.15 (Encoding in Practice). Consider a type registry:

```
# Location L1: Class definition
class PNGLoader(ImageLoader):
    format = "png"

# Location L2: Registry entry
LOADERS = {"png": PNGLoader, "jpg": JPGLoader}
```

A configuration value
A type-to-handler mapping
A method signature
An inheritance relationship
A dataclass field

- L_3 : The documentation comment (tertiary encoding)
- If F changes (e.g., to “PNGLoader handles png and apng”), all three locations must be updated for correctness. The program is incorrect if L_2 still says $\{ "png" : \text{PNGLoader}\}$ when the class now handles both formats.

E. Modification Complexity

Definition IV.16 (Modification Complexity).

$$M(C, \delta_F) = |\{L \in C : \text{encodes}(L, F)\}|$$

The number of locations that must be updated when fact F changes.

Modification complexity is the central metric of this paper. It measures the *cost* of changing a fact. A codebase with $M(C, \delta_F) = 47$ requires 47 edits to correctly implement a change to fact F . A codebase with $M(C, \delta_F) = 1$ requires only 1 edit.

Theorem IV.17 (Correctness Forcing). $M(C, \delta_F)$ is the **minimum** number of edits required for correctness. Fewer edits imply an incorrect program.

Proof. Suppose $M(C, \delta_F) = k$, meaning k locations encode F . By Definition IV.14, each encoding location must be updated when F changes. If only $j < k$ locations are updated, then $k - j$ locations still reflect the old value of F . These locations create inconsistencies:

- 1) The specification says F has value v' (new)
- 2) Locations L_1, \dots, L_j reflect v'
- 3) Locations L_{j+1}, \dots, L_k reflect v (old)

By Definition IV.14, the program is incorrect. Therefore, all k locations must be updated, and k is the minimum. ■

F. Independence and Degrees of Freedom

Not all encoding locations are created equal. Some are derived from others.

Definition IV.18 (Independent Locations). Locations L_1, L_2 are *independent* for fact F iff they can diverge. Updating L_1 does not automatically update L_2 , and vice versa.

Formally: L_1 and L_2 are independent iff there exists a sequence of edits that makes L_1 and L_2 encode different values for F .

Definition IV.19 (Derived Location). Location L_{derived} is *derived from* L_{source} iff updating L_{source} automatically updates L_{derived} . Derived locations are not independent of their sources.

Example IV.20 (Independent vs. Derived). Consider two architectures for the type registry:

Architecture A (independent locations):

```
# L1: Class definition
class PNGLoader(ImageLoader) : ...

# L2: Manual registry (independent of L1)
LOADERS = {"png": PNGLoader}
```

Here L_1 and L_2 are independent. A developer can change L_1 without updating L_2 , causing inconsistency.

Architecture B (derived location):

```
# L1: Class definition with registration
class PNGLoader(ImageLoader) :
    format = "png"

# L2: Derived registry (computed from L1)
LOADERS = {cls.format: cls for cls in ImageLoader.__subclasses__()}
```

Here L_2 is derived from L_1 . Updating the class definition automatically updates the registry. They cannot diverge.

Definition IV.21 (Degrees of Freedom).

$$\text{DOF}(C, F) = |\{L \in C : \text{encodes}(L, F) \wedge \text{independent}(L)\}|$$

The number of *independent* locations encoding fact F .

DOF is the key metric. Modification complexity M counts all encoding locations. DOF counts only the independent ones. If all but one encoding location is derived, DOF = 1 even though M may be large.

Theorem IV.22 (DOF = Incoherence Potential). $\text{DOF}(C, F) = k$ implies k different values for F can coexist in C simultaneously. With $k > 1$, incoherent states are reachable.

Proof. Each independent location can hold a different value. By Definition IV.18, no constraint forces agreement between independent locations. Therefore, k independent locations can hold k distinct values. This is an instance of Theorem IV.7 applied to software. ■

Corollary IV.23 (DOF > 1 Implies Incoherence Risk). $\text{DOF}(C, F) > 1$ implies incoherent states are reachable. The codebase can enter a state where different locations encode different values for the same fact.

G. The DOF Lattice

DOF values form a lattice with distinct information-theoretic meanings:

DOF	Encoding Status
0	Fact F is not encoded (no representation)
1	Coherence guaranteed (optimal rate under coherence constraint)
$k > 1$	Incoherence possible (redundant independent encodings)

Theorem IV.24 (DOF = 1 is Uniquely Coherent). For any fact F that must be encoded, $\text{DOF}(C, F) = 1$ is the unique value guaranteeing coherence:

- 1) $\text{DOF} = 0$: Fact is not represented
- 2) $\text{DOF} = 1$: Coherence guaranteed (by Theorem IV.6)
- 3) $\text{DOF} > 1$: Incoherence reachable (by Theorem IV.7)

Proof. This is a direct instantiation of Corollary IV.8 to computational systems:

- 1) $\text{DOF} = 0$ means no location encodes F . The fact is unrepresented.
- 2) $\text{DOF} = 1$ means exactly one independent location. All other encodings are derived. Divergence is impossible. Coherence is guaranteed at optimal rate.
- 3) $\text{DOF} > 1$ means multiple independent locations. By Corollary IV.23, they can diverge. Incoherence is reachable.

Only $\text{DOF} = 1$ achieves coherent representation. This is an information-theoretic optimality condition, not a design preference. ■

V. OPTIMAL REPRESENTATION: DOF = 1

VI. OPTIMAL ENCODING RATE (DOF = 1)

Having established the encoding model (Section IV-A), we now prove that $\text{DOF} = 1$ is the unique optimal rate guaranteeing coherence under modification constraints.

A. DOF = 1 as Optimal Rate

DOF = 1 is not a design guideline. It is the information-theoretically optimal rate guaranteeing coherence for facts encoded in systems with modification constraints.

Definition VI.1 (Optimal Encoding (DOF = 1)). Encoding system C achieves *optimal encoding rate* for fact F iff:

$$\text{DOF}(C, F) = 1$$

Equivalently: exactly one independent encoding location exists for F . All other encodings are derived.

This generalizes the “Single Source of Truth” (SSOT) principle from software engineering to universal encoding theory.

Encoding-theoretic interpretation:

- $\text{DOF} = 1$ means exactly one independent encoding location
- All other locations are derived (cannot diverge from source)
- Incoherence is *impossible*, not merely unlikely

- The encoding rate is minimized subject to coherence constraint

Theorem VI.2 (DOF = 1 Guarantees Determinacy). *If $\text{DOF}(C, F) = 1$, then for all reachable states of C , the value of F is determinate: all encodings agree.*

Proof. By Theorem IV.6, DOF = 1 guarantees coherence. Coherence means all encodings hold the same value. Therefore, the value of F is uniquely determined by the single source. ■

Hunt & Thomas's "single, unambiguous, authoritative representation" [3] (SSOT principle) corresponds precisely to this encoding-theoretic structure:

- **Single:** DOF = 1 (exactly one independent encoding)
- **Unambiguous:** No incoherent states possible (Theorem IV.6)
- **Authoritative:** The source determines all derived values (Definition IV.19)

Our contribution is proving that SSOT is not a heuristic but an information-theoretic optimality condition.

Theorem VI.3 (DOF = 1 Achieves $O(1)$ Update). *If $\text{DOF}(C, F) = 1$, then coherence restoration requires $O(1)$ updates: modifying the single source maintains coherence automatically via derivation.*

Proof. Let $\text{DOF}(C, F) = 1$. Let L_s be the single independent encoding location. All other encodings L_1, \dots, L_k are derived from L_s .

When fact F changes:

- 1) Update L_s (1 edit)
- 2) By Definition IV.19, L_1, \dots, L_k are automatically updated
- 3) Coherence is maintained: all locations agree on the new value

Coherence restoration requires exactly 1 manual update. The number of encoding locations k is irrelevant. Complexity is $O(1)$. ■

Theorem VI.4 (Uniqueness of Optimal Rate). *DOF = 1 is the unique rate guaranteeing coherence. DOF = 0 fails to represent F ; DOF > 1 permits incoherence.*

Proof. By Theorem IV.6, DOF = 1 guarantees coherence. By Theorem IV.7, DOF > 1 permits incoherence.

This leaves only DOF = 1 as coherence-guaranteeing rate. DOF = 0 means no independent location encodes F —the fact is not represented.

Therefore, DOF = 1 is uniquely optimal. This is information-theoretic necessity, not design choice. ■

Corollary VI.5 (Incoherence Under Redundancy). *Multiple independent sources encoding the same fact permit incoherent states. $\text{DOF} > 1 \Rightarrow$ incoherence reachable.*

Proof. Direct application of Theorem IV.7. With $\text{DOF} > 1$, independent locations can be modified separately, reaching states where they disagree. ■

B. Rate-Complexity Tradeoff

The DOF metric creates a fundamental tradeoff between encoding rate and modification complexity.

Question: When fact F changes, how many manual updates are required to restore coherence?

- **DOF = 1:** $O(1)$ updates. The single source determines all derived locations automatically.
- **DOF = $n > 1$:** $\Omega(n)$ updates. Each independent location must be synchronized manually.

This is a *rate-distortion* analog: higher encoding rate ($\text{DOF} > 1$) incurs higher modification complexity. DOF = 1 achieves minimal complexity under the coherence constraint.

Key insight: Many locations may encode F (high total encoding locations), but if DOF = 1, coherence restoration requires only 1 manual update. The derivation mechanism handles propagation automatically.

Example VI.6 (Encoding Rate vs. Modification Complexity). Consider an encoding system where a fact F = "all processors must implement operation P " is encoded at 51 locations:

- 1 abstract specification location
- 50 concrete implementation locations

Architecture A (DOF = 51): All 51 locations are independent.

- Modification complexity: Changing F requires 51 manual updates
- Coherence risk: After $k < 51$ updates, system is incoherent (partial updates)
- Only after all 51 updates is coherence restored

Architecture B (DOF = 1): The abstract specification is the single source; implementations are derived.

- Modification complexity: Changing F requires 1 update (the specification)
- Coherence guarantee: Derived locations update automatically via enforcement mechanism
- The *specification* has a single authoritative source

Computational realization (software): Abstract base classes with enforcement (type checkers, runtime validation) achieve DOF = 1 for contract specifications. Changing the abstract method signature updates the contract; type checkers flag non-compliant implementations.

Note: Implementations are separate facts. DOF = 1 for the contract specification does not eliminate implementation updates—it ensures the specification itself is determinate.

C. Derivation: The Coherence Mechanism

Derivation is the mechanism by which DOF is reduced without losing encodings. A derived location cannot diverge from its source, eliminating it as a source of incoherence.

Definition VI.7 (Derivation). Location L_{derived} is *derived from* L_{source} for fact F iff:

$$\text{updated}(L_{\text{source}}) \rightarrow \text{automatically_updated}(L_{\text{derived}})$$

No manual intervention is required. Coherence is maintained automatically.

Derivation can occur at different times depending on the encoding system:

Derivation Time	Examples Across Domains
Compile/Build time	C++ templates, Rust macros, database infrastructure-as-code compilation
Definition time	Python metaclasses, ORM model registration, schema creation
Query/Access time	Database views, computed columns, lazy evaluation

Structural facts require definition-time derivation. Structural facts (class existence, schema structure, service topology) are fixed when defined. Compile-time may be too early (declarative source not yet parsed). Runtime is too late (structure already immutable). Definition-time is the unique opportunity for structural derivation.

Theorem VI.8 (Derivation Preserves Coherence). *If L_{derived} is derived from L_{source} , then L_{derived} cannot diverge from L_{source} and does not contribute to DOF.*

Proof. By Definition VI.7, derived locations are automatically updated when the source changes. Let L_d be derived from L_s . If L_s encodes value v , then L_d encodes $f(v)$ for some function f . When L_s changes to v' , L_d automatically changes to $f(v')$.

There is no reachable state where $L_s = v'$ and $L_d = f(v)$ with $v' \neq v$. Divergence is impossible. Therefore, L_d does not contribute to DOF. ■

Corollary VI.9 (Derivation Achieves Coherence). *If all encodings of F except one are derived from that one, then $\text{DOF}(C, F) = 1$ and coherence is guaranteed.*

Proof. Let L_s be the non-derived encoding. All other encodings L_1, \dots, L_k are derived from L_s . By Theorem VI.8, none can diverge. Only L_s is independent. Therefore, $\text{DOF}(C, F) = 1$, and by Theorem IV.6, coherence is guaranteed. ■

D. Computational Realizations of $\text{DOF} = 1$

$\text{DOF} = 1$ is achieved across computational domains using definition-time derivation mechanisms. We show examples from software, databases, and configuration systems.

Software: Subclass Registration (Python)

```
class Registry:
    _registry = {}
    def __init_subclass__(cls, **kwargs):
        Registry._registry[cls.__name__] = cls

class PNGHandler(Registry): # Automatically registered
    pass
```

Encoding structure:

- Source: Class definition (declared once)
- Derived: Registry dictionary entry (computed at definition time via `__init_subclass__`)
- DOF = 1: Registry cannot diverge from class hierarchy

Databases: Materialized Views

```
CREATE TABLE users (id INT, name TEXT, created_at DATE);
```

```
CREATE MATERIALIZED VIEW user_count AS
SELECT COUNT(*) FROM users;
```

Encoding structure:

- Source: Base table `users`
- Derived: Materialized view `user_count` (updated on schema generation, refresh)
- DOF = 1: View cannot diverge from base table (consistency guaranteed by DBMS)

Configuration: Infrastructure as Code (Terraform)

```
resource "aws_instance" "app" {
    ami = "ami-12345"
    instance_type = "t2.micro"
}

output "instance_ip" {
    value = aws_instance.app.public_ip
}
```

Encoding structure:

- Source: Resource declaration (authoritative configuration)
- Derived: Output value (computed from resource state)
- DOF = 1: Output cannot diverge from actual resource (computed at apply time)

Common pattern: In all cases, the source is declared once, and derived locations are computed automatically at definition/build/query time. Manual synchronization is eliminated. Coherence is guaranteed by the system, not developer discipline.

VII. INFORMATION-THEORETIC REQUIREMENTS

VIII. REALIZABILITY REQUIREMENTS FOR COMPUTATIONAL SYSTEMS

We now derive the capabilities necessary and sufficient for computational systems to achieve $\text{DOF} = 1$ (optimal encoding rate). This section answers: *What must a computational system provide for $\text{DOF} = 1$ to be realizable?*

The requirements are information-theoretic necessities, not implementation preferences. Programming languages are one domain where these requirements apply; distributed databases and configuration systems are others.

The requirements are derived from the $\text{DOF} = 1$ optimality condition (Theorem IV.24). The proofs establish logical necessity.

As The Foundational Axiom

The derivation rests on one axiom, which follows from how computational systems with declarative structures work:

Axiom VIII.1 (Structural Fixation). *Structural facts are fixed at definition time. After a declarative structure (class, schema, resource) is defined, its structural properties cannot be retroactively changed.*

This axiom applies across computational domains:

Programming languages:

- Once class `Foo` extends `Bar` is defined, `Foo`'s parent is `Bar`, not `Baz`

- Once `def process(self, x: int)` is defined, the signature cannot retroactively become `(self, x: str)`
- Once trait Handler is implemented for PNGDecoder, that relationship is permanent

Databases:

- Once `CREATE TABLE users (id INT, name TEXT)` executes, the table structure is fixed
- Altering the schema requires explicit `ALTER TABLE` (not retroactive modification)
- Foreign key constraints are fixed at creation time

Configuration systems:

- Once `resource "aws_instance" "app"` is declared, the resource type is fixed
- The dependency graph is determined at declaration time

Systems that allow modification (Python's `__bases__`, `ALTER TABLE`) are modifying *future* behavior, not *past* structure. The fact that “table X was created with schema S at time T” is fixed at the moment of creation.

All subsequent theorems are logical consequences of this axiom. Rejecting the axiom requires demonstrating a computational system where structural facts can be retroactively modified—which does not exist.

B. The Timing Constraint

The key insight is that structural facts have a *timing constraint*. Unlike configuration values (which can be changed at any time), structural facts are fixed at specific moments:

Definition VIII.2 (Structural Timing). A structural fact F (class existence, inheritance relationship, method signature) is *fixed* when its defining construct is executed. After that point, the structure cannot be retroactively modified.

In Python, classes are defined when the `class` statement executes:

```
class Detector(Processor): # Structure fixed HERE
    def detect(self, img): ...
```

After this point, Detector's inheritance

In Java, classes are defined at compile time:

```
public class Detector extends Processor {
    public void detect(Image img) { ... }
}
```

Critical Distinction: Compile-Time vs. Definition-Time

These terms are often confused. We define them precisely:

Definition VIII.3 (Compile-Time). *Compile-time* is when source code is translated to an executable form (bytecode, machine code). Compile-time occurs *before the program runs*.

Definition VIII.4 (Definition-Time). *Definition-time* is when a class/type definition is *executed*. In Python, this is *at runtime* when the `class` statement runs. In Java, this is *at compile-time* when `javac` processes the file.

The key insight: **Python's `class` statement is executable code.** When Python encounters:

```
class Foo(Bar):
    x = 1
```

It *executes* code that:

- Creates a new namespace
- Executes the class body in that namespace
- Calls the metaclass to create the class object
- Calls `__init_subclass__` on parent classes
- Binds the name `Foo` to the new class

This is why Python has “definition-time hooks”—they execute when the definition runs.

Java's `class` declaration is *not* executable—it is a static declaration processed by the compiler. No user code can hook into this process.

The timing constraint has profound implications for derivation:

Theorem VIII.5 (Timing Forces Definition-Time Derivation). *Derivation for structural facts must occur at or before the moment the structure is fixed.*

Proof. Let F be a structural fact. Let t_{fix} be the moment F is fixed. Any derivation D that depends on F must execute at some time t_D .

Case 1: $t_D < t_{\text{fix}}$. Then D executes before F is fixed. D cannot derive from F because F does not yet exist.

Case 2: $t_D > t_{\text{fix}}$. Then D executes after F is fixed. D can read F but cannot modify structure derived from F —the structure is already fixed.

Case 3: $t_D = t_{\text{fix}}$. Then D executes at the moment F is fixed. D can both read F and modify derived structures before they are fixed.

Therefore, derivation for structural facts must occur at definition time ($t_D = t_{\text{fix}}$). ■

C. Requirement 1: Definition-Time Computation

Definition VIII.6 (Definition-Time Computation Capability). A computational system has *definition-time computation capability* iff it can execute arbitrary computation when a declarative structure is *created*, not when it is *used*.
cannot be changed

This requirement has theoretical foundations in metaobject protocols [8] and computational reflection [9]. Programming language realizations include Python's metaclasses, CLOS initialization hooks, and Smalltalk's class creation protocol.

Examples across domains:

Programming languages (Python):

Hook	When it executes
<code>__init_subclass__</code>	When a subclass is defined
Metaclass <code>__new__</code> / <code>__init__</code>	When a class using that metaclass is created
Class decorator	Immediately after class body executes
<code>__set_name__</code>	When a descriptor is assigned to a class

Example: `__init_subclass__` registration

```
class Registry:
    _handlers = {}
```

```

def __init_subclass__(cls, format=None,
                     super().__init_subclass__(**kwargs)
                     if format:
                         Registry._handlers[format] = cls

class PNGHandler(Registry, format="png"):
    pass # Automatically registered when class is defined

class JPGHandler(Registry, format="jpg"):
    pass # Automatically registered when class is defined

# Registry._handlers == {"png": PNGHandler,
#                         "jpg": JPGHandler}

```

The registration happens at definition time, not at first use. When the class `PNGHandler` statement executes, `__init_subclass__` runs and adds the handler to the registry.

Theorem VIII.7 (Definition-Time Computation is Necessary). *Achieving $DOF = 1$ for structural facts requires definition-time computation capability.*

Proof. By Theorem VIII.5, derivation for structural facts must occur at definition time. Without definition-time computation, no computation can execute at that moment. Therefore, derivation is impossible. Without derivation, secondary encodings cannot be automatically updated. $DOF > 1$ is unavoidable.

Contrapositive: If a computational system lacks definition-time computation, $DOF = 1$ for structural facts is unrealizable. ■

Computational systems lacking this capability:

Programming languages:

- **Java:** Annotations are metadata, not executable computation. They are processed by external tools (annotation processors), not by the runtime at class definition.
- **C++:** Templates expand at compile time but do not execute arbitrary code. SFINAE and `constexpr` if select branches, not execute computation.
- **Go:** No hook mechanism. Interfaces are implicit. No code runs at type definition.
- **Rust:** Procedural macros run at compile time but are opaque at runtime (see Theorem VIII.14).

Databases (most systems):

- Traditional SQL databases do not execute arbitrary user code at schema definition time
- Triggers execute on *data* operations, not schema operations
- PostgreSQL event triggers are a partial exception but limited in scope

D. Requirement 2: Introspectable Derivation Results

Definition-time computation enables derivation. But achieving verifiable $DOF = 1$ also requires *verification*—the ability to confirm that exactly one independent encoding exists. This requires *computational reflection*—the ability of a system to reason about its own structure [9].

Definition VIII.8 (Introspectable Derivation). Derivation is *introspectable* iff the program can query:

- 1) What structures were derived
- 2) From which source each derived structure came
- 3) What the current state of derived structures is

Python's introspection capabilities:

Query	Python Mechanism
What subclasses exist?	<code>cls.__subclasses__()</code>
What is the inheritance chain?	<code>cls.__mro__</code>
What attributes does a class have?	<code>dir(cls), vars(cls)</code>
What type is this object?	<code>type(obj), isinstance(obj, ...)</code>
What methods are abstract?	<code>cls.__abstractmethods__</code>

Example: Verifying registration completeness

```

def verify_registration():
    """Verify all subclasses are registered."""
    all_subclasses = set(ImageLoader.__subclasses__)
    registered = set(LOADER_REGISTRY.values())

    unregistered = all_subclasses - registered
    if unregistered:
        raise RuntimeError(f"Unregistered loaders: {unregistered}")

```

This verification is only possible because Python provides `__subclasses__()`. In systems without this capability, operators cannot enumerate what structures exist.

Theorem VIII.9 (Introspection is Necessary for Verifiable $DOF = 1$). *Verifying that $DOF = 1$ holds requires introspection capability.*

Proof. Verification of $DOF = 1$ requires confirming exactly one independent encoding exists. This requires:

- 1) Enumerating all locations encoding fact F
- 2) Determining which are independent vs. derived
- 3) Confirming exactly one is independent

Step (1) requires introspection: the system must query what structures exist and what they encode. Without introspection, the system cannot enumerate encodings. Verification is impossible.

Without verifiable $DOF = 1$, system operators cannot confirm optimal encoding holds. They must trust implementation correctness without runtime confirmation. Bugs in derivation logic go undetected until coherence violations occur. ■

Computational systems lacking introspection:

Programming languages:

- **C++:** Cannot query “what types instantiated template `Foo<T>?`”
- **Rust:** Procedural macro expansion is opaque at runtime. Cannot query what was generated.
- **TypeScript:** Types are erased at runtime. Cannot query type relationships.
- **Go:** No type registry. Cannot enumerate types implementing an interface.

Configuration systems:

- Many infrastructure-as-code systems lack introspection: cannot query “what resources derive from this template?”

- Kubernetes Operators can query derived resources via API server (introspection present)

E. Independence of Requirements

The two requirements—definition-time hooks and introspection—are independent. Neither implies the other.

- Theorem VIII.10** (Requirements are Independent). 1)
 A language can have definition-time hooks without introspection
 2) A language can have introspection without definition-time hooks

Proof. (1) **Hooks without introspection:** Rust procedural macros execute at compile time (a form of definition-time hook) but the generated code is opaque at runtime. The program cannot query what the macro generated.

(2) **Introspection without hooks:** Java provides `Class.getMethod()`, `Class.getInterfaces()`, etc. (introspection) but no code executes when a class is defined. Annotations are metadata, not executable hooks.

Therefore, the requirements are independent. ■

F. The Realizability Theorem

Theorem VIII.11 (Necessary and Sufficient Realizability Conditions). A computational system S can achieve verifiable $\text{DOF} = 1$ for structural facts if and only if:

- 1) S provides definition-time computation capability, AND
- 2) S provides introspectable derivation results

Proof. (\Rightarrow) **Necessity:** Suppose S achieves verifiable $\text{DOF} = 1$ for structural facts.

- By Theorem VIII.7, S must provide definition-time computation
- By Theorem VIII.9, S must provide introspection

(\Leftarrow) **Sufficiency:** Suppose S provides both capabilities.

- Definition-time computation enables derivation at the right moment (when structure is fixed)
- Introspection enables verification that all secondary encodings are derived
- Therefore, $\text{DOF} = 1$ is achievable: create one source, derive all others, verify completeness

The if-and-only-if follows. ■

Corollary VIII.12 (DOF-1-Complete Systems). A computational system is DOF-1-complete (can achieve optimal encoding rate for structural facts) iff it satisfies both requirements. Otherwise it is DOF-1-incomplete.

G. The Logical Chain (Summary)

For clarity, we summarize the complete derivation from axiom to realizability theorem:

Axiom VIII.1: Structural facts are fixed at definition time.

↓ (definitional)

Theorem VIII.5: Derivation for structural facts must occur at definition time.

↓ (logical necessity)

Theorem VIII.7: Definition-time computation is necessary for $\text{DOF} = 1$.

Theorem VIII.9: Introspection is necessary for verifiable $\text{DOF} = 1$.

↓ (conjunction)

Theorem VIII.11: A computational system can achieve $\text{DOF} = 1$ iff it

↓ (evaluation)

Corollary: Among programming languages, Python, CLOS, Smalltalk and TypeScript are not.

Every step is machine-checked in Lean 4. The proofs compile with zero `sorry` placeholders. Rejecting this chain requires identifying a specific flaw in the axiom, the logic, or the Lean formalization.

H. Concrete Impossibility Demonstration

We now demonstrate *exactly why* DOF-1-incomplete systems cannot achieve $\text{DOF} = 1$ for structural facts. This is not about system quality—it is about fundamental capabilities that certain systems *cannot provide* within their design constraints.

The Structural Fact: “`PNGHandler` handles `.png` files.”

This fact must be encoded in two places:

- 1) The class definition (where the handler is defined)
- 2) The registry/dispatcher (where format→handler mapping lives)

Python achieves $\text{DOF} = 1$:

```
class ImageHandler:
    _registry = {}

    def __init_subclass__(cls, format=None, **kwargs):
        super().__init_subclass__(**kwargs)
        if format:
            ImageHandler._registry[format] = cls

class PNGHandler(ImageHandler, format="png"):
    def load(self, path): ...
```

$\text{DOF} = 1$. The `format="png"` in the class definition is the *single source*. The registry entry is *derived* automatically by `__init_subclass__`. Adding a new handler requires changing exactly one location.

Java cannot achieve $\text{DOF} = 1$:

```
// File 1: PNGHandler.java
@Handler(format = "png") // Annotation is METADATA!
public class PNGHandler implements ImageHandler {
    public BufferedImage load(String path) { ... }
}
```

```
// File 2: HandlerRegistry.java (SEPARATE SOURCE!)
public class HandlerRegistry {
    static {
        register("png", PNGHandler.class); // Must
        register("jpg", JPGHandler.class);
        // Forgot to add TIFFHandler? Runtime error!
```

```

    }
}

```

$\text{DOF} = 2$. The `@Handler(format = "png")` annotation is *data*, not code. It does not execute when the class is defined. The registry must be maintained separately.

Theorem VIII.13 (Generated Files Are Independent Encodings). *A generated source file constitutes an independent encoding, not a derivation. Therefore, code generation does not achieve $\text{DOF} = 1$.*

Proof. Let F be a structural fact (e.g., “`PNGHandler` handles `.png` files”).

Let E_1 be the annotation: `@Handler(format="png")` on `PNGHandler.java`.

Let E_2 be the generated file: `HandlerRegistry.java` containing `register("png", PNGHandler.class)`.

By Definition IV.21, E_1 and E_2 are both encodings of F iff modifying either can change the system’s behavior regarding F .

Test: If we delete or modify `HandlerRegistry.java`, does the system’s behavior change? **Yes**—the handler will not be registered.

Test: If we modify the annotation, does the system’s behavior change? **Yes**—the generated file will have different content.

Therefore, E_1 and E_2 are independent encodings. $\text{DOF} = 2$.

Formally: if an artifact r is absent from the system’s runtime state (cannot be queried or mutated in-process), then $\text{encodes}(r, F)$ introduces an independent degree of freedom.

The fact that E_2 was *generated from* E_1 does not make it a derivation, because:

- 1) E_2 exists as a separate artifact that can be edited, deleted, or fail to generate
- 2) E_2 must be separately compiled/processed
- 3) The generation process is external to the runtime and can be bypassed

Contrast with Python, where the registry entry exists only in memory, created by the class statement itself. There is no second file. \blacksquare

Why Rust proc macros don’t help:

Theorem VIII.14 (Opaque Expansion Prevents Verification). *If macro/template expansion is opaque at runtime, $\text{DOF} = 1$ cannot be verified.*

Proof. Verification of $\text{DOF} = 1$ requires answering: “Is every encoding of F derived from the single source?”

This requires enumerating all encodings. If expansion is opaque, the system cannot query what was generated.

In Rust, after `#[derive(Handler)]` expands, the program cannot ask “what did this macro generate?” The expansion is compiled into the binary but not introspectable.

Without introspection, the system cannot verify $\text{DOF} = 1$. Optimal encoding may hold but cannot be confirmed. \blacksquare

The Gap is Fundamental:

The distinction is not about syntax quality or language preference. The distinction is about fundamental capabilities:

- **Python:** Class definition *executes code* that creates derived structures *in runtime state*
- **Java:** Class definition *produces metadata* that external tools process into *separate files*
- **Rust:** Macro expansion is *invisible at runtime*—verification impossible

This is a design choice with permanent consequences for realizability. No amount of engineering in Java can make the registry *derived from* the class definition in the $\text{DOF} = 1$ sense, because Java’s execution model provides no mechanism for arbitrary code execution at class definition time.

IX. REALIZATION IN COMPUTATIONAL SYSTEMS

X. REALIZATION EVALUATION: COMPUTATIONAL SYSTEMS

We now evaluate computational systems against the $\text{DOF} = 1$ realizability requirements established in Section VIII. We focus on programming languages as the primary computational domain with detailed evaluation. The criteria apply to any computational system with declarative structures.

This evaluation is exhaustive for mainstream programming languages: we check every mainstream language against formally-defined criteria derived from Theorem VIII.11.

A. Evaluation Criteria

We evaluate systems on four criteria, derived from the realizability requirements:

Criterion	Abbrev	Test
Definition-time hooks	DEF	Can arbitrary code execute when is defined?
Introspectable results	INTRO	Can the program query what was created?
Structural modification	STRUCT	Can hooks modify the structure being modified?
Hierarchy queries	HIER	Can the program enumerate structures/implementers?

DEF and **INTRO** are the two requirements from Theorem VIII.11. **STRUCT** and **HIER** are refinements that distinguish partial from complete realizability.

Scoring (Precise Definitions):

- \checkmark = Full support: The feature is available, usable for $\text{DOF} = 1$, and does not require external tools
- \times = No support: The feature is absent or fundamentally cannot achieve $\text{DOF} = 1$
- \triangle = Partial/insufficient: Feature exists but fails a realizability requirement (e.g., needs external tooling or lacks runtime reach)

Methodology note (tooling exclusions): We exclude capabilities that require external build tools or libraries (annotation processors, Lombok, `reflect-metadata+ts-transformer`, `ts-json-schema-generator`, etc.). Only language-native, runtime-verifiable features count toward realizability.

Note: We use Δ sparingly for mainstream languages only when a built-in mechanism exists but fails realizability requirements (e.g., requires compile-time tooling or lacks runtime reach). For non-mainstream languages in Section X-D, we note partial support where relevant since these languages are not our primary focus. For INTRO, we require *subclass enumeration*: the ability to answer “what classes inherit from X?” at runtime. Java’s `getMethods()` does not satisfy this because it cannot enumerate subclasses without classpath scanning via external libraries.

B. Mainstream Language Definition

Definition X.1 (Mainstream Language). A language is *mainstream* iff it appears in the top 20 of at least two of the following indices consistently over 5+ years:

- 1) TIOBE Index [10] (monthly language popularity)
- 2) Stack Overflow Developer Survey (annual)
- 3) GitHub Octoverse (annual repository statistics)
- 4) RedMonk Programming Language Rankings (quarterly)

This definition excludes niche languages (Haskell, Erlang, Clojure) while including all languages a typical software organization might consider. The 5-year consistency requirement excludes flash-in-the-pan languages.

C. Mainstream Language Evaluation

Language	DEF	INTRO	STRUCT	HIER	DOF-1?
Python	✓	✓	✓	✓	YES
JavaScript	✗	✗	✗	✗	NO
Java	✗	✗	✗	✗	NO
C++	✗	✗	✗	✗	NO
C#	✗	✗	✗	✗	NO
TypeScript	Δ	Δ	✗	✗	NO
Go	✗	✗	✗	✗	NO
Rust	✗	✗	✗	✗	NO
Kotlin	✗	✗	✗	✗	NO
Swift	✗	✗	✗	✗	NO

TypeScript earns Δ for DEF/INTRO because decorators (aligned with ES decorators since TypeScript 5.0, 2023) plus `reflect-metadata` can run at class decoration time and expose limited metadata, but (a) they require `experimentalDecorators` or specific `tsconfig` flags instead of being always-on language features, (b) they cannot enumerate implementers at runtime (no `__subclasses__()` equivalent), and (c) type information is erased at compile time. Consequently DOF = 1 remains unrealizable without external tooling, so the overall verdict stays NO. **Note (as of 2026):** While ES decorators are now standard in JavaScript, they still lack subclass enumeration—the fundamental barrier to DOF = 1 remains.

1) *Python: Full DOF-1 Support:* Python provides all four capabilities:

DEF (Definition-time hooks):

- `__init_subclass__`: Executes when a subclass is defined

- Metaclasses: `__new__` and `__init__` execute at class creation
- Class decorators: Execute immediately after class body

INTRO (Introspection):

- `__subclasses__()`: Returns list of direct subclasses
- `__mro__`: Returns method resolution order
- `type()`, `isinstance()`, `issubclass()`: Type queries
- `dir()`, `vars()`, `getattr()`: Attribute introspection

STRUCT (Structural modification):

- Metaclasses can add/remove/modify class attributes
- `__init_subclass__` can modify the subclass being defined
- Decorators can return a different class entirely

HIER (Hierarchy queries):

- `__subclasses__()`: Enumerate subclasses
 - `__bases__`: Query parent classes
 - `__mro__`: Full inheritance chain
- 2) *JavaScript: No DOF-1 Support:* JavaScript lacks definition-time hooks:

DEF: ✗. No code executes when a class is defined. The `class` syntax is declarative. Decorators (TC39 Stage 3, finalized 2024) exist but cannot access or enumerate subclasses at runtime.

INTRO: ✗. `Object.getPrototypeOf()`, `instanceof` exist but *cannot enumerate subclasses*. No equivalent to `__subclasses__()`.

STRUCT: ✗. Cannot modify class structure at definition time.

HIER: ✗. Cannot enumerate subclasses. No equivalent to `__subclasses__()`.

3) *Java: No DOF-1 Support:* Java’s annotations are metadata, not executable hooks [11]:

DEF: ✗. Annotations are processed by external tools (annotation processors), not by the JVM at class loading. The class is already fully defined when annotation processing occurs.

INTRO: ✗. `Class.getMethods()`, `Class.getInterfaces()`, `Class.getSuperclass()` exist but *cannot enumerate subclasses*. The JVM does not track subclass relationships. External libraries (Reflections, ClassGraph) provide this via classpath scanning, but that is external tooling, not a language feature.

STRUCT: ✗. Cannot modify class structure at runtime. Bytecode manipulation (ASM, ByteBuddy) is external tooling, not language-level support.

HIER: ✗. Cannot enumerate subclasses without external libraries (Reflections, ClassGraph).

Why annotation processors don’t count:

- 1) They run at compile time, not definition time. The class being processed is already fixed
- 2) They cannot modify the class being defined; they generate *new* classes
- 3) Generated classes are separate compilation units, not derived facts within the source
- 4) Results are not introspectable at runtime. You cannot query “was this method generated?”

Why Lombok doesn't count: Lombok approximates SSOT but violates it: the Lombok configuration becomes a second source of truth. Changes require updating both source and Lombok annotations. The tool can fail, be misconfigured, or be bypassed.

4) *C++: No DOF-1 Support:* C++ templates are compile-time, not definition-time [12]:

DEF: ✗. Templates expand at compile time but do not execute arbitrary code. `constexpr` functions are evaluated at compile time but cannot hook into class definition.

INTRO: ✗. No runtime type introspection. RTTI (`typeid`, `dynamic_cast`) provides minimal information. Cannot enumerate template instantiations.

STRUCT: ✗. Cannot modify class structure after definition.

HIER: ✗. Cannot enumerate subclasses. No runtime class registry.

5) *Go: No DOF-1 Support:* Go's design philosophy explicitly rejects metaprogramming [13]:

DEF: ✗. No hook mechanism. Types are defined declaratively. No code executes at type definition.

INTRO: ✗. `reflect` package provides limited introspection but cannot enumerate types implementing an interface.

STRUCT: ✗. Cannot modify type structure.

HIER: ✗. Interfaces are implicit (structural typing). Cannot enumerate implementers.

6) *Rust: No DOF-1 Support:* Rust's procedural macros are compile-time and opaque [14]:

DEF: ✗. Procedural macros execute at compile time, not definition time. The generated code is not introspectable at runtime.

INTRO: ✗. No runtime type introspection. `std::any::TypeId` provides minimal information.

STRUCT: ✗. Cannot modify type structure at runtime.

HIER: ✗. Cannot enumerate trait implementers.

Why procedural macros don't count:

- 1) They execute at compile time, not definition time. The generated code is baked into the binary
- 2) `#[derive(Debug)]` generates code, but you cannot query "does this type derive Debug?" at runtime
- 3) Verification requires source inspection or documentation, not runtime query
- 4) No equivalent to Python's `__subclasses__()`. You cannot enumerate trait implementers

Consequence: Rust achieves *compile-time* DOF = 1 but not *runtime* DOF = 1. For applications requiring runtime reflection (ORMs, serialization frameworks, dependency injection), Rust requires manual synchronization or external codegen tools.

Theorem X.2 (Python Uniqueness in Mainstream). *Among mainstream programming languages, Python is the only language satisfying all DOF-1 realizability requirements.*

Proof. By exhaustive evaluation. We checked all 10 mainstream languages against the four criteria derived from Theorem VIII.11. Only Python satisfies all four. The evaluation is complete. No mainstream language is omitted. ■

D. Non-Mainstream Languages

Three non-mainstream languages also satisfy DOF-1 realizability requirements:

Language	DEF	INTRO	STRUCT	HIER	DOF-1
Common Lisp (CLOS)	✓	✓	✓	✓	YES
Smalltalk	✓	✓	✓	✓	YES
Ruby	✓	✓	Partial	✓	Partially

1) *Common Lisp (CLOS):* CLOS (Common Lisp Object System) provides the most powerful metaobject protocol:

DEF: ✓. The MOP (Metaobject Protocol) allows arbitrary code execution at class definition via `:metaclass` and method combinations.

INTRO: ✓. `class-direct-subclasses`, `class-precedence-list`, `class-slots` provide complete introspection.

STRUCT: ✓. MOP allows complete structural modification.

HIER: ✓. `class-direct-subclasses` enumerates subclasses.

CLOS is arguably more powerful than Python for metaprogramming. However, it is not mainstream by our definition.

2) *Smalltalk:* Smalltalk pioneered many of these concepts:

DEF: ✓. Classes are objects. Creating a class sends messages that can be intercepted.

INTRO: ✓. `subclasses`, `allSubclasses`, `superclass` provide complete introspection.

STRUCT: ✓. Classes can be modified at any time.

HIER: ✓. `subclasses` enumerates subclasses.

3) *Ruby:* Ruby provides hooks but with limitations [15]:

DEF: ✓. `inherited`, `included`, `extended` hooks execute at definition time.

INTRO: ✓. `subclasses`, `ancestors`, `instance_methods` provide introspection.

STRUCT: Partial. Ruby's `inherited` hook receives the subclass *after* its body has been parsed, meaning the hook cannot intercept or transform the class definition as it is being constructed. Contrast with Python's `__init_subclass__`, which executes *during* class creation with access to keyword arguments passed in the class definition (e.g., `class Foo(Base, key=val)`). Ruby can add methods post-hoc via `define_method`, but cannot parameterize class creation or inject attributes before the class body executes.

HIER: ✓. `subclasses` enumerates subclasses (added in Ruby 3.1).

Ruby is close to full DOF-1 support but the structural modification limitations (no parameterized class creation, no pre-body hook) prevent complete realizability for use cases requiring definition-time configuration.

Theorem X.3 (Three-Language Theorem). *Exactly three programming languages in common use satisfy complete DOF-1 realizability requirements: Python, Common Lisp (CLOS), and Smalltalk.*

Proof. By exhaustive evaluation of mainstream and notable non-mainstream languages. Python, CLOS, and Smalltalk satisfy all four criteria. Ruby satisfies three of four (partial

STRUCT). All other evaluated languages fail at least two criteria. ■

E. Implications for System Selection

The evaluation has practical implications for computational system selection:

1. If $\text{DOF} = 1$ for structural facts is required:

- Among programming languages, Python is the only mainstream option
- CLOS and Smalltalk are alternatives if mainstream status is not required
- Ruby is a partial option with workarounds needed
- Consider whether the domain truly requires runtime $\text{DOF} = 1$ or whether compile-time guarantees (Rust, Haskell) suffice

2. If using a $\text{DOF}=1$ -incomplete system:

- External tooling (code generators, linters) can help maintain coherence
- But tooling is not equivalent to native realizability
- Tooling-based coherence cannot be verified at runtime
- Tooling adds build complexity and introduces failure modes

3. For system designers:

- Definition-time computation and introspection should be considered if coherence guarantees are a design goal
- These capabilities have costs (implementation complexity, runtime overhead, security surface) that must be weighed
- The absence of these capabilities is a deliberate design choice with information-theoretic consequences for encodability

XI. COMPLEXITY BOUNDS

XII. RATE-COMPLEXITY BOUNDS

We now prove the rate-complexity bounds that make $\text{DOF} = 1$ optimal. The key result: the gap between $\text{DOF}=1$ -complete and $\text{DOF}=1$ -incomplete architectures is *unbounded*—it grows without limit as encoding systems scale.

A. Cost Model

Definition XII.1 (Modification Cost Model). Let δ_F be a modification to fact F in encoding system C . The *effective modification complexity* $M_{\text{effective}}(C, \delta_F)$ is the number of syntactically distinct edit operations that must be performed manually. Formally:

$$M_{\text{effective}}(C, \delta_F) = |\{L \in \text{Locations}(C) : \text{requires_manual_edit}(L, \delta_F)\}|$$

where $\text{requires_manual_edit}(L, \delta_F)$ holds iff location L must be updated manually (not by automatic derivation) to maintain coherence after δ_F .

Unit of cost: One edit = one syntactic modification to one location. We count locations, not keystrokes or characters. This abstracts over edit complexity to focus on the scaling behavior.

What we measure: Manual edits only. Derived locations that update automatically have zero cost. This distinguishes $\text{DOF} = 1$ systems (where derivation handles propagation) from $\text{DOF} > 1$ systems (where all updates are manual).

B. Upper Bound: $\text{DOF} = 1$ Achieves $O(1)$

Theorem XII.2 ($\text{DOF} = 1$ Upper Bound). *For an encoding system with $\text{DOF} = 1$ for fact F :*

$$M_{\text{effective}}(C, \delta_F) = O(1)$$

Effective modification complexity is constant regardless of system size.

Proof. Let $\text{DOF}(C, F) = 1$. By Definition VI.1, C has exactly one independent encoding location. Let L_s be this single independent location.

When F changes:

- 1) Update L_s (1 manual edit)
- 2) All derived locations L_1, \dots, L_k are automatically updated by the derivation mechanism
- 3) Total manual edits: 1

The number of derived locations k may grow with system size, but the number of *manual* edits remains 1. Therefore, $M_{\text{effective}}(C, \delta_F) = O(1)$. ■

Note on “effective” vs. “total” complexity: Total modification complexity $M(C, \delta_F)$ counts all locations that change. Effective modification complexity counts only manual edits. With $\text{DOF} = 1$, total complexity may be $O(n)$ (many derived locations change), but effective complexity is $O(1)$ (one manual edit).

C. Lower Bound: $\text{DOF} > 1$ Requires $\Omega(n)$

Theorem XII.3 ($\text{DOF} > 1$ Lower Bound). *For an encoding system with $\text{DOF} > 1$ for fact F , if F is encoded at n independent locations:*

$$M_{\text{effective}}(C, \delta_F) = \Omega(n)$$

Proof. Let $\text{DOF}(C, F) = n$ where $n > 1$.

By Definition IV.18, the n encoding locations are independent—updating one does not automatically update the others. When F changes:

- 1) Each of the n independent locations must be updated manually
- 2) No automatic propagation exists between independent locations
- 3) Total manual edits: n

Therefore, $M_{\text{effective}}(C, \delta_F) = \Omega(n)$. ■

D. The Unbounded Gap

Theorem XII.4 (Unbounded Gap). *The ratio of modification complexity between $\text{DOF}=1$ -incomplete and $\text{DOF}=1$ -complete architectures grows without bound:*

$$\lim_{n \rightarrow \infty} \frac{M_{\text{DOF}>1}(n)}{M_{\text{DOF}=1}} = \lim_{n \rightarrow \infty} \frac{n}{1} = \infty$$

Proof. By Theorem XII.2, $M_{\text{DOF}=1} = O(1)$. Specifically, $M_{\text{DOF}=1} = 1$ for any system size.

By Theorem XII.3, $M_{\text{DOF}>1}(n) = \Omega(n)$ where n is the number of independent encoding locations.

The ratio is:

$$\frac{M_{\text{DOF}>1}(n)}{M_{\text{DOF}=1}} = \frac{n}{1} = n$$

As $n \rightarrow \infty$, the ratio $\rightarrow \infty$. The gap is unbounded. ■

Corollary XII.5 (Arbitrary Reduction Factor). *For any constant k , there exists a system size n such that $\text{DOF} = 1$ provides at least $k \times$ reduction in modification complexity.*

Proof. Choose $n = k$. Then $M_{\text{DOF}>1}(n) = n = k$ and $M_{\text{DOF}=1} = 1$. The reduction factor is $k/1 = k$. ■

E. Practical Implications

The unbounded gap has practical implications:

1. DOF = 1 matters more at scale. For small systems ($n = 3$), the difference between 3 edits and 1 edit is minor. For large systems ($n = 50$), the difference between 50 edits and 1 edit is significant.

2. The gap compounds over time. Each modification to fact F incurs the complexity cost. If F changes m times over the system lifetime, total cost is $O(mn)$ with $\text{DOF} > 1$ vs. $O(m)$ with $\text{DOF} = 1$.

3. The gap affects error rates. Each manual edit is an opportunity for error. With n edits, the probability of at least one error is $1 - (1-p)^n$ where p is the per-edit error probability. As n grows, this approaches 1.

Example XII.6 (Error Rate Calculation). Assume a 1% error rate per edit ($p = 0.01$).

Edits (n)	P(at least one error)	Architecture
1	1.0%	DOF = 1
10	9.6%	DOF = 10
50	39.5%	DOF = 50
100	63.4%	DOF = 100

With 50 independent encoding locations ($\text{DOF} = 50$), there is a 39.5% chance of introducing an error when modifying fact F . With $\text{DOF} = 1$, the chance is 1%.

F. Amortized Analysis

The complexity bounds assume a single modification. Over the lifetime of an encoding system, facts are modified many times.

Theorem XII.7 (Amortized Complexity). *Let fact F be modified m times over the system lifetime. Let n be the number of independent encoding locations. Total modification cost is:*

- $\text{DOF} = 1: O(m)$
- $\text{DOF} = n > 1: O(mn)$

Proof. Each modification costs $O(1)$ with $\text{DOF} = 1$ and $O(n)$ with $\text{DOF} = n$. Over m modifications, total cost is $m \cdot O(1) = O(m)$ with $\text{DOF} = 1$ and $m \cdot O(n) = O(mn)$ with $\text{DOF} = n$. ■

For a fact modified 100 times with 50 independent encoding locations:

- $\text{DOF} = 1: 100 \text{ edits total}$

- $\text{DOF} = 50: 5,000 \text{ edits total}$

The $50\times$ reduction factor applies to every modification, compounding over the system lifetime.

XIII. EMPIRICAL VALIDATION

XIV. PRACTICAL DEMONSTRATION

We demonstrate the theoretical results with concrete before/after examples from OpenHCS [5], a production bioimage analysis platform. These examples show how Python’s definition-time hooks achieve SSOT for structural facts.

Methodology: This case study follows established guidelines for software engineering case studies [16]. We use a single-case embedded design with multiple units of analysis (DOF measurements, code changes, maintenance metrics).

The value of these examples is *qualitative*: they show the pattern, not aggregate statistics. Each example demonstrates a specific SSOT mechanism. Readers can verify the pattern applies to their own codebases.

A. SSOT Patterns

Three patterns recur in SSOT architectures:

- 1) **Contract enforcement via ABC:** Replace scattered `hasattr()` checks with a single abstract base class. The ABC is the single source; `isinstance()` checks are derived.
- 2) **Automatic registration via `__init_subclass__`:** Replace manual registry dictionaries with automatic registration at class definition time. The class definition is the single source; the registry entry is derived.
- 3) **Automatic discovery via `__subclasses__()`:** Replace explicit import lists with runtime enumeration of subclasses. The inheritance relationship is the single source; the plugin list is derived.

B. Detailed Examples

We present three examples showing before/after code for each pattern.

1) *Pattern 1: Contract Enforcement (PR #44 [6]):* This example is from a publicly verifiable pull request [6]. The PR eliminated 47 scattered `hasattr()` checks by introducing ABC contracts, reducing DOF from 47 to 1.

The Problem: The codebase used duck typing to check for optional capabilities:

```
# BEFORE: 47 scattered hasattr() checks (DOF = 47)

# In pipeline.py
if hasattr(processor, 'supports_gpu'):
    if processor.supports_gpu():
        use_gpu_path(processor)

# In serializer.py
if hasattr(obj, 'to_dict'):
    return obj.to_dict()

# In validator.py
```

```

if hasattr(config, 'validate'):
    config.validate()

# ... 44 more similar checks across 12 files

Each hasattr() check is an independent encoding of the fact "this type has capability X." If a capability is renamed or removed, all 47 checks must be updated.

The Solution: Replace duck typing with ABC contracts:

# AFTER: 1 ABC definition (DOF = 1)

class GPUCapable(ABC):
    @abstractmethod
    def supports_gpu(self) -> bool: ...

class Serializable(ABC):
    @abstractmethod
    def to_dict(self) -> dict: ...

class Validatable(ABC):
    @abstractmethod
    def validate(self) -> None: ...

# Usage: isinstance() checks are derived from ABC
if isinstance(processor, GPUCapable):
    if processor.supports_gpu():
        use_gpu_path(processor)

```

The ABC is the single source. The `isinstance()` check is derived. It queries the ABC's `__subklasshook__` or MRO, not an independent encoding.

DOF Analysis:

- Pre-refactoring: 47 independent `hasattr()` checks
- Post-refactoring: 1 ABC definition per capability
- Reduction: 47×

2) Pattern 2: Automatic Registration: This pattern applies whenever classes must be registered in a central location.

The Problem: Type converters were registered in a manual dictionary:

```

# BEFORE: Manual registry (DOF = n, where n = number of converters)
# In converters.py
class NumpyConverter:
    def convert(self, data): ...

class TorchConverter:
    def convert(self, data): ...

# In registry.py (SEPARATE FILE - independent)
CONVERTERS = {
    'numpy': NumpyConverter,
    'torch': TorchConverter,
    # ... more entries that must be maintained
}

```

Adding a new converter requires: (1) defining the class, (2) adding to the registry. Two independent edits, violating SSOT.

The Solution: Use `__init_subclass__` for automatic registration:

```

# AFTER: Automatic registration (DOF = 1)

class Converter(ABC):
    _registry = {}

    def __init_subclass__(cls, format=None, **kwargs):
        super().__init_subclass__(**kwargs)
        if format:
            Converter._registry[format] = cls

    @abstractmethod
    def convert(self, data): ...

class NumpyConverter(Converter, format='numpy'):
    def convert(self, data): ...

class TorchConverter(Converter, format='torch'):
    def convert(self, data): ...

# Registry is automatically populated
# Converter._registry == {'numpy': NumpyConverter, 'torch': TorchConverter}

DOF Analysis:

- Pre-refactoring: n manual registry entries (1 per converter)
- Post-refactoring: 1 base class with __init_subclass__
- The single source is the class definition; the registry entry is derived



3) Pattern 3: Automatic Discovery: This pattern applies whenever all subclasses of a type must be enumerated.



The Problem: Plugins were discovered via explicit imports:



```

BEFORE: Explicit plugin list (DOF = n, where n = number of converters)
In plugin_loader.py
from plugins import (
 DetectorPlugin,
 SegmenterPlugin,
 FilterPlugin,
 # ... more imports that must be maintained
)

PLUGINS = [
 DetectorPlugin,
 SegmenterPlugin,
 FilterPlugin,
 # ... more entries that must match the imports
]

```


```

Adding a plugin requires: (1) creating the plugin file, (2) adding the import, (3) adding to the list. Three edits for one fact, violating SSOT.

The Solution: Use `__subclasses__()` for automatic discovery:

```
# AFTER: Automatic discovery (DOF = 1)
```

```

class Plugin(ABC):
    @abstractmethod
    def execute(self, context): ...

# In plugin_loader.py
def discover_plugins():
    return Plugin.__subclasses__()

# Plugins just need to inherit from Plugin
class DetectorPlugin(Plugin):
    def execute(self, context): ...

```

DOF Analysis:

- Pre-refactoring: n explicit entries (imports + list)
- Post-refactoring: 1 base class definition
- The single source is the inheritance relationship; the plugin list is derived

4) *Pattern 4: Introspection-Driven Code Generation:* This pattern demonstrates why both SSOT requirements (definition-time hooks *and* introspection) are necessary. The code is from `openhcs/debug/pickle_to_python.py`, which converts serialized Python objects to runnable Python scripts.

The Problem: Given a runtime object (dataclass instance, enum value, function with arguments), generate valid Python code that reconstructs it. The generated code must include:

- Import statements for all referenced types
- Default values for function parameters
- Field definitions for dataclasses
- Module paths for enums

Without SSOT: Manual maintenance lists

```

# Hypothetical non-introspectable language
IMPORTS = {
    "sklearn.filters": ["gaussian", "sobel"],
    "numpy": ["array"],
    # Must manually update when types change
}

DEFAULT_VALUES = {
    "gaussian": {"sigma": 1.0, "mode": "reflect"},
    # Must manually update when signatures change
}

```

Every type, every function parameter, every enum. Each requires a manual entry. When a function signature changes, both the function *and* the metadata list must be updated. DOF > 1 .

With SSOT (Python): Derive everything from introspection

```

def collect_imports_from_data(data_obj):
    """Traverse structure, derive imports from metadata."""
    if isinstance(obj, Enum):
        # Enum definition is single source
        module = obj.__class__.__module__
        name = obj.__class__.__name__
        enum_imports[module].add(name)

    elif is_dataclass(obj):

```

```

        # Dataclass definition is single source
        function_imports[obj.__class__.__module__]
            .obj.__class__.__name__)
        # Fields are derived via introspection
        for f in fields(obj):
            register_imports(getattr(obj, f.name))

def generate_dataclass_repr(instance):
    """Generate constructor call from field metadata"""
    for field in dataclasses.fields(instance):
        current_value = getattr(instance, field.name)
        # Field name, type, default all come from __init__
        lines.append(f'{field.name}={repr(current_value)}')

```

The Key Insight: The class definition at definition-time establishes facts:

- `@dataclass` decorator → `dataclasses.fields()` returns field metadata
- `Enum` definition → `__module__`, `__name__` attributes exist
- Function signature → `inspect.signature()` returns parameter defaults

Each manual metadata entry is replaced by an introspection query. The definition is the single source; the generated code is derived.

Why This Requires Both SSOT Properties:

- 1) **Definition-time hooks:** The `@dataclass` decorator executes at class definition time, storing field metadata that didn't exist before. Without this hook, `fields()` would have nothing to query.
- 2) **Introspection:** The `fields()`, `__module__`, `inspect.signature()` APIs query the stored metadata. Without introspection, the metadata would exist but be inaccessible.

Impossibility in Non-SSOT Languages:

- **Go:** No decorator hooks, no field introspection. Would require external code generation (separate tool maintaining parallel metadata).
- **Rust:** Procedural macros can inspect at compile-time but metadata is erased at runtime. Cannot query field names from a runtime struct instance.
- **Java:** Reflection provides introspection but no mechanism to store arbitrary metadata at definition-time without annotations (which themselves require manual specification).

The pattern is simple: traverse an object graph, query definition-time metadata via introspection, emit Python code. But this simplicity *depends* on both SSOT requirements. Remove either, and the pattern breaks.

C. Summary

These four patterns (contract enforcement, automatic registration, automatic discovery, and introspection-driven generation) demonstrate how Python's definition-time hooks achieve SSOT for structural facts:

- **PR #44 is verifiable:** The $47 \rightarrow 1$ reduction can be confirmed by inspecting the public pull request.

- **The patterns are general:** Each pattern applies whenever the corresponding structural relationship exists (capability checking, type registration, subclass enumeration, code generation from metadata).
- **The mechanism is the same:** In all cases, the class definition becomes the single source, and secondary representations (registry entries, plugin lists, capability checks, generated code) become derived via Python’s definition-time hooks and introspection.

The theoretical prediction (that SSOT requires definition-time hooks and introspection) is confirmed by these examples. The patterns shown here are instances of the general mechanism proved in Section VIII.

XV. RELATED WORK

XVI. RELATED WORK

This section surveys related work across four areas: the DRY principle, metaprogramming, software complexity metrics, and formal methods in software engineering.

A. The DRY Principle

Hunt & Thomas [3] articulated DRY (Don’t Repeat Yourself) as software engineering guidance in *The Pragmatic Programmer* (1999):

“Every piece of knowledge must have a single, unambiguous, authoritative representation within a system.”

This principle has been widely adopted but never formalized. Our work provides:

- 1) A formal definition of SSOT as $\text{DOF} = 1$
- 2) Proof of what language features are necessary and sufficient
- 3) Machine-checked verification of the core theorems

Comparison: Hunt & Thomas provide guidance; we provide a decision procedure. Their principle is aspirational; our formalization is testable.

B. Metaprogramming and Reflection

Metaobject Protocols: Kiczales et al. [8] established the theoretical foundations for metaobject protocols (MOPs) in *The Art of the Metaobject Protocol* (1991). MOPs allow programs to inspect and modify their own structure at runtime.

Our analysis explains *why* languages with MOPs (CLOS, Smalltalk, Python) are uniquely capable of achieving SSOT: MOPs provide both definition-time hooks and introspection, the two requirements we prove necessary.

Reflection: Smith [9] introduced computational reflection in Lisp. Reflection enables programs to reason about themselves, which is essential for introspectable derivation.

Python Metaclasses: Van Rossum [17] unified types and classes in Python 2.2, enabling the metaclass system that powers Python’s SSOT capabilities. The `__init_subclass__` hook [18] (Python 3.6) simplified definition-time hooks, making SSOT patterns accessible without metaclass complexity.

C. Software Complexity Metrics

Cyclomatic Complexity: McCabe [19] introduced cyclomatic complexity as a measure of program complexity based on control flow. Our DOF metric is orthogonal: it measures *modification* complexity, not *execution* complexity.

Coupling and Cohesion: Stevens et al. [20] introduced coupling and cohesion as design quality metrics. High DOF indicates high coupling (many locations must change together) and low cohesion (related information is scattered).

Code Duplication: Fowler [21] identified code duplication as a “code smell” requiring refactoring. Our DOF metric formalizes this: $\text{DOF} > 1$ is the formal definition of duplication for a fact. Roy & Cordy [22] survey clone detection techniques; Juergens et al. [23] empirically demonstrated that code clones lead to maintenance problems—our DOF metric provides a theoretical foundation for why this occurs.

D. Information Hiding

Parnas [24] established information hiding as a design principle: modules should hide design decisions likely to change. SSOT is compatible with information hiding:

- The single source may be encapsulated within a module
- Derivation exposes only what is intended (the derived interface)
- Changes to the source propagate automatically without exposing internals

SSOT and information hiding are complementary: information hiding determines *what* to hide; SSOT determines *how* to avoid duplicating what is exposed.

E. Formal Methods in Software Engineering

Type Theory: Pierce [25] formalized type systems with machine-checked proofs. Our work applies similar rigor to software engineering principles.

Program Semantics: Winskel [26] formalized programming language semantics. Our formalization of SSOT is in the same tradition: making informal concepts precise.

Verified Software: The CompCert project [27] demonstrated that production software can be formally verified. Our Lean 4 [4] proofs are in this tradition, though at a higher level of abstraction.

Generative Programming: Czarnecki & Eisenecker [28] established generative programming as a paradigm for automatic program generation. Our SSOT patterns are a specific application: generating derived structures from single sources at definition time.

F. Information-Theoretic Foundations

Minimum Description Length: Rissanen [1] established the MDL principle: the best model minimizes total description length (model + data given model). Our $\text{DOF} = 1$ criterion is the software instantiation: the single source is the model; derived locations have zero marginal description length. Grünwald [2] provides comprehensive treatment of MDL optimality and uniqueness—our Theorem IV.24 establishes the analogous uniqueness for software representations.

Generative Complexity: Heering [7], [29] formalized *generative complexity* as the Kolmogorov complexity of the shortest generator for a program family. SSOT architectures achieve minimal generative complexity by design: the single source is the shortest generator, and derived structures are its output. This connects our formalization to algorithmic information theory while remaining constructive—we exhibit generators, not just prove their existence.

Comparison: Prior work on generative complexity remained largely theoretical (Kolmogorov complexity is uncomputable). Our contribution is the *constructive realization*: specific language features (DEF, INTRO) that achieve the theoretical minimum in practice. The DOF metric operationalizes generative complexity for software engineering.

G. Language Comparison Studies

Programming Language Pragmatics: Scott [30] surveys programming language features systematically. Our evaluation criteria (DEF, INTRO, STRUCT, HIER) could be added to such surveys.

Empirical Studies: Prechelt [31] compared programming languages empirically. Our case studies follow a similar methodology but focus on a specific metric (DOF).

H. Novelty of This Work

To our knowledge, this is the first work to:

- 1) Formally define SSOT as $\text{DOF} = 1$
- 2) Prove necessary and sufficient language features for SSOT
- 3) Provide machine-checked proofs of these results
- 4) Exhaustively evaluate mainstream languages against formal criteria
- 5) Measure DOF reduction in a production codebase

The insight that metaprogramming helps with DRY is not new. What is new is the *formalization and proof* that specific features are necessary, and the *machine-checked verification* of these proofs.

XVII. CONCLUSION

XVIII. CONCLUSION

Methodology and Disclosure

Role of LLMs in this work. This paper was developed through human-AI collaboration. The author provided the core intuitions (the DOF formalization, the DEF+INTRO conjecture, the language evaluation criteria), while large language models (Claude, GPT-4) served as implementation partners for drafting proofs, formalizing definitions, and generating LaTeX.

The Lean 4 proofs were iteratively developed: the author specified theorems to prove, the LLM proposed proof strategies, and the Lean compiler verified correctness. This is epistemically sound: a Lean proof that compiles is correct regardless of generation method. The proofs are *costly signals* (per the companion paper on credibility) whose validity is independent of their provenance.

What the author contributed: The $\text{DOF} = 1$ formalization of SSOT, the DEF+INTRO language requirements, the

claim that Python uniquely satisfies these among mainstream languages, the OpenHCS case studies, and the complexity bounds.

What LLMs contributed: LaTeX drafting, Lean tactic exploration, prose refinement, and literature search assistance.

Transparency about this methodology reflects our belief that the contribution is the insight and the verified proof, not the typing labor.

We have provided the first formal foundations for the Single Source of Truth principle. The key contributions are:

1. Formal Definition: SSOT is defined as $\text{DOF} = 1$, where DOF (Degrees of Freedom) counts independent encoding locations for a fact. This definition is derived from the structure of the problem, not chosen arbitrarily.

2. Uniqueness Theorem: We prove that SSOT ($\text{DOF}=1$) is the **unique** minimal representation for structural facts (Theorem VI.4). Any system with $\text{DOF} > 1$ contains redundancy and is therefore non-minimal. This follows from the general uniqueness theorem for minimal complete representations (Paper 1).

3. Language Requirements: We prove that SSOT for structural facts requires (1) definition-time hooks AND (2) introspectable derivation. Both are necessary; both together are sufficient. This is an if-and-only-if theorem.

4. Language Evaluation: Among mainstream languages, only Python satisfies both requirements. CLOS and Smalltalk also satisfy them but are not mainstream. This is proved by exhaustive evaluation.

5. Complexity Bounds: SSOT achieves $O(1)$ modification complexity; non-SSOT requires $\Omega(n)$. The gap is unbounded: for any constant k , there exists a codebase size where SSOT provides at least $k \times$ reduction.

6. Mathematical Necessity: The uniqueness theorem (Theorem VI.4) establishes that $\text{DOF}=1$ is the unique minimal representation: $|\{r : \text{minimal}(r)\}| = 1$. This singleton solution space eliminates design freedom. Claiming “SSOT is a valid design choice among alternatives” while accepting uniqueness instantiates $P \wedge \neg P$: uniqueness entails $\neg \exists$ alternatives with equal minimality; preference presupposes \exists such alternatives. Given minimality as a requirement, the mathematics forces DRY. This is not a guideline—it is the unique solution to the stated constraints.

7. Practical Demonstration: Concrete before/after examples from OpenHCS demonstrate the patterns in practice. PR #44 provides a verifiable example: migration from 47 `hasattr()` checks to ABC contracts, achieving $\text{DOF} 47 \rightarrow 1$.

Implications:

- 1) **For practitioners:** If SSOT for structural facts is required, Python (or CLOS/Smalltalk) is necessary. Other mainstream languages cannot achieve SSOT within the language.
- 2) **For language designers:** Definition-time hooks and introspection should be considered if DRY is a design goal. Their absence is a deliberate choice with consequences.

- 3) **For researchers:** Software engineering principles can be formalized and machine-checked. This paper demonstrates the methodology.

Limitations:

- Results apply to *structural* facts. Configuration values and runtime state have different characteristics.
- The complexity bounds are asymptotic. Small codebases may not benefit significantly.
- Examples are from a single codebase. The patterns are general, but readers should verify applicability to their domains.

Future Work:

- Extend the formalization to non-structural facts
- Develop automated DOF measurement tools
- Study the relationship between DOF and other software quality metrics
- Investigate SSOT in multi-language systems

Connection to Leverage Framework:

SSOT achieves *infinite leverage* in the framework of the companion paper on leverage-driven architecture:

$$L(\text{SSOT}) = \frac{|\text{Derivations}|}{1} \rightarrow \infty$$

A single source derives arbitrarily many facts. This is the theoretical maximum—no architecture can exceed infinite leverage. The leverage framework provides a unified view: this paper (SSOT) and the companion paper on typing discipline selection are both instances of leverage maximization. The metatheorem—“maximize leverage”—subsumes both results.

A. Data Availability

OpenHCS Codebase: The OpenHCS platform (45K LoC Python) is available at <https://github.com/trissim/openhcs> [5]. The codebase demonstrates the SSOT patterns described in Section XIV.

PR #44: The migration from duck typing (`hasattr()`) to ABC contracts is documented in a publicly verifiable pull request [6]: <https://github.com/trissim/openhcs/pull/44>. Readers can inspect the before/after diff to verify the DOF $47 \rightarrow 1$ reduction.

Lean 4 Proofs: The complete Lean 4 formalization (1,753 lines across 13 files, 0 `sorry` placeholders) [32] is included as supplementary material. Reviewers can verify the proofs by running `lake build` in the proof directory.

REFERENCES

- [1] J. Rissanen, “Modeling by shortest data description,” *Automatica*, vol. 14, no. 5, pp. 465–471, 1978, foundational paper on Minimum Description Length (MDL) principle: total description = model + data given model; optimal model minimizes this sum.
- [2] P. D. Grünwald, *The Minimum Description Length Principle*. MIT Press, 2007, comprehensive treatment of MDL; DOF=1 (SSOT) can be framed as the MDL-optimal representation for redundant facts.
- [3] A. Hunt and D. Thomas, *The Pragmatic Programmer: From Journeyman to Master*. Addison-Wesley Professional, 1999.
- [4] L. de Moura and S. Ullrich, “The Lean 4 theorem prover and programming language,” in *Automated Deduction – CADE 28*. Springer, 2021, pp. 625–635.
- [5] T. Simas, “OpenHCS: Open-source high-content screening platform,” 2025, version 0.1.0. Available at: <https://github.com/trissim/openhcs>. [Online]. Available: <https://github.com/trissim/openhcs>
- [6] ———, “UI Anti-Duck-Typing Refactor: ABC-based architecture (PR #44),” GitHub pull request, 2025, merged November 29, 2025. Available at: <https://github.com/trissim/openhcs/pull/44>. [Online]. Available: <https://github.com/trissim/openhcs/pull/44>
- [7] J. Heering, “Generative software complexity,” *Science of Computer Programming*, vol. 97, pp. 82–85, 2015, proposes Kolmogorov complexity to measure software structure; the shortest generator for a set of programs indicates maximal complexity reduction.
- [8] G. Kiczales, J. des Rivière, and D. G. Bobrow, *The Art of the Metaobject Protocol*. MIT press, 1991.
- [9] B. C. Smith, “Reflection and semantics in lisp,” in *Proceedings of the 11th ACM SIGACT-SIGPLAN symposium on Principles of programming languages*, 1984, pp. 23–35.
- [10] TIOBE Software BV, “TIOBE index for programming languages,” TIOBE Programming Community Index, 2024, online: tiobe.com/tiobe-index/.
- [11] J. Gosling, B. Joy, G. Steele, G. Bracha, A. Buckley, D. Smith, and G. Bierman, *The Java Language Specification, Java SE 17 Edition*. Oracle America, Inc., 2021, online: docs.oracle.com/javase/specs/.
- [12] B. Stroustrup, *The C++ Programming Language*, 4th ed. Addison-Wesley Professional, 2013.
- [13] The Go Authors, “The Go programming language specification,” Language specification, 2024, online: go.dev/ref/spec.
- [14] The Rust Team, “The Rust reference,” Language reference, 2024, online: doc.rust-lang.org/reference/.
- [15] D. Flanagan and Y. Matsumoto, *The Ruby Programming Language*. O’Reilly Media, 2008.
- [16] P. Runeson and M. Höst, “Guidelines for conducting and reporting case study research in software engineering,” *Empirical Software Engineering*, vol. 14, no. 2, pp. 131–164, 2009.
- [17] G. van Rossum, “Unifying types and classes in python 2.2,” <https://www.python.org/doc/newstyle/>, 2003.
- [18] M. Teich, R. Hettinger, and Y. Selivanov, “PEP 487 – simpler customisation of class creation,” Python Enhancement Proposals, 2016, online: peps.python.org/pep-0487/.
- [19] T. J. McCabe, “A complexity measure,” *IEEE Transactions on software Engineering*, vol. SE-2, no. 4, pp. 308–320, 1976.
- [20] W. P. Stevens, G. J. Myers, and L. L. Constantine, “Structured design,” *IBM systems journal*, vol. 13, no. 2, pp. 115–139, 1974.
- [21] M. Fowler, *Refactoring: improving the design of existing code*. Addison-Wesley Professional, 1999.
- [22] C. K. Roy and J. R. Cordy, “A survey on software clone detection research,” *School of Computing TR 2007-541, Queen’s University*, vol. 115, pp. 64–68, 2007.
- [23] E. Juergens, F. Deissenboeck, B. Hummel, and S. Wagner, “Do code clones matter?” in *Proceedings of the 31st International Conference on Software Engineering*. IEEE, 2009, pp. 485–495.
- [24] D. L. Parnas, “On the criteria to be used in decomposing systems into modules,” *Communications of the ACM*, vol. 15, no. 12, pp. 1053–1058, 1972.
- [25] B. C. Pierce, *Types and programming languages*. MIT press, 2002.
- [26] G. Winskel, *The Formal Semantics of Programming Languages: An Introduction*. MIT press, 1993.
- [27] X. Leroy, “Formal verification of a realistic compiler,” *Communications of the ACM*, vol. 52, no. 7, pp. 107–115, 2009.
- [28] K. Czarnecki and U. W. Eisenecker, *Generative Programming: Methods, Tools, and Applications*. Addison-Wesley Professional, 2000.
- [29] J. Heering, “Software architecture and software configuration management,” in *Software Configuration Management*, ser. Lecture Notes in Computer Science. Springer, 2003, vol. 2649, pp. 1–16, introduces generative complexity as structural complexity measure for software.
- [30] M. L. Scott, *Programming language pragmatics*. Morgan Kaufmann, 2015.
- [31] L. Prechelt, “An empirical comparison of seven programming languages,” *Computer*, vol. 33, no. 10, pp. 23–29, 2000.
- [32] T. Simas, “Lean 4 formalization: SSOT requirements theorems,” Supplementary material, 2025, 1753 lines across 13 files, 0 `sorry` placeholders. Included with paper submission.
- [33] B. H. Liskov and J. M. Wing, “A behavioral notion of subtyping,” *ACM Transactions on Programming Languages and Systems (TOPLAS)*, vol. 16, no. 6, pp. 1811–1841, 1994.
- [34] W. R. Cook, W. Hill, and P. S. Canning, “Inheritance is not subtyping,” in *Proceedings of the 17th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. ACM, 1989, pp. 125–135.

- [35] J. C. Reynolds, “Types, abstraction and parametric polymorphism,” *Information Processing*, vol. 83, pp. 513–523, 1983.

APPENDIX

All theorems are machine-checked in Lean 4 (9,351 lines across 26 files, 0 `sorry` placeholders, 541 theorems/lemmas). Complete source available at: `proofs/`.

This appendix presents the actual Lean 4 source code from the repository. Every theorem compiles without `sorry`. The proofs can be verified by running `lake build` in the `proofs/` directory.

A. Model Correspondence

What the formalization models: The Lean proofs operate at the level of *abstract language capabilities*, not concrete language semantics. We do not model Python’s specific execution semantics or Java’s bytecode. Instead, we model:

- 1) **DOF as a natural number:** $\text{DOF}(C, F) \in \mathbb{N}$ counts independent encoding locations
- 2) **Language capabilities as propositions:** `HasDefinitionHooks` and `HasIntrospection` are *propositions derived from operational semantics*, not boolean flags. For example, `Python.HasDefinitionHooks` is proved by showing `init_subclass_in_class_definition`, which derives from the modeled `execute_class_statement`.
- 3) **Derivation as a relation:** derives(L_s, L_d) holds when L_d ’s value is determined by L_s

Soundness argument: The formalization is sound if:

- The abstract predicates correspond to actual language features (verified by the evaluation in Section X)
- The derivation relation correctly captures automatic propagation (verified by concrete examples in Section XIV)

What we do NOT model: Performance characteristics, type safety properties, concurrency semantics, or any property orthogonal to SSOT. The model is intentionally narrow: it captures exactly what is needed to prove SSOT requirements, and nothing more.

B. On the Nature of Foundational Proofs

Before presenting the proof listings, we address a potential misreading: a reader examining the Lean source code will notice that many proofs are remarkably short, sometimes a single tactic like `omega` or `exact h`. This brevity is not a sign of triviality. It is characteristic of *foundational* work, where the insight lies in the formalization, not the derivation.

Definitional vs. derivational proofs. Our core theorems establish *definitional* properties and impossibilities, not complex derivations. For example, Theorem VIII.7 (definition-time hooks are necessary for SSOT) is proved by showing that without hooks, updates to derived locations cannot be triggered at definition time. The proof is short because it follows directly from the definition of “definition-time.” If no code executes when a type is defined, then no derivation can occur at that

moment. This is not a complex chain of reasoning; it is an unfolding of what “definition-time” means.

Precedent in foundational CS. This pattern appears throughout foundational computer science:

- **Turing’s Halting Problem (1936):** The proof is a simple diagonal argument, perhaps 10 lines in modern notation. Yet it establishes a fundamental limit on computation that no future algorithm can overcome.
- **Brewer’s CAP Theorem (2000):** The impossibility proof is straightforward: if a partition occurs, a system cannot be both consistent and available. The insight is in the *formalization* of what consistency, availability, and partition-tolerance mean, not in the proof steps.
- **Rice’s Theorem (1953):** Most non-trivial semantic properties of programs are undecidable. The proof follows from the Halting problem via reduction, a few lines. The profundity is in the *generality*, not the derivation.

Why simplicity indicates strength. A definitional requirement is *stronger* than an empirical observation. When we prove that definition-time hooks are necessary for SSOT (Theorem VIII.7), we are not saying “all languages we examined need hooks.” We are saying something universal: *any* language achieving SSOT for structural facts must have hooks, because the logical structure of the problem forces this requirement. The proof is simple because the requirement is forced by the definitions. There is no wiggle room.

Where the insight lies. The semantic contribution of our formalization is:

- 1) **Precision forcing.** Formalizing “degrees of freedom” and “independent locations” in Lean requires stating exactly what it means for two locations to be independent (Definition IV.18). This precision eliminates ambiguity that plagues informal DRY discussions.
- 2) **Completeness of requirements.** Theorem VIII.11 is an if-and-only-if theorem: hooks AND introspection are both necessary and sufficient. This is not “we found two helpful features.” This is “these are the *only* two requirements.” The formalization proves completeness.
- 3) **Universal applicability.** The SSOT requirements apply to *any* language, not just those we evaluated. A future language designer can check their language against these requirements. If it lacks hooks or introspection, SSOT for structural facts is impossible. Not hard, not inconvenient, but *impossible*.

What machine-checking guarantees. The Lean compiler verifies that every proof step is valid, every definition is consistent, and no axioms are added beyond Lean’s foundations. Zero `sorry` placeholders means zero unproven claims. The 8,916 lines across 25 files (519 theorems/lemmas) establish a verified chain from basic definitions (edit space, facts, encoding) through grounded operational semantics (AbstractClassSystem, AxisFramework, NominalResolution, SSOTGrounded) to the final theorems (SSOT requirements, complexity bounds, language evaluation). Reviewers need not trust our informal explanations. They can run `lake build` and verify the proofs themselves.

Comparison to informal DRY guidance. Hunt & Thomas's *Pragmatic Programmer* [3] introduced DRY as a principle 25 years ago, but without formalization. Prior work treats DRY as a guideline, not a mathematical property. Our contribution is making DRY *formal*: defining what it means ($\text{DOF} = 1$), deriving what it requires (hooks + introspection), and proving the claims machine-checkable. The proofs are simple because the formalization makes the structure clear.

This follows the tradition of metatheory: Liskov & Wing [33] formalized behavioral subtyping, Cook et al. [34] formalized inheritance semantics, Reynolds [35] formalized parametricity. In each case, the contribution was not complex proofs, but *precise formalization* that made previously informal ideas mechanically verifiable. Simple proofs from precise definitions are the goal, not a limitation.

C. Basic.lean: Core Definitions (48 lines)

This file establishes the core abstractions. We model DOF as a natural number whose properties we prove directly, avoiding complex type machinery.

```
/-
SSOT Formalization - Basic Definitions
Paper 2: Formal Foundations for the Single-Source-of-Truth Definition
```

Design principle: Keep definitions simple
 DOF and modification complexity are modeled
whose properties we prove abstractly.

```
-/
-- Core abstraction: Degrees of Freedom as a natural number not encoded (missing)
--  $\text{DOF}(C, F) = \text{number of independent locations}$  (optimal)
-- We prove properties about DOF values directly
```

```
-- Key definitions stated as documentation: E. Requirements.lean: Necessity Proofs (113 lines)
-- EditSpace: set of syntactically valid modifications
-- Fact: atomic unit of program specification
-- Encodes(L, F): L must be updated when F changes
-- Independent(L): L can diverge (not derived from another location)
--  $\text{DOF}(C, F) = |\{L : \text{encodes}(L, F) \wedge \text{independent}(L)\}|$  (Language Requirements (Necessary))
-- Theorem 1.6: Correctness Forcing
--  $M(C, \delta_F)$  is the MINIMUM number of edits required for correctness
-- Fewer edits than M leaves at least one encoding location inconsistent
```

Theorem 1.6: $\text{CorrectnessForcing}(M, \delta_F) \rightarrow \forall C, \exists F, \text{DOF}(C, F) \leq M$

Corollary 1.10: $\text{DOF} > 1 \rightarrow \exists C, \exists F, \text{DOF}(C, F) > 1$

```
theorem correctness_forcing (M : Nat) (edits : EditSpace) : edits < M :=
```

```
import Ssot.Derivation
```

```
-- Language feature predicates
```

```
structure LanguageFeatures where
  has_definition_hooks : Bool -- Code executes with hooks
  has_introspection : Bool -- Can query what is known about the state
  has_structural_modification : Bool -- Can change the state
  has_hierarchy_queries : Bool -- Can enumerate subtypes
```

```
theorem dof_inconsistency_potential (k : Nat) : k > 1 := by
  exact kk
```

```
theorem dof_gt_one_inconsistent (dof : Nat) : dof > 1 := by
  exact kk
```

```
theorem dof_gt_one_inconsistent (dof : Nat) : dof > 1 := by
  exact kk
```

```
inductive FactKind where
```

D. SSOT.lean: SSOT Definition (38 lines)

This file defines SSOT and proves its optimality using a simple Nat-based formulation.

```
/-
SSOT Formalization - Single Source of Truth Definition
Paper 2: Formal Foundations for the Single-Source-of-Truth Definition
```

```
-- Definition 2.1: Single Source of Truth
-- SSOT holds for fact F iff  $\text{DOF}(C, F) = 1$ 
def satisfies_SSOT (dof : Nat) : Prop := dof = 1
```

```
-- Theorem 2.2: SSOT Optimality
theorem ssot_optimality (dof : Nat) (h : satisfies_SSOT dof) :
  dof = 1 := by
  exact h
```

```
-- Corollary 2.3: SSOT implies O(1) modification complexity
theorem ssot_implies_constant_complexity (dof : Nat) :
  dof <= 1 := by
  unfold satisfies_SSOT at h
  omega
```

```
Theorem 1.6: Non-SSOT implies potential inconsistency
theorem non_ssot_inconsistency (dof : Nat) (h : Non-SSOT dof) :
  dof > 1 := by
  unfold satisfies_SSOT at h
  omega
```

```
-- Key insight: SSOT is the unique sweet spot
-- for potential inconsistency
-- Lean 4: <= is notation for < or =
-- Lean 4: \/ is not a logical operator
-- Lean 4: /\ is not a logical operator
```

```
-- Theorem 1.6: Correctness Forcing
-- M(C, delta_F) is the MINIMUM number of edits required for correctness
-- Fewer edits than M leaves at least one encoding location inconsistent
```

```
theorem correctness_forcing (M : Nat) (edits : EditSpace) : edits < M :=
```

```
import Ssot.Derivation
```

```
-- Language feature predicates
```

```
structure LanguageFeatures where
  has_definition_hooks : Bool -- Code executes with hooks
  has_introspection : Bool -- Can query what is known about the state
  has_structural_modification : Bool -- Can change the state
  has_hierarchy_queries : Bool -- Can enumerate subtypes
```

```
theorem dof_inconsistency_potential (k : Nat) : k > 1 := by
  exact kk
```

```

| structural -- Fixed at definition time -/
| runtime -- Can be modified at runtime SSOT Formalization - Complexity Bounds
deriving DecidableEq                                     Paper 2: Formal Foundations for the Single Source
                                                     -/
inductive Timing where
| definition -- At class/type definition import Ssot.SSOT
| runtime -- After program starts      import Ssot.Completeness
deriving DecidableEq

-- Theorem 6.1: SSOT Upper Bound (O(1))
def structural_timing : FactKind → Timing           dof = 1 := by
| FactKind.structural => Timing.definition exact h
| FactKind.runtime => Timing.runtime

-- Axiom: Structural facts are fixed at definition time
def structural_timing : FactKind → Timing           dof = 1 := by
| FactKind.structural => Timing.definition exact h
| FactKind.runtime => Timing.runtime

-- Theorem 6.2: Non-SSOT Lower Bound (Omega(n))
def can_derive_at (L : LanguageFeatures) (t : Timing) := nBoelby=
match t with
| Timing.definition => L.has_definition_hooks
| Timing.runtime => true -- All languages-each theorem 6.2.1 Unbounded Complexity Gap
                           theorem complexity_gap_unbounded :
-- Theorem 3.2: Definition-Time Hooks are NECESSARY all bound : Nat, exists n : Nat, n > bound
theorem definition_hooks_necessary (L : LanguageFeatures) :
can_derive_at L Timing.definition = false exact <bound + 1, Nat.lt_succ_self bound>
L.has_definition_hooks = false := by
intro h
simp [can_derive_at] at h
exact h

-- Corollary: The gap between O(1) and O(n) is unbounded
theorem gap_ratio_unbounded (n : Nat) (hn : n > 0)
n / 1 = n := by
simp

-- Theorem 3.4: Introspection is NECESSARY for Verifiable SSOT
def can_enumerate_encodings (L : LanguageFeatures) :
L.has_introspection
theorem introspection_necessary_for_verification-{:per fact change, n = 1}:
can_enumerate_encodings L = false → True := by
L.has_introspection = false := by trivial
intro h
simp [can_enumerate_encodings] at h
exact h

-- Key insight: This is not about "slightly better"
-- It's about constant vs linear complexity - fundamental difference

-- THE KEY THEOREM: Both requirements are independently necessary
theorem both_requirements_independent :
forall L : LanguageFeatures,
(L.has_definition_hooks = true \and L.has_introspection = false) →
can_enumerate_encodings L = false := by
intro L ⟨_, h_no_intro⟩
simp [can_enumerate_encodings, h_no_intro]

theorem both_requirements_independent' :
forall L : LanguageFeatures,
(L.has_definition_hooks = false \and L.has_introspection = true) →
can_derive_at L Timing.definition = false exact by hooks_and_introspection_enable_ssot
intro L ⟨h_no_hooks, _⟩
simp [can_derive_at, h_no_hooks]

-- The language capability claims are derived from formalized operational semantics, not declared as boolean flags. This is the key innovation that forecloses the "trivial proofs" critique.
1) The Proof Chain (Non-Triviality Argument): Consider the claim "Python can achieve SSOT." In the formalization, this is not a tautology. It is the conclusion of a multi-step proof chain:
theorem python_can_achieve_ssot :
canAchieveSSOT Python HasDefinitionHooks Python
Python.python_has_hooks
Python.python_has_introspection


```

F. Bounds.lean: Complexity Bounds (56 lines)

This file proves the $O(1)$ upper bound and $\Omega(n)$ lower bound.

Where `python_has_hooks` is proved from operational semantics:

```
-- From LangPython.lean: __init_subclass__ executed
theorem python_has_hooks : HasDefinitionHooks := by
```



```

def consistent (cfg : MultiEncodingConfig) : Prop
-- DOF > 1 implies there exists an inconsistency
theorem dof_gt_one_implies_inconsistency_possible (n : Nat) (h : n > 1) :
  exists c : ConfigSystem, dof c = n /\ inconsistent (cfg : MultiEncodingConfig) : Prop
-- Contrapositive: guaranteed consistency requires DOF 1 <= 1 in cfg.locations /\ 12 in cfg.locations
theorem consistency_requires_dof_le_one (n : Nat)
  (hall : forall c : ConfigSystem, dof c = n /\ inconsistent (cfg : MultiEncodingConfig) : Prop)
-- DOF = 0 means the fact is not encoded
theorem dof_zero_means_not_encoded (c : ConfigSystem) (h : dof c = 0) :
  Not (encodes_fact c)
-- Independence: updating one location doesn't affect others
theorem update_preserves_other_locations (c : ConfigSystem) (loc other : LocationId)
  (new_val : Value) (h : other != loc) :
  (update_location c loc new_val).value_at = same_as_heir_of_different_provenance :
  exists T1 T2 : Typ, shapeEquivalent T1 T2 /\ typeIdentityEncoding T1 != typeIdentityEncoding T2
-- Oracle necessity: valid oracles can disagree
theorem resolution_requires_external_choice :
  exists o1 o2 : Oracle, valid_oracle o1 /\ valid_oracle o2 /\ o1 c 11 12 != o2 c 11 12
theorem ssot_unique_complete_consistent :
  forall dof : Nat,
  dof != 0 → -- Complete: fact is encoded
  (forall cfg : MultiEncodingConfig, cfg.dof = dof) satisfies_SSOT dof
-- The trichotomy: every DOF is incomplete, optimal
theorem dof_trichotomy : forall dof : Nat,
  dof = 0 \wedge satisfies_SSOT dof \vee
  (exists cfg : MultiEncodingConfig, cfg.dof = dof) satisfies_SSOT dof

```

J. SSOTGrounded.lean: Bridging SSOT to Operational Semantics (184 lines)

This file is the key innovation addressing the “trivial proofs” critique. It bridges the abstract SSOT definition ($\text{DOF} = 1$) to concrete operational semantics from AbstractClassSystem. The central insight: SSOT failures arise when the same fact has multiple independent encodings that can diverge.

/-

SSOTGrounded: Connecting SSOT to Operational Semantics

This file bridges the abstract SSOT definition ($\text{DOF} = 1$) to the concrete operational semantics from AbstractClassSystem. The key insight: SSOT failures arise when the same fact has multiple independent encodings that can diverge.

-/

```

import Ssot.AbstractClassSystem
import Ssot.SSOT

namespace SSOTGrounded

-- A fact encoding location in a configuration
structure EncodingLocation where
  id : Nat
  value : Nat
  deriving DecidableEq

-- A configuration with potentially multiple encodings
structure MultiEncodingConfig where
  locations : List EncodingLocation
  dof : Nat := locations.length

-- All encodings agree on the value

```

Why this matters: The `ssot_unique_complete_consistent` theorem proves that $\text{DOF} = 1$ is the *unique* configuration class that is both complete (fact is encoded) and guarantees consistency (no observer can see different values). This is not a tautology—it is a constructive proof that any $\text{DOF} \geq 2$ admits an inconsistent configuration.

The `same_shape_different_provenance` theorem connects to Paper 1’s capability analysis: shape-based typing loses the Bases axis, so two types with identical shapes can have different provenance. This is precisely the information loss that causes SSOT violations when type identity facts have $\text{DOF} > 1$.

K. AbstractClassSystem.lean: Operational Semantics (3,276 lines)

This file provides the grounded operational semantics that make the SSOT proofs non-trivial. It imports directly from Paper 1’s formalization, ensuring consistency across the paper sequence. Key definitions include:

- **Typ:** Types with namespace (Σ) and bases list, modeling both structural and nominal information.
- **shapeEquivalent:** Two types are shape-equivalent iff they have the same namespace (structural view).

- **Capability enumeration:** Identity, provenance, enumeration, conflict resolution, interface checking.
- **Language instantiations:** Python, Java, Rust, TypeScript with their specific capability profiles.

The central result is the *capability gap theorem*: shape-based observers cannot distinguish types that differ only in their bases. This formally establishes that structural typing loses information, which is the root cause of SSOT violations for type identity facts.

L. *AxisFramework.lean*: Axis-Parametric Theory (1,721 lines)

This file establishes the mathematical foundations of axis-parametric type systems. Key results include:

- **Domain-driven impossibility:** Given any domain D , `requiredAxesOf D` computes the axes D needs. Missing any derived axis implies impossibility—not implementation difficulty, but information-theoretic impossibility.
- **Fixed vs. parameterized asymmetry:** Fixed-axis systems guarantee failure for some domains; parameterized systems guarantee success for all domains.
- **Capability lattice:** Formal ordering of type systems by capability inclusion with Python at the top (full capabilities) and duck typing at the bottom.

M. *NominalResolution.lean*: Resolution Algorithm (609 lines)

Machine-checked proofs for the dual-axis resolution algorithm:

- **Resolution completeness** (Theorem 7.1): The algorithm finds a value if one exists.
- **Provenance preservation** (Theorem 7.2): Uniqueness and correctness of provenance tracking.
- **Normalization idempotence** (Invariant 4): Repeated normalization is identity.

N. *ContextFormalization.lean*: Greenfield/Retrofit (215 lines)

Proves that the greenfield/retrofit classification is decidable and that provenance requirements are detectable from system queries. This eliminates potential circularity concerns by deriving requirements from observable behavior.

O. *DisciplineMigration.lean*: Discipline vs Migration (142 lines)

Formalizes the distinction between discipline optimality (abstract capability comparison, universal) and migration optimality (practical cost-benefit, context-dependent). This clarifies that capability dominance is separate from migration cost analysis.

P. Verification Summary

File	Lines	Key Theorems
<i>Core SSOT Framework</i>		
Basic.lean	47	3
SSOT.lean	37	3
Derivation.lean	66	2
Requirements.lean	112	5
Completeness.lean	167	11
Bounds.lean	80	5
<i>Grounded Semantics (from Paper 1)</i>		
AbstractClassSystem.lean	3,276	45
AxisFramework.lean	1,721	89
NominalResolution.lean	609	31
ContextFormalization.lean	215	8
DisciplineMigration.lean	142	7
<i>SSOT Bridge</i>		
SSOTGrounded.lean	184	6
Foundations.lean	364	15
Inconsistency.lean	224	12
Coherence.lean	264	8
CaseStudies.lean	148	4
<i>Language Instantiations</i>		
Languages.lean	108	6
LangPython.lean	234	10
LangRust.lean	254	8
LangStatic.lean	187	5
LangEvaluation.lean	160	12
Dof.lean	82	4
PythonInstantiation.lean	249	8
JavaInstantiation.lean	63	2
RustInstantiation.lean	64	2
TypeScriptInstantiation.lean	65	2
Total (26 files)	9,351	541

All 541 theorems/lemmas compile without `sorry` placeholders. The proofs can be verified by running `lake build` in the `proofs/` directory. Every theorem in the paper corresponds to a machine-checked proof.

Grounding note: The formalization includes five major proof files from Paper 1 (`AbstractClassSystem`, `AxisFramework`, `NominalResolution`, `ContextFormalization`, `DisciplineMigration`) that provide the grounded operational semantics. This ensures that SSOT claims are not “trivially true by definition” but rather derive from a substantial formal model of type system capabilities.

Key grounded results:

- 1) **Capability gap theorem** (`AbstractClassSystem`): Shape-based observers cannot distinguish types with different bases.
- 2) **Axis impossibility theorems** (`AxisFramework`): Missing axes guarantee incompleteness for some domains.
- 3) **Resolution completeness** (`NominalResolution`): Dual-axis resolution is complete and provenance-preserving.
- 4) **Consistency is non-trivial:** $\text{DOF} \geq 2$ admits inconsistent configurations (constructive witness in `Inconsis-`

tency.lean).

- 5) **SSOT is uniquely optimal:** No other DOF value is both complete and guaranteed-consistent.
- 6) **Language claims derive from semantics:** `python_can_achieve_ssot` chains through `python_has_hooks` to `init_subclass_in_class_definition` to `execute_class_statement`—not boolean flags.
- 7) **Rust impossibility is substantive:** `rust_lacks_introspection` is a 40-line proof by contradiction, not definitional unfolding.

These grounded proofs connect the abstract DOF formalization to concrete operational semantics, ensuring the SSOT theorems have substantial content that cannot be dismissed as definitional tautologies.

This appendix addresses anticipated objections, organized thematically. Each objection is stated in its strongest form, then refuted.

Q. Objection: The Model Doesn't Capture Real Semantics

Objection: “You’ve formalized a toy model and proved properties about it. But the model doesn’t capture real Python/Rust semantics. The proofs are valid but vacuously true about artificial constructs.”

Response: The model is validated through *instantiation proofs* that bridge abstract theorems to concrete language semantics. This is the standard methodology for programming language formalization [25].

a) *The Two-Layer Architecture*.: Paper 1 establishes this methodology:

- 1) **Abstract layer:** Define the (B, S) model for any language with explicit inheritance
- 2) **Instantiation layer:** Prove that concrete language features map to the abstract model
- 3) **Theorem transfer:** Abstract theorems apply to the instantiation

b) *Python Instantiation Proofs*.: The file `PythonInstantiation.lean` (250 LOC) proves:

```
-- Python's type() factors into (B, S)
theorem python_type_is_two_axis (pt : PythonType) :
    exists B S, pythonTypeAxes pt = (B, S)

-- All observables factor through axes
lemma observables_factor_through_axes {p q} :
    (h : sameAxes p q) (attr : AttrName) :
    metaclassOf p = metaclassOf q /\ 
    getAttrHas p attr = getAttrHas q attr
```

The second theorem is the key: if two types have identical `__bases__` and `__dict__`, they are observationally indistinguishable. This proves the model captures Python’s observable behavior.

c) *Semantic Correspondence*.: The `LangPython.lean` file (235 LOC) directly transcribes Python’s documented semantics for class creation:

```
-- Class definition events (from Python datamodel)
```

```
inductive ClassDefEvent where
| metacall_start : PyId -> ClassDefEvent
| new_called : PyId -> ClassDefEvent
| namespace_populated : PyId -> ClassDefEvent
| init_subclass_called : PyId -> PyId -> ClassDefEvent
| subclasses_updated : PyId -> PyId -> ClassDefEvent
| init_called : PyId -> ClassDefEvent
| class_bound : PyId -> ClassDefEvent
```

The theorem `init_subclass_in_class_definition` is then *derived* from this semantics—not assumed. The model is a direct encoding of Python’s specification.

d) *Falsifiability*.: The model makes testable predictions. To falsify it, produce Python code where:

- Two types with identical `__bases__` and `__dict__` behave differently, or
- A subclass exists that is not in `__subclasses__()`, or
- `__init_subclass__` does not fire during class definition

The model is empirically vulnerable. No counterexample has been produced.

e) *The Interpretation Gap*.: Every formalization eventually requires interpretation to connect symbols to reality. The claim is not “this Lean code IS Python” but “this Lean code models Python’s observable behavior with sufficient fidelity that theorems transfer.” The instantiation proofs establish this transfer.

R. Objection: The SSOT Definition is Too Narrow

Objection: “Your definition of SSOT as $\text{DOF} = 1$ is too restrictive. Real-world systems have acceptable levels of duplication.”

Response: The definition is **derived**, not chosen. $\text{DOF} = 1$ is the unique optimal point:

DOF	Meaning
0	Fact is not encoded (underspecification)
1	Single source of truth (optimal)
>1	Multiple sources can diverge (inconsistency risk)

$\text{DOF} = 2$ means two locations can hold different values for the same fact. The *possibility* of inconsistency exists. The definition is mathematical: SSOT requires $\text{DOF} = 1$. Systems with `__DOFType` may be pragmatically acceptable but do not satisfy SSOT.

S. External Tools vs Language-Level SSOT

External tools (annotation processors, code generators, build systems) can approximate SSOT behavior. These differ from language-level SSOT in three dimensions:

- 1) **External to language semantics:** Build tools can fail, be misconfigured, or be bypassed. They operate outside the language model.
- 2) **No runtime verification:** The program cannot *conform* that derivation occurred correctly. Python’s

`__subclasses__()` verifies registration completeness at runtime. External tools provide no runtime guarantee.

- 3) **Configuration-dependent:** External tools require project-specific setup. Python's `__init_subclass__` works in any environment without configuration.

The analysis characterizes SSOT *within language semantics*, where $\text{DOF} = 1$ holds at runtime.

T. Objection: The Requirements Are Circular

Objection: “You define ‘structural fact’ as ‘fixed at definition time,’ then prove you need definition-time hooks. The conclusion is embedded in the definition—this is circular.”

Response: The definition does not assume definition-time hooks; it defines what structural facts *are*. The derivation has three distinct steps:

- 1) **Definition:** A fact F is *structural* iff it is encoded in the syntactic structure of type definitions (class existence, method signatures, inheritance relationships). This is a classification, not a requirement.
- 2) **Observation:** Structural facts are fixed when types are defined. This follows from what “syntactic structure” means—you cannot change a class’s bases after the `class` statement completes.
- 3) **Theorem:** Coherent derivation of structural facts requires hooks that execute at definition time. This is the actual result—it follows from the observation, not from the definition.

The circularity objection mistakes a *consequence* for a *premise*. We do not define structural facts as “requiring definition-time hooks.” We define them by their syntactic locus, observe when they become fixed, and derive the necessary language features.

To reject this, you would need to show either:

- Structural facts are NOT fixed at definition time (provide a counterexample), or
- Coherent derivation can occur without definition-time hooks (provide the mechanism)

Neither objection has been sustained.

U. Derivation Order

The analysis proceeds from definition to language evaluation:

- 1) Define SSOT mathematically ($\text{DOF} = 1$)
- 2) Prove necessary language features (definition-time hooks + introspection)
- 3) Evaluate languages against derived criteria
- 4) Result: Python, CLOS, and Smalltalk satisfy both requirements

Three languages satisfy the criteria. Two (CLOS, Smalltalk) are not mainstream. This validates that the requirements characterize a genuine language capability class. The requirements are derived from SSOT’s definition, independent of any particular language’s feature set.

V. Empirical Validation

The case studies demonstrate patterns, with publicly verifiable instances:

- PR #44: `47 hasattr()` checks $\rightarrow 1$ ABC definition (verifiable via GitHub diff)
- Three general patterns: contract enforcement, automatic registration, automatic discovery
- Each pattern represents a mechanism, applicable to codebases exhibiting similar structure

The theoretical contribution is the formal proof. The examples demonstrate applicability.

W. Asymptotic Analysis

The complexity bounds are derived from the mechanism:

- SSOT: changing a fact requires 1 edit (the single source)
- Non-SSOT: changing a fact requires n edits (one per encoding location)
- The ratio $n/1$ grows unbounded as n increases

PR #44 demonstrates the mechanism at $n = 47$: `47 hasattr()` checks $\rightarrow 1$ ABC definition. The $47 \times$ reduction is observable via GitHub diff. The gap widens as codebases grow.

X. Cost-Benefit Analysis

SSOT involves trade-offs:

- **Benefit:** Modification complexity $O(1)$ vs $\Omega(n)$
- **Cost:** Metaprogramming complexity, potential performance overhead

The analysis characterizes what SSOT requires. The decision to use SSOT depends on codebase scale and change frequency.

Y. Machine-Checked Formalization

The proofs formalize definitions precisely. Machine-checked proofs provide:

- 1) **Precision:** Lean requires every step to be explicit
- 2) **Verification:** Computer-checked, eliminating human error
- 3) **Reproducibility:** Anyone can run the proofs and verify results

The contribution is formalization itself: converting informal principles into machine-verifiable theorems. Simple proofs from precise definitions are the goal.

Z. Build Tool Analysis

External build tools shift the SSOT problem:

- 1) **DOF ≥ 2 :** Build tool configuration becomes a second source. Let C be codebase, T be tool. Then $\text{DOF}(C \cup T, F) \geq 2$ because both source and config encode F .
- 2) **No runtime verification:** Generated code lacks derivation provenance. Cannot query “was this method generated or hand-written?”
- 3) **Cache invalidation:** Build tools must track dependencies. Stale caches cause bugs absent from language-native derivation.

- 4) **Build latency:** Every edit requires build step. Language-native SSOT (Python metaclasses) executes during import.

External tools reduce DOF from n to k where k is the number of tool configurations. Since $k > 1$, SSOT (DOF = 1) is not satisfied.

Cross-language code generation (e.g., protobuf) requires external tools. The analysis characterizes single-language SSOT.

. Objection: Inconsistency Is Only in Comments

Objection: “The proofs don’t formalize ‘inconsistency’—it only appears in comments. The heavy lifting is done by the comments, not by the formal system.”

Response: This critique was valid for earlier versions. We have since added `Ssot/Inconsistency.lean` (216 LOC, zero sorry), which formalizes inconsistency as a Lean Prop:

```
structure ConfigSystem where
  num_locations : Nat
  value_at : LocationId -> Value

def inconsistent (c : ConfigSystem) : Prop
  exists l1 l2, l1 < c.num_locations /\ l2
    l1 != l2 /\ c.value_at l1 != c.value_at l2
```

The file proves:

- 1) **DOF > 1 implies inconsistency possible:**
`dof_gt_one_implies_inconsistency_possible`
 we constructively exhibit an inconsistent configuration for any $n > 1$.
- 2) **Guaranteed consistency requires DOF ≤ 1:**
`consistency_requires_dof_le_one`
 contrapositive of the above.
- 3) **DOF = 0 means the fact is not encoded:**
`dof_zero_means_not_encoded`—no locations implies the system cannot represent the value.
- 4) **Independence formalized:**
`update_preserves_other_locations`—
 updating one location does not affect others, formalizing what “independent” means.
- 5) **Oracle necessity:** `resolution_requires_external`
 when locations disagree, there exist valid oracles that give different resolutions. Therefore, resolving disagreement requires an external, arbitrary choice. The system itself provides no basis to prefer one value over another.

This addresses the critique directly: inconsistency is now a formal property that Lean knows about, not a comment. The interpretation “this models real configuration systems” still requires mapping to reality, but every formalization eventually bottoms out in interpretation. The contribution is making assumptions *explicit and attackable*, not eliminating interpretation entirely.

. Objection: What About the Type’s Name?

Objection: “Your two-axis model (B, S) ignores the type’s name. Isn’t N (the name) a third independent axis?”

Response: No. N is not an independent axis—it is a slot on the type object, set at definition time and immutable thereafter. Technically, `__name__` is stored on the `PyTypeObject` struct (a C-level slot), not in `__dict__`. However, this does not make it independent:

- 1) **N is fixed at definition time.** The name is set by the `class` statement and cannot be changed without creating a new type.
- 2) **N does not affect behavior.** Two classes with identical `__bases__` (B) and `__dict__` (S) behave identically. The name is a label, not an axis of variation.
- 3) **N is observable but not discriminating.** You can query `cls.__name__`, but no Python code changes behavior based on it (except for debugging/logging).

The Lean formalization (`AbstractClassSystem.lean`) captures this:

```
-- N is just a label for a (B, S) pair
-- N contributes no observables beyond B
-- Theorem obs_eq_bs proves: (B, S) equality suffi
```

The operational test: given two classes with identical `__bases__` (B) and identical `__dict__` (S), can any Python code distinguish them behaviorally? No. The name is metadata, not a degree of freedom for the type’s semantics. This is why the model is (B, S) and not (B, S, N). N is a fixed label assigned at definition, not an independent axis that can vary.

. Objection: Model Doesn’t Mirror Compiler Internals

Objection: “Your Rust model (`RuntimeItem`, erasure) doesn’t mirror rustc’s actual HIR→MIR phases. You haven’t modeled proc-macro hygiene, # [link_section] retention, or the actual expander traces.”

Response: We model *observable behavior*, not compiler implementation. The claim is:

At runtime, you cannot distinguish hand-written code from macro-generated code.

This is empirically testable. Challenge: produce Rust code that, at runtime, recovers whether a given struct was written by a human or expanded by a macro—with external metadata, build logs, or source access.

The model’s `RuntimeItem` having no source field is *observationally accurate*: real Rust binaries contain no such field. Whether rustc internally tracks provenance during compilation is irrelevant; what matters is that this information is not preserved in the final artifact.

If the model is wrong, show the Rust code that falsifies it. The burden is on the critic to produce the counterexample.

. Objection: Rust Proc Macros + Static Registries Achieve DOF = 1

Objection: “Rust can achieve DOF = 1 using proc macros and static registries. Example:

```
# [derive(AutoRegister)]
struct MyHandler;
static HANDLERS: &[&dyn Handler] = &[&MyHandler, ...]
```

The macro generates the registry at compile time. There's no second DOF that can diverge."

Response: This conflates *enabling* a pattern with *enforcing* it. The critical distinction:

- 1) **Proc macros are per-item isolated.** When `# [derive(AutoRegister)]` executes on `MyHandler`, it cannot see `OtherHandler`. Each macro invocation is independent—there is no shared state during compilation. Therefore, no single macro can generate a complete registry.

- 2) **Registration is bypassable.** You can write:

```
struct SneakyHandler;
impl Handler for SneakyHandler { ... }
```

The `impl` exists; the registry entry does not. **DOF = 2:** the `impl` and the registry are independent locations that can disagree.

- 3) **The inventory crate uses linker tricks, not language semantics.** It works by emitting items into special linker sections and collecting them at link time. This is:

- Platform-specific (different on Linux, macOS, Windows)
- Not enforced—you can `impl Trait` without `# [inventory::submit]`
- External to language semantics (depends on linker behavior)

Contrast Python:

```
class SneakyHandler(Registry):
    pass # Cannot create unregistered subclass
```

In Python, the hook is *unforgeable*. The language semantics guarantee that creating a subclass triggers `__init_subclass__`. There is no syntax to bypass this. **DOF = 1** by construction.

The objection confuses “can create a registry” with “can guarantee all items are in the registry.” Rust enables the former; Python enforces the latter.

. Objection: You Just Need Discipline

Objection: “Real teams maintain consistency through code review, documentation, and discipline. You don't need language features.”

Response: Discipline is the human oracle. The theorem states:

With $\text{DOF} > 1$, consistency requires an external oracle to resolve disagreements.

“Discipline” is exactly that oracle—human memory, review processes, documentation conventions. This is not a counter-argument; it is the theorem restated in different words.

The question is whether the oracle is:

- **Internal** (language-enforced, automatic, unforgeable), or
- **External** (human-maintained, fallible, bypassable)

Language-level SSOT provides an internal oracle. Discipline provides an external one. Both satisfy consistency when they work. The difference is failure mode: language enforcement cannot be forgotten; human discipline can.

. Objection: The Proofs Are Trivial

Objection: “Most of your proofs are just `rfl` (reflexivity). That means they're trivial tautologies, not real theorems.”

Response: When you model correctly, theorems become definitional. This is a feature, not a bug.

Consider: “The sum of two even numbers is even.” In a well-designed formalization, this might be `rfl`—not because it's trivial, but because the definition of “even” was chosen to make the property structural.

That said, not all proofs are `rfl`. The `rust_lacks_introspection` theorem is 40 lines of actual reasoning:

- /1) Assume a hypothetical introspection function exists
- 2) Use `erasure_destroys_source` to show user-written and macro-expanded code produce identical `RuntimeItems`
 - 3) Derive that the function would need to return two different sources for the same item
 - 4) Contradiction

The proof structure (assumption → lemma application → contradiction) is genuine mathematical reasoning, not tautology. The `rfl` proofs establish the scaffolding; the substantive proofs build on that scaffolding.

. Objection: Real Codebases Don't Need Formal DOF

Objection: “Nobody actually needs Lean-enforced DOF guarantees. Conventions work fine in practice.”

Response: This is an interpretation gap, not a flaw in the proof. We prove:

IF you encode a fact in multiple locations AND require guaranteed consistency, THEN you need either $\text{DOF} = 1$ or an external oracle.

Whether real codebases “need” guaranteed consistency is an engineering judgment outside the scope of formal verification. The same gap exists for:

- **CAP theorem:** Proves partition tolerance forces trade-off. Whether your system needs strong consistency is judgment.
- **Rice's theorem:** Proves semantic properties are undecidable. Whether you need decidable analysis is judgment.
- **Halting problem:** Proves general termination is undecidable. Whether your programs need termination guarantees is judgment.

The theorem characterizes what is *logically required*. Application to specific codebases requires human interpretation. This is philosophy, not mathematics, and lies outside the proof's scope.